

지난 시간 리뷰

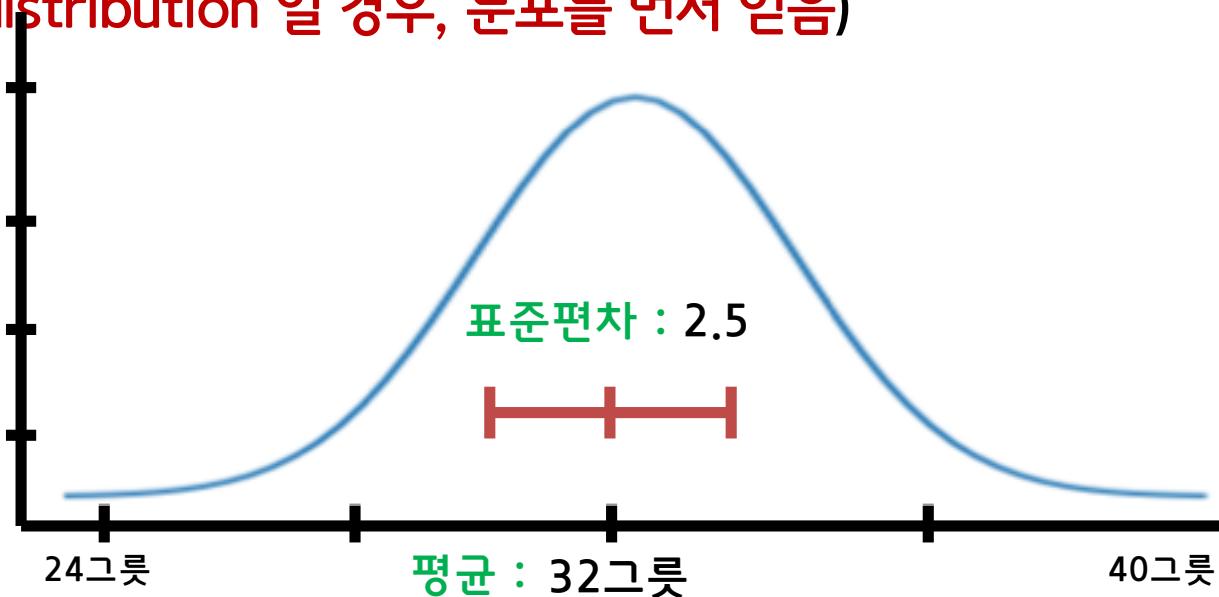
모두의연구소
박은수 책임연구원

Probability vs Likelihood



Probability

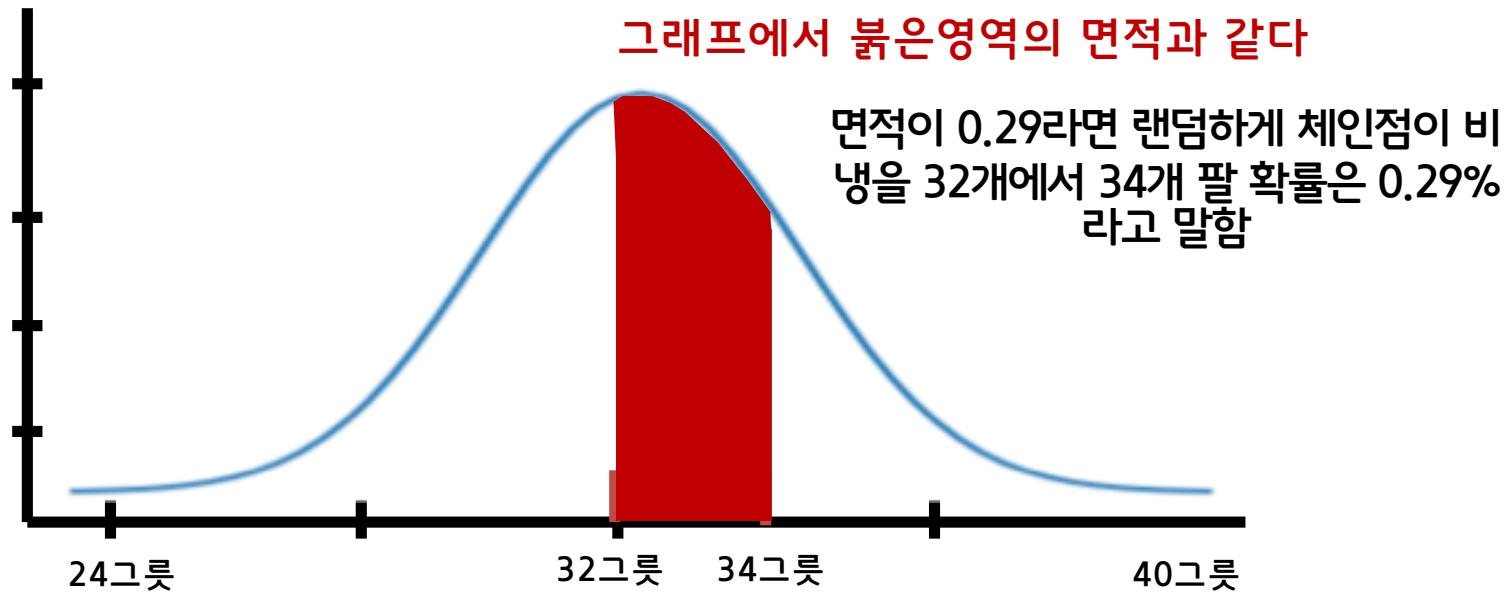
- ‘딥러닝 면옥’ 체인점들의 하루 비빔냉면 판매량 (**normal distribution** 일 경우, **분포를 먼저 얻음**)



Probability vs Likelihood

Probability

랜덤하게 선택한 체인점이 하루에 비빔면을 32개에서 34개 판매할 확률 :



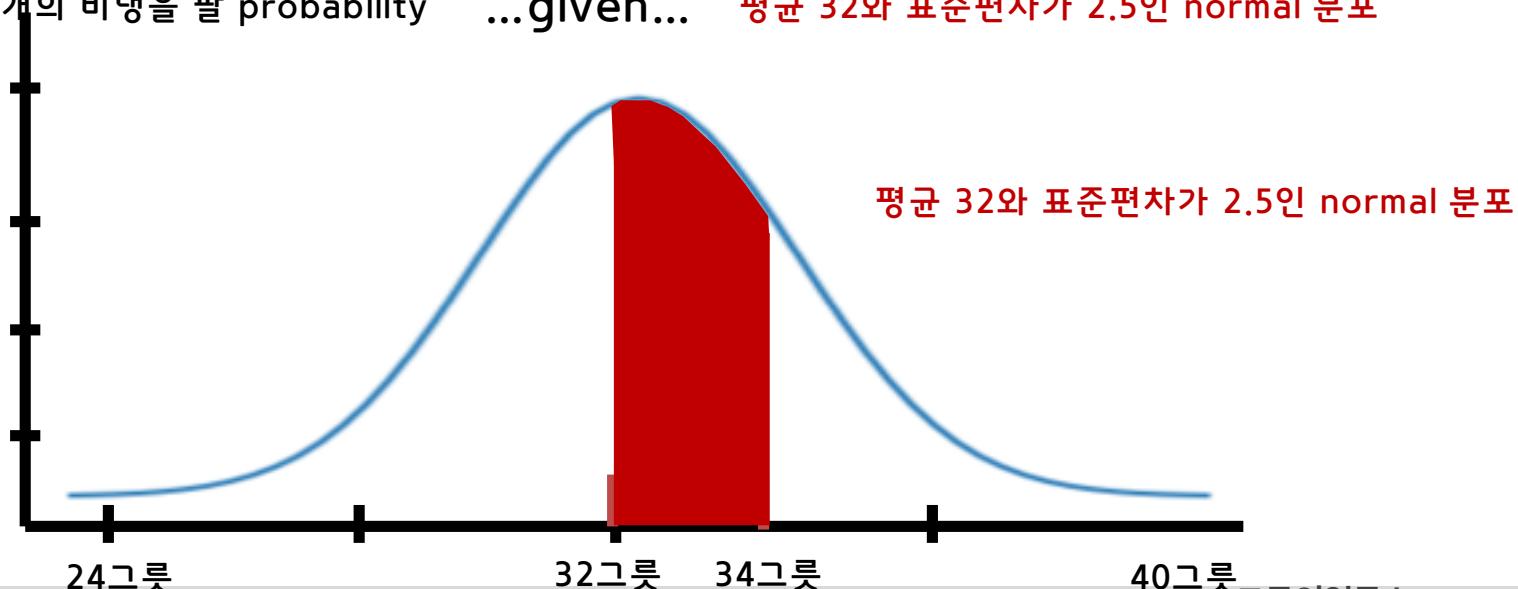
Probability vs Likelihood

Probability

수학적으로 표현해 보면

$$P(\text{32개에서 34개의 비냉을 팜} \mid \text{mean} = 32 \text{ and deviation} = 2.5)$$

하루 32개에서 34개의 비냉을 팔 probability ...given... 평균 32와 표준편차가 2.5인 normal 분포



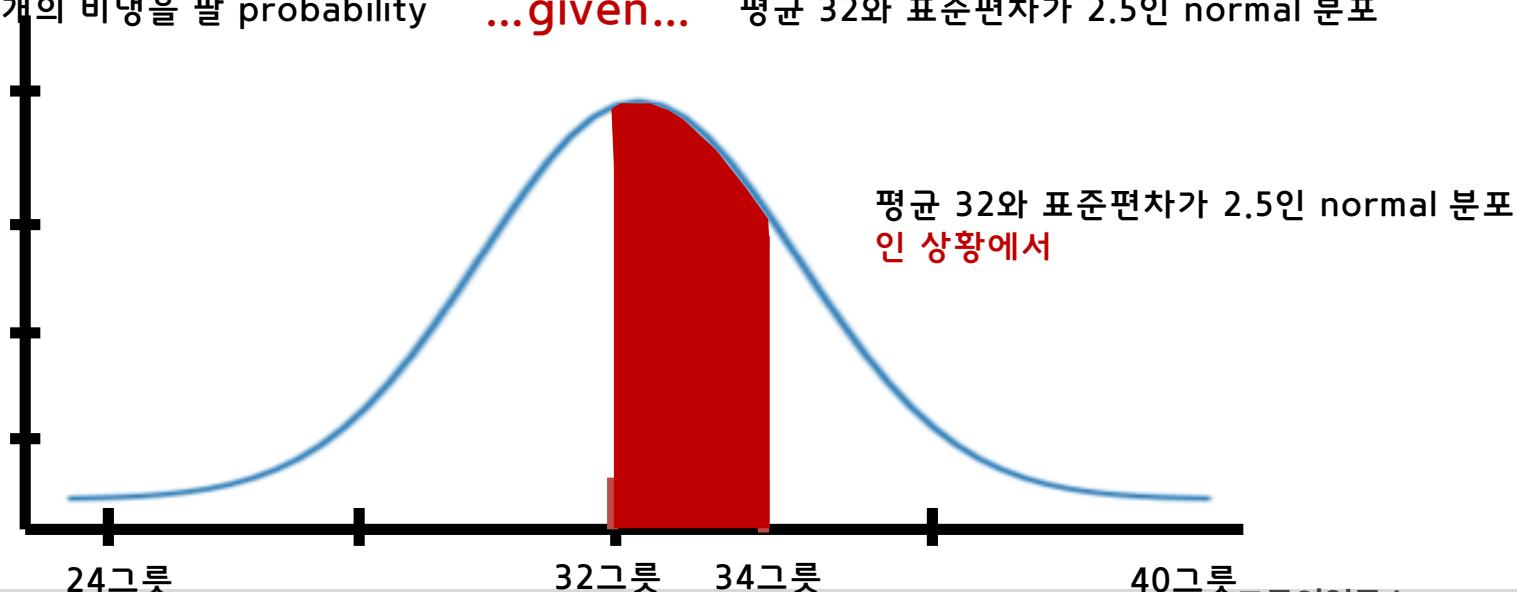
Probability vs Likelihood

Probability

수학적으로 표현해 보면

$$P(\text{32개에서 34개의 비냉을 팜} \mid \text{mean} = 32 \text{ and deviation} = 2.5)$$

하루 32개에서 34개의 비냉을 팔 probability ...given... 평균 32와 표준편차가 2.5인 normal 분포



Probability vs Likelihood

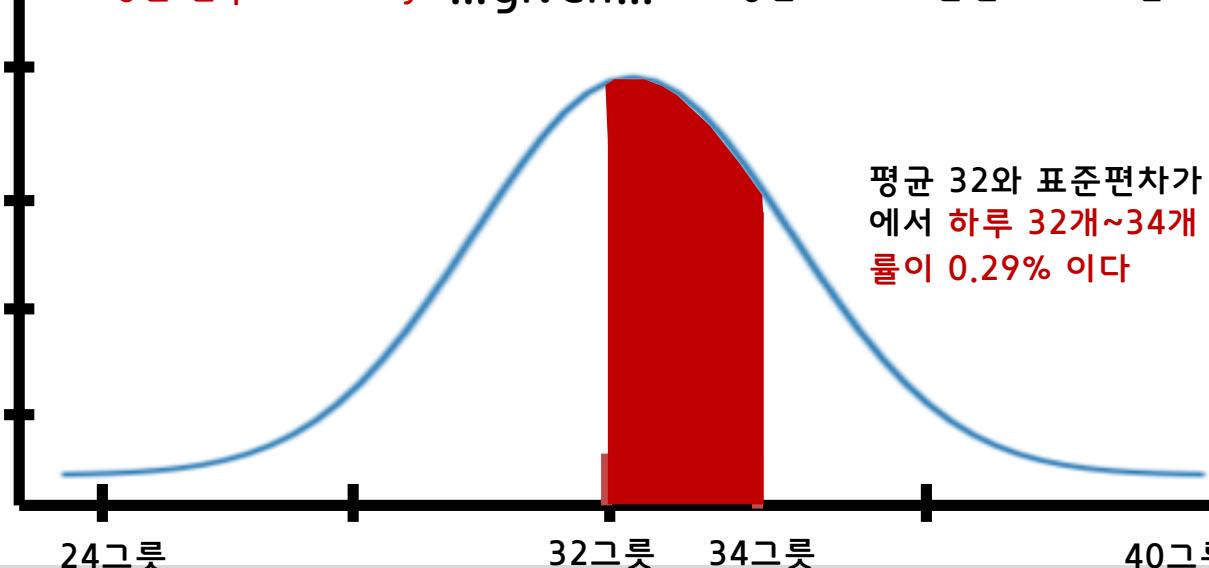
Probability

수학적으로 표현해 보면

$$P(\text{32개에서 34개의 비냉을 팜} \mid \text{mean} = 32 \text{ and deviation} = 2.5)$$

하루 32개에서 34개의 비냉을 팔 probability ... given...

평균 32와 표준편차가 2.5인 normal 분포



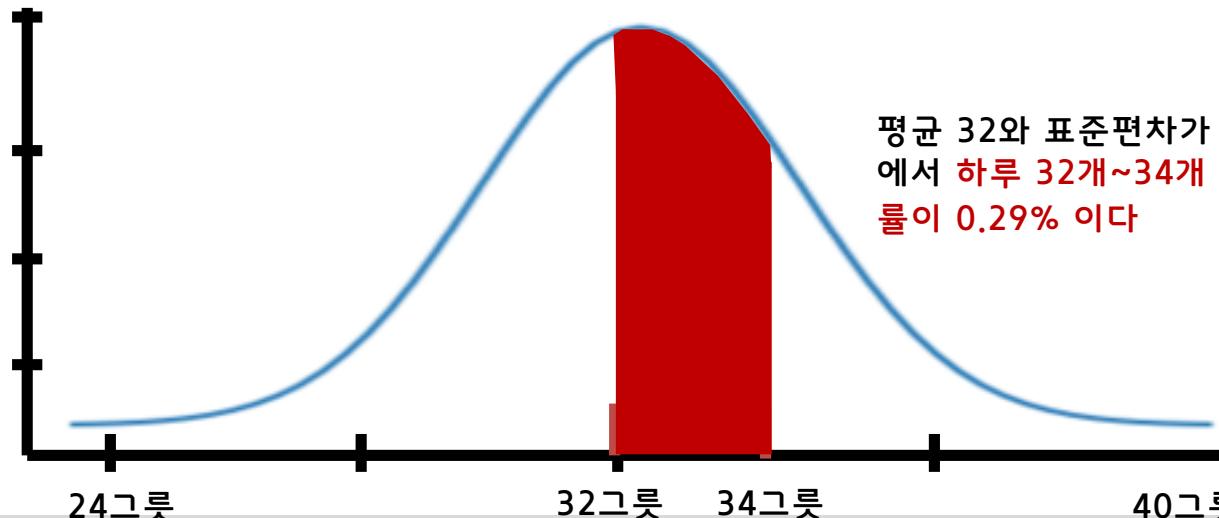
Probability vs Likelihood

Probability

수학적으로 표현해 보면

$$P(32\text{개에서 } 34\text{개의 비냉을 팜} \mid \text{mean} = 32 \text{ and deviation} = 2.5)$$

우리가 다른 비냉판매 수에 관심이 있다면 변경 가능한 부분



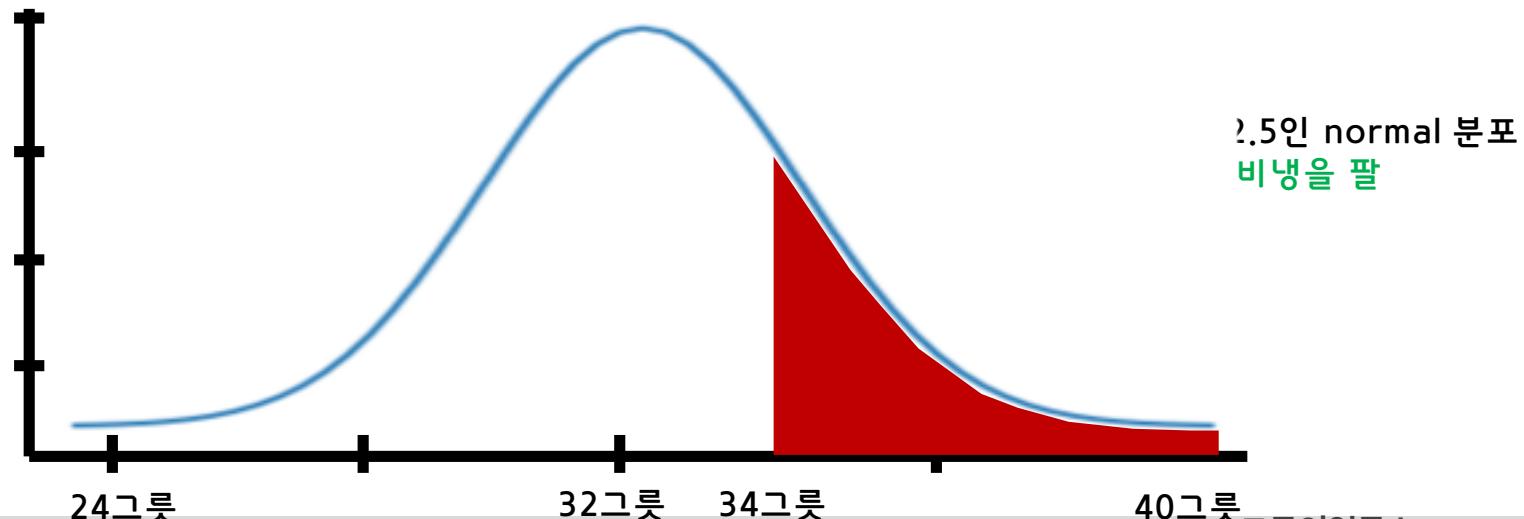
Probability vs Likelihood

Probability

수학적으로 표현해 보면

$$P(\text{비냉판매 수} > 34 \mid \text{mean} = 32 \text{ and deviation} = 2.5)$$

34개 이상 비냉을 판매할 수 체인점의 probability



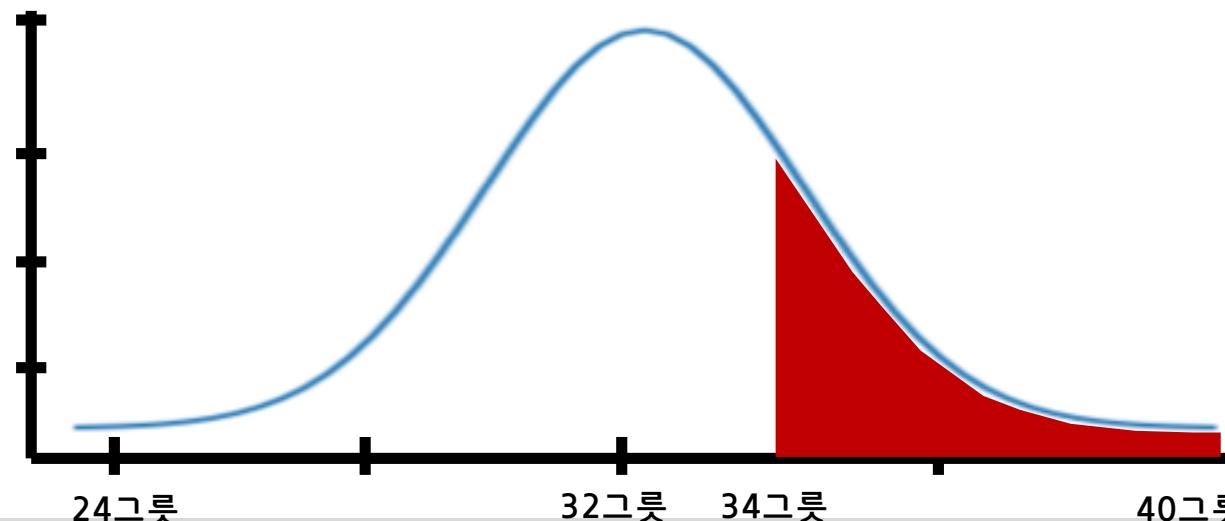
Probability vs Likelihood

Probability

수학적으로 표현해 보면

$$P(\text{비냉판매 수} > 34 \mid \text{mean} = 32 \text{ and } \text{deviation} = 2.5)$$

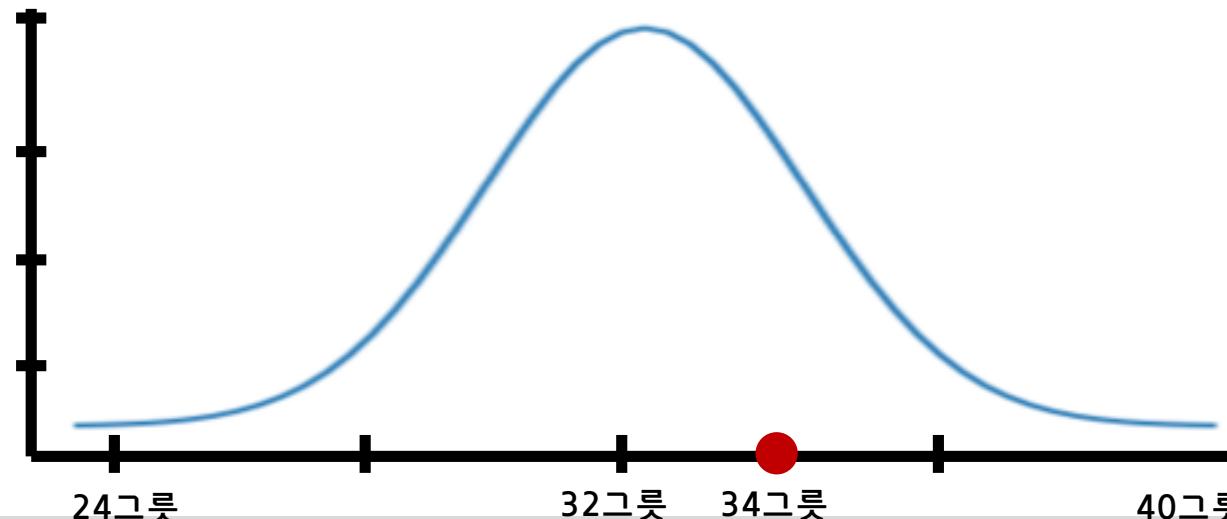
분포를 오른쪽에 고정 시켜 두고 왼쪽을 변경 시켜 새로운 **probability**를 얻을 수 있다



Probability vs Likelihood

Likelihood

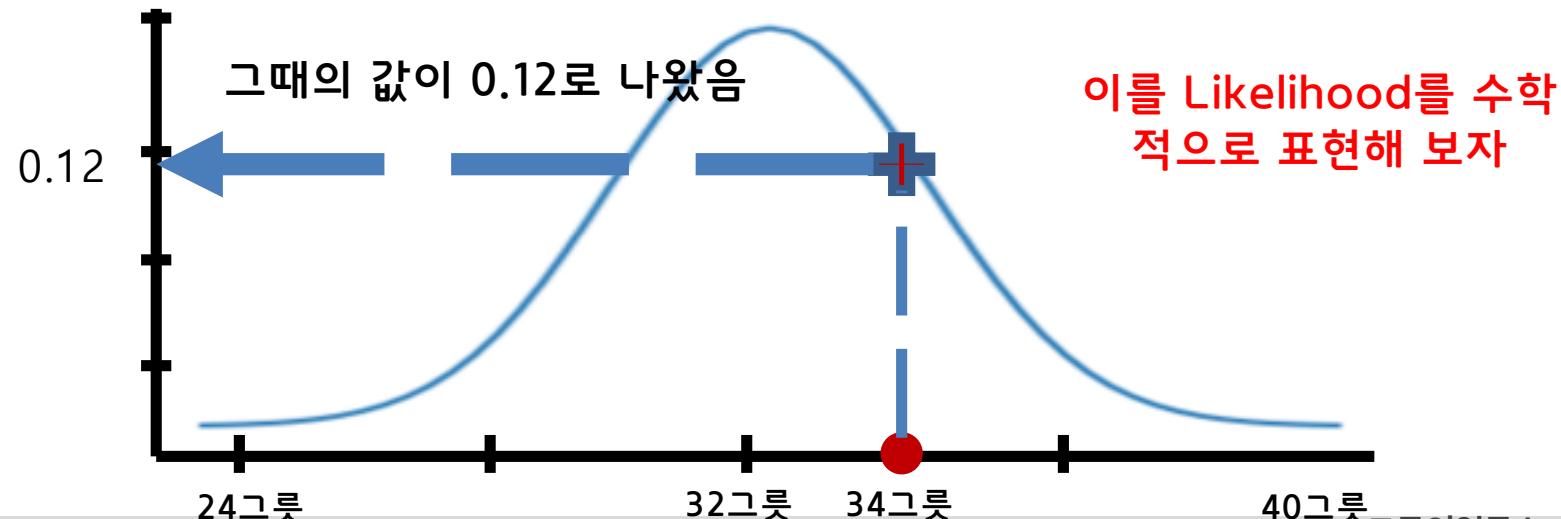
먼저 체인점의 하루 평균 비냉판매량을 측정함 : 34그릇이 나왔다고 가정



Probability vs Likelihood

Likelihood

먼저 체인점의 하루 평균 비냉판매량을 측정함 : 34그릇이 나왔다고 가정

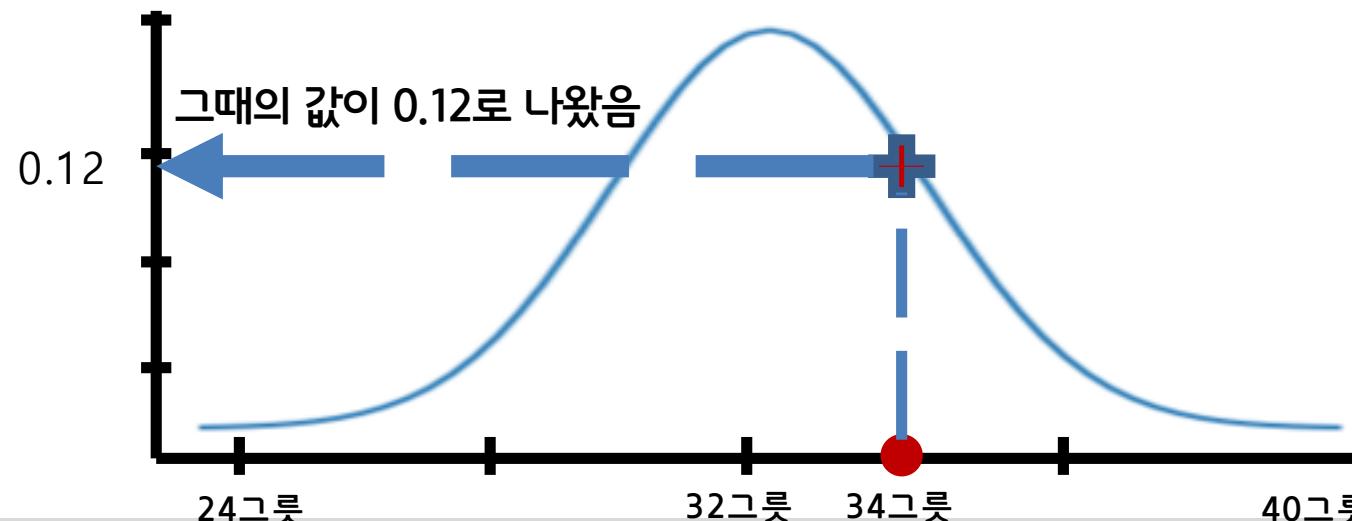


Probability vs Likelihood

Likelihood

$L(\text{mean} = 32 \text{ and } \text{deviation} = 2.5 | \text{비냉판매 수 } 34 \text{ 그릇})$

비냉이 34그릇 팔렸을

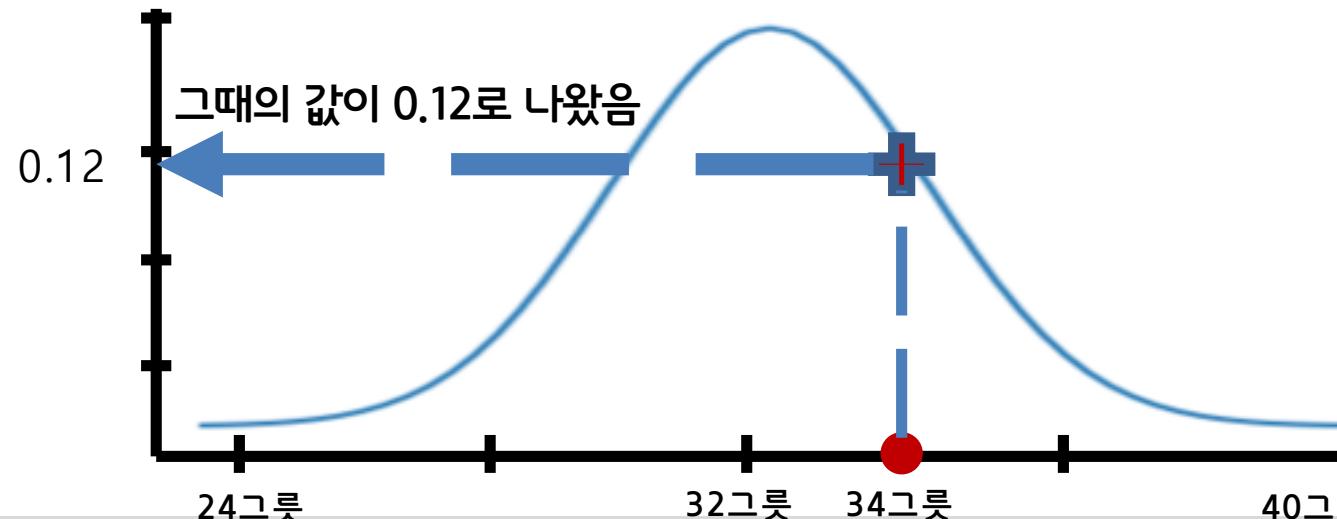


Probability vs Likelihood

Likelihood

$L(\text{mean} = 32 \text{ and } \text{deviation} = 2.5 | \text{비냉판매 수 } 34 \text{ 그릇})$

비냉이 34그릇 팔렸을 때

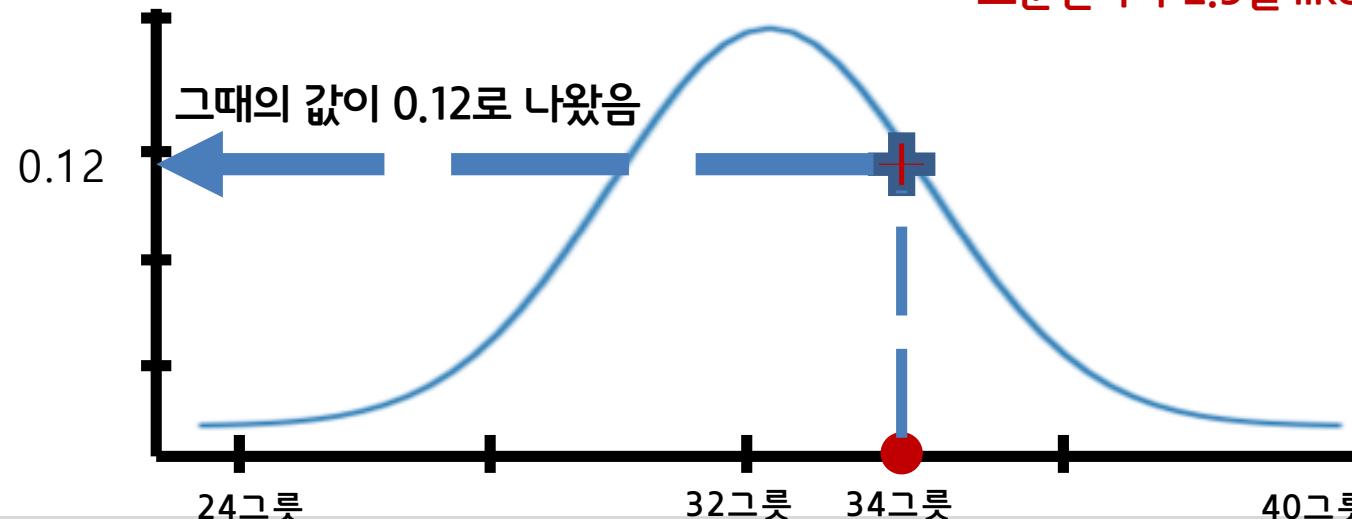


Probability vs Likelihood

Likelihood

$L(\text{mean} = 32 \text{ and } \text{deviation} = 2.5 | \text{비냉판매 수 } 34 \text{ 그릇})$

비냉이 34그릇 팔렸을 때, 평균이 32 표준편차가 2.5일 likelihood는 0.12임.

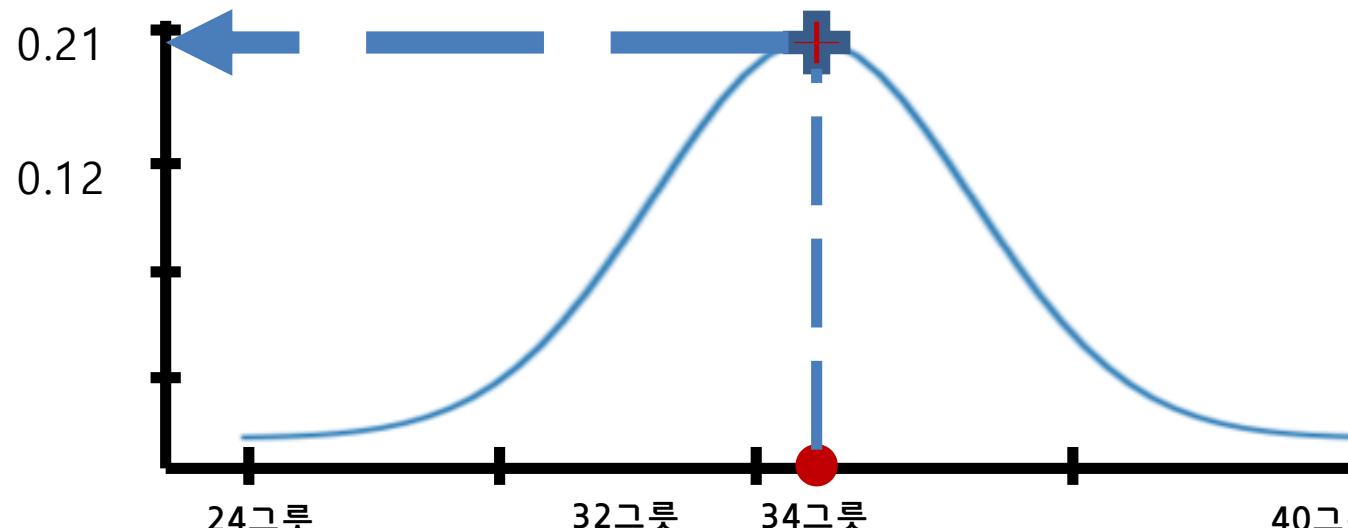


Probability vs Likelihood

Likelihood

$L(\text{mean} = 34 \text{ and deviation} = 2.5 | \text{비냉판매 수 } 34 \text{ 그릇})$

평균이 34로 변경 될 경우의 likelihood는 0.21로 증가 함



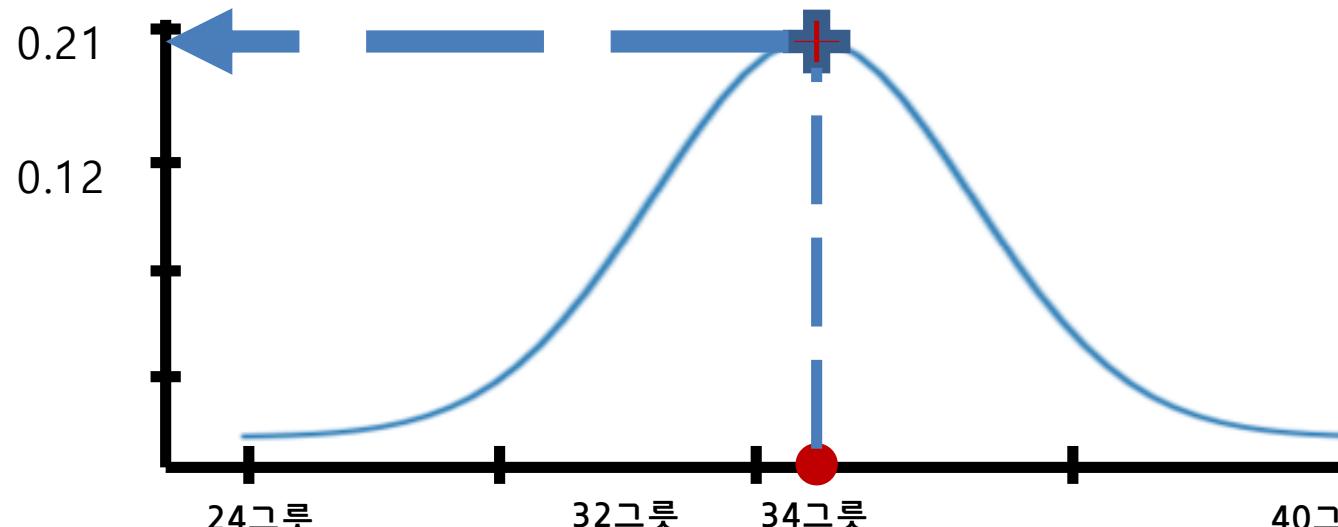
Probability vs Likelihood

Likelihood

Measurement(측정)은 고정

$L(\text{mean} = 34 \text{ and deviation} = 2.5 | \text{비냉판매 수 } 34 \text{ 그릇})$

Distribution을 바꾸면 likelihood가 변경 됨

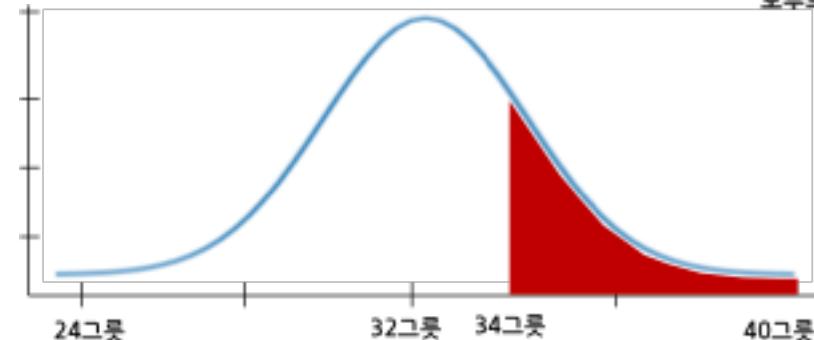


Probability vs Likelihood

Probability : 고정된 분포에서의 우리가 원하는 영역의 면적

$$P(\text{ data } | \text{ distribution })$$

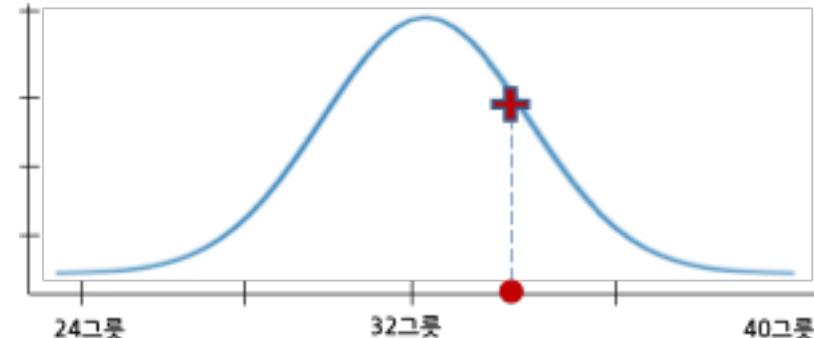
Probability of data given distribution



Likelihood : 고정된 데이터 포인트에서 변형 가능한 분포를 이용해 측정한 Y축 값

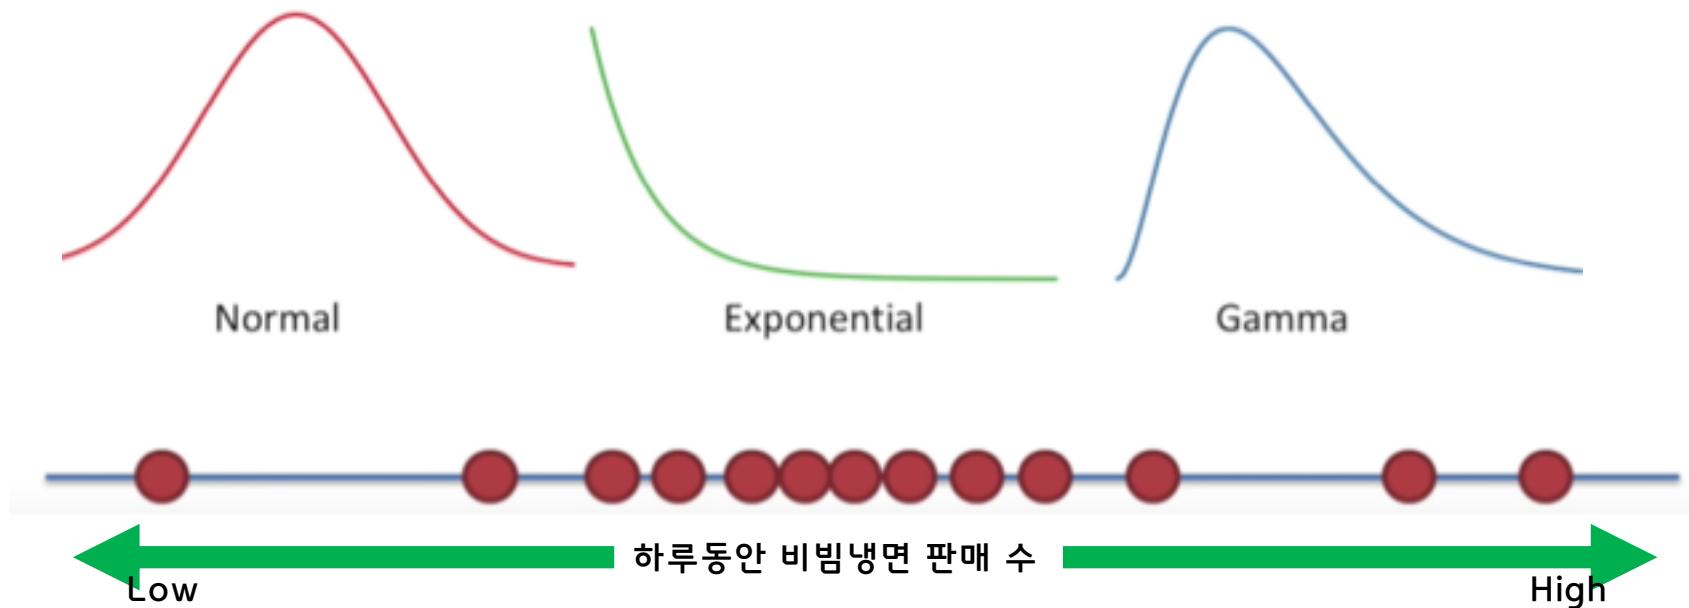
$$L(\text{ distribution } | \text{ data })$$

Likelihood of distribution given data



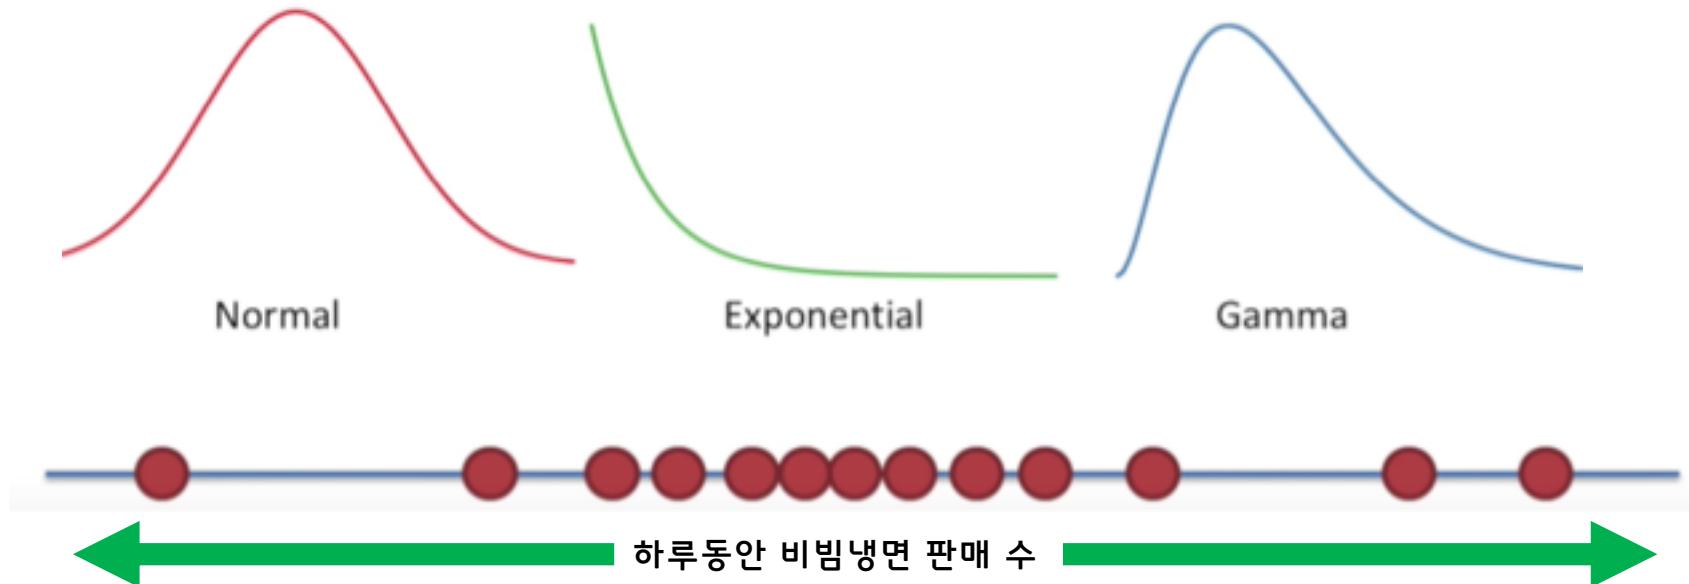
Maximum Likelihood

Maximum likelihood의 목표는 데이터를 가장 잘 표현하는 분포를 찾는 것이다



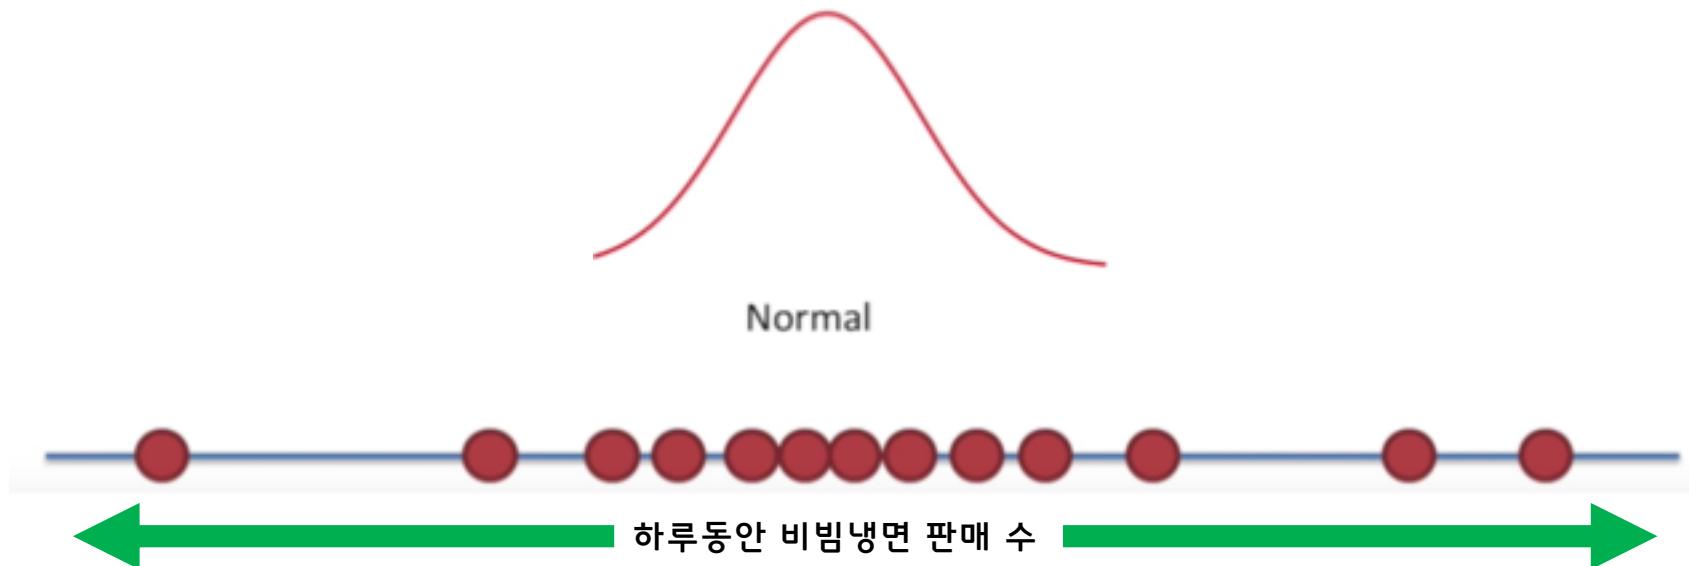
Maximum Likelihood

분포를 찾는 이유 : 다루기 쉬워지고 일반화 될 수 있다



Maximum Likelihood

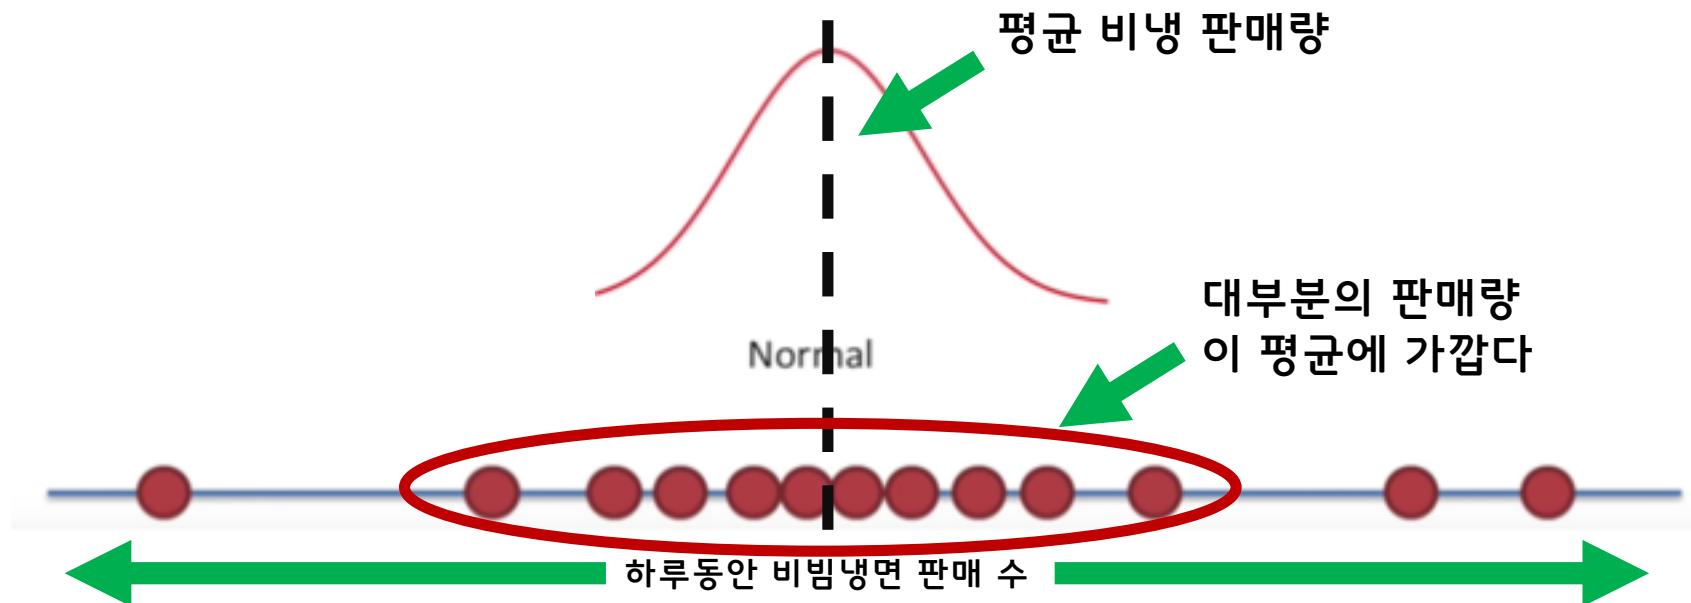
하루 동안의 비빔냉면 판매 수가 normal distribution이라 가정하면



Maximum Likelihood

Normally distributed의 의미 :

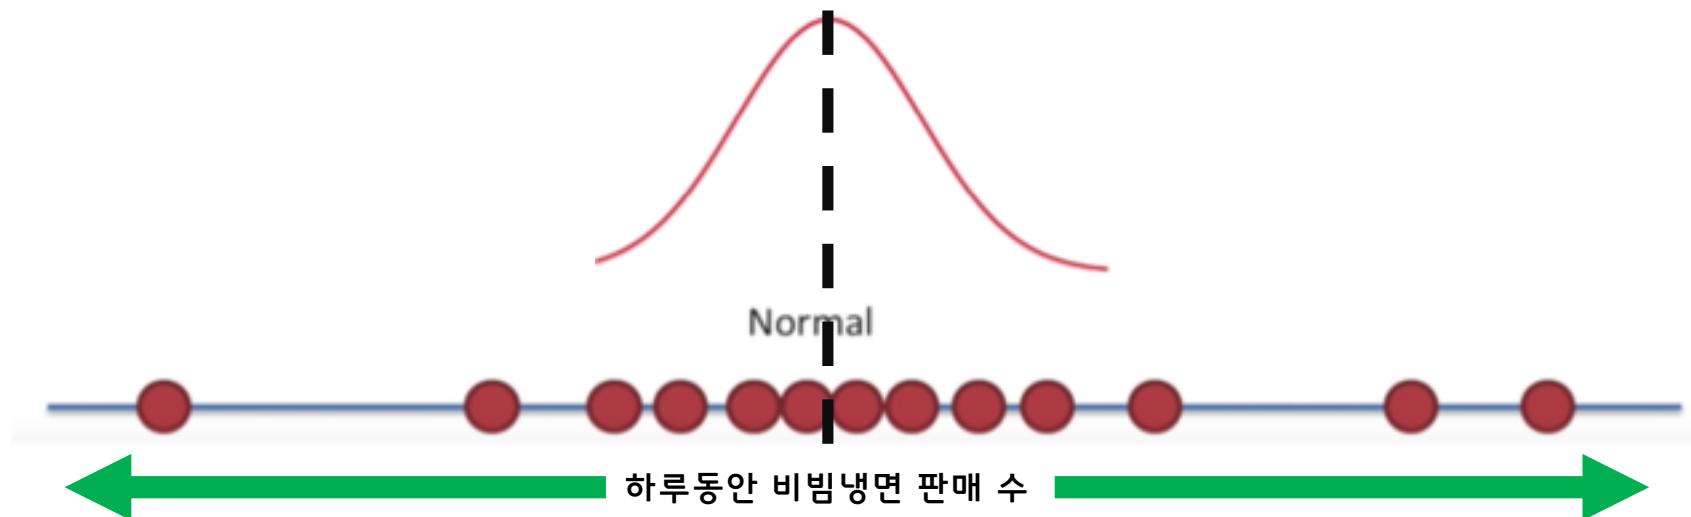
- 1) 대부분의 측정이 평균(mean)에 가깝다



Maximum Likelihood

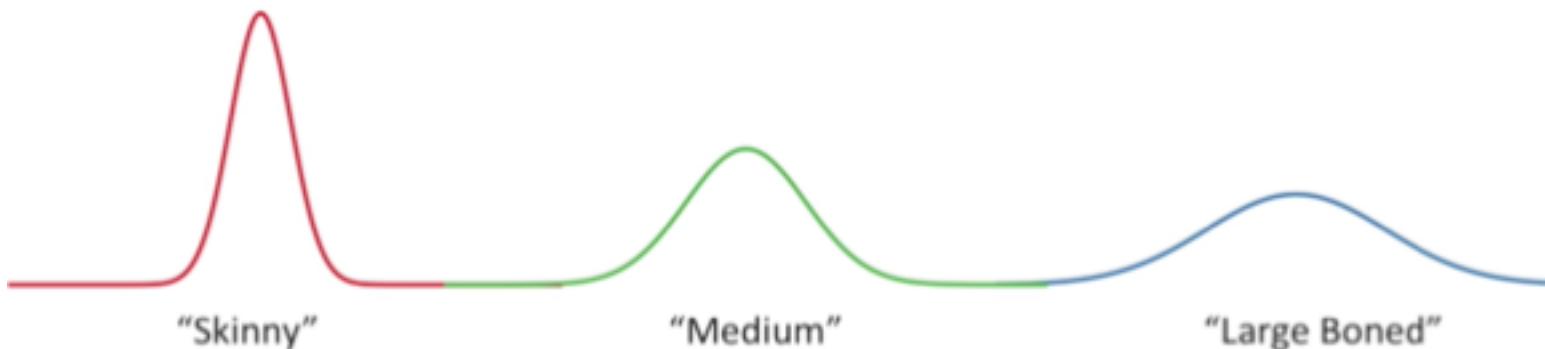
Normally distributed의 의미 :

- 1) 대부분의 측정이 평균(mean)에 가깝다
- 2) 측정결과가 평균을 중심으로 대략 대칭형이다



Maximum Likelihood

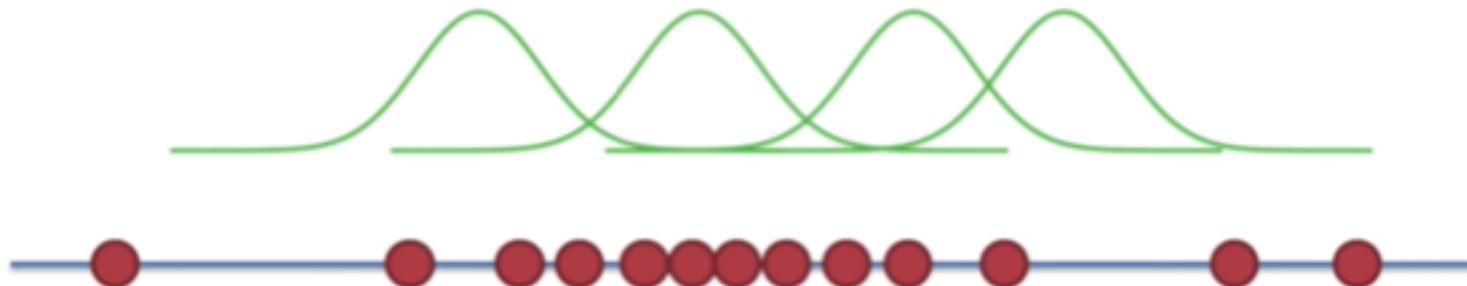
Normally distribution은 여러가지 형태를 가질 수 있습니다



Maximum Likelihood

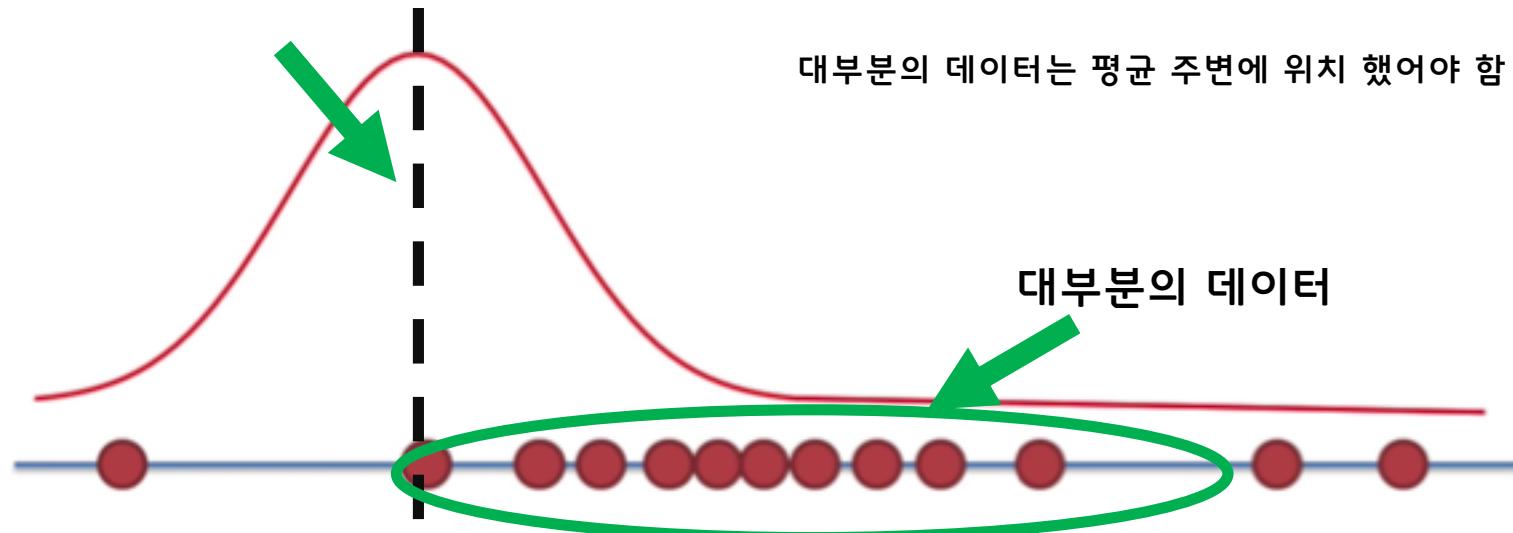


형태를 찾았으면 이 분포를 어디에다 위치시킬지 결정해야 합니다



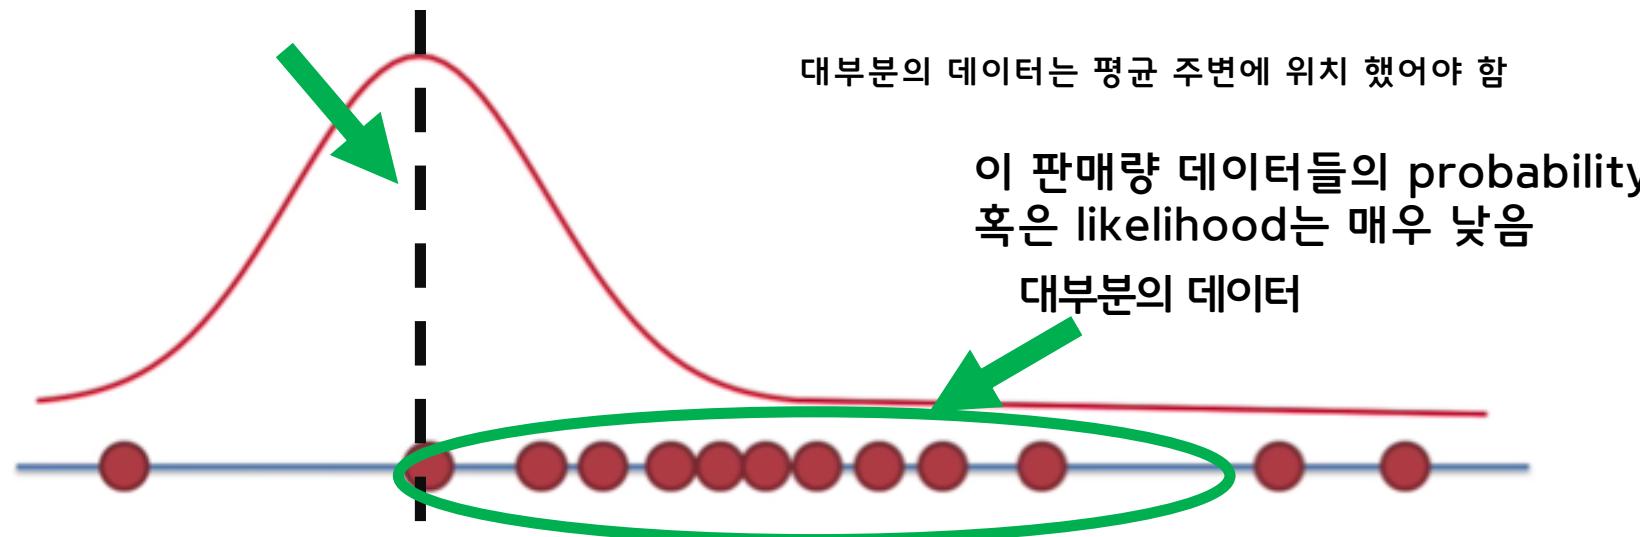
Maximum Likelihood

분포의 평균 (mean value of the distribution)

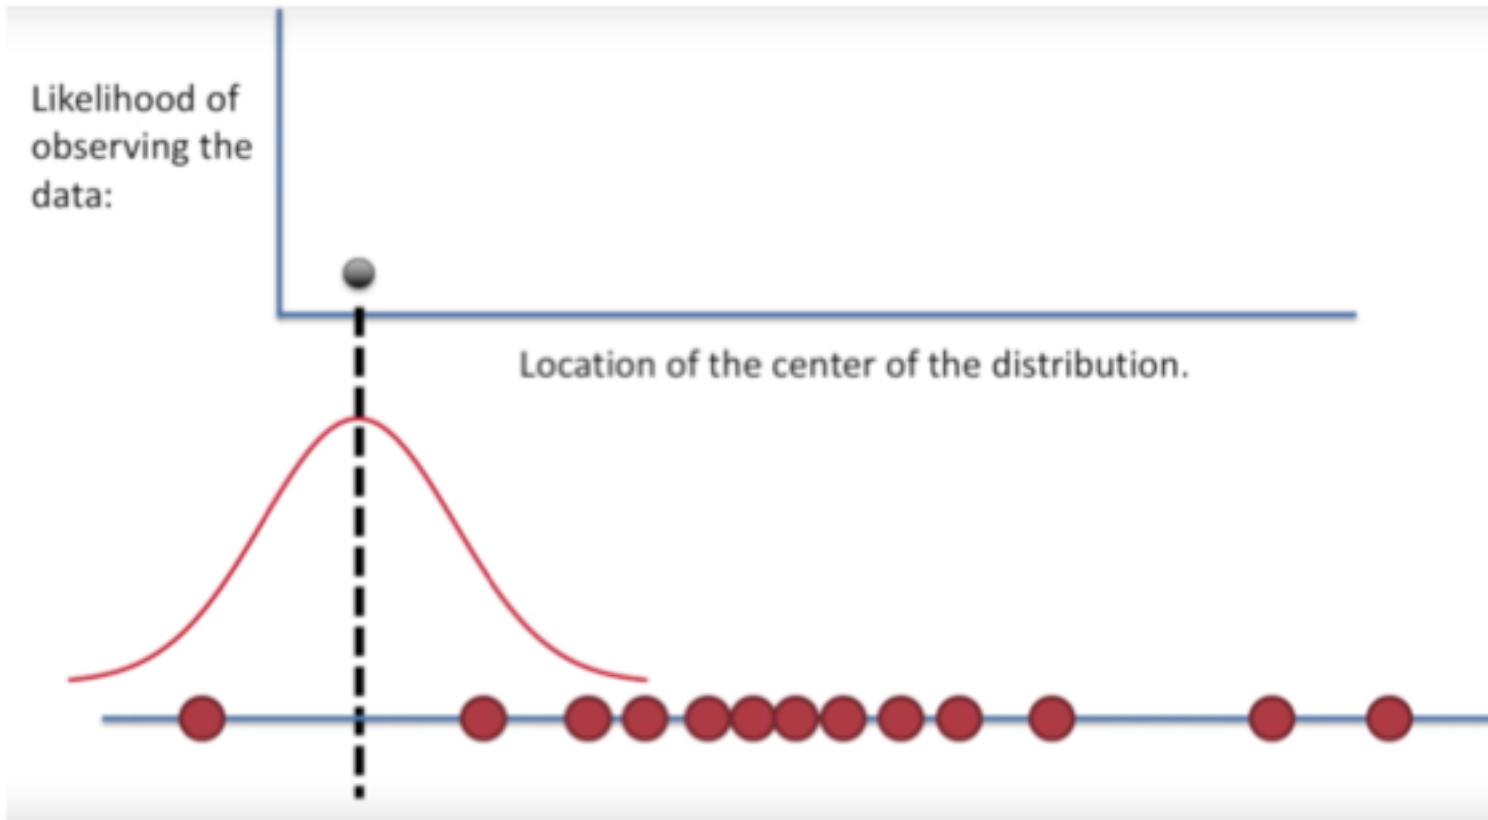


Maximum Likelihood

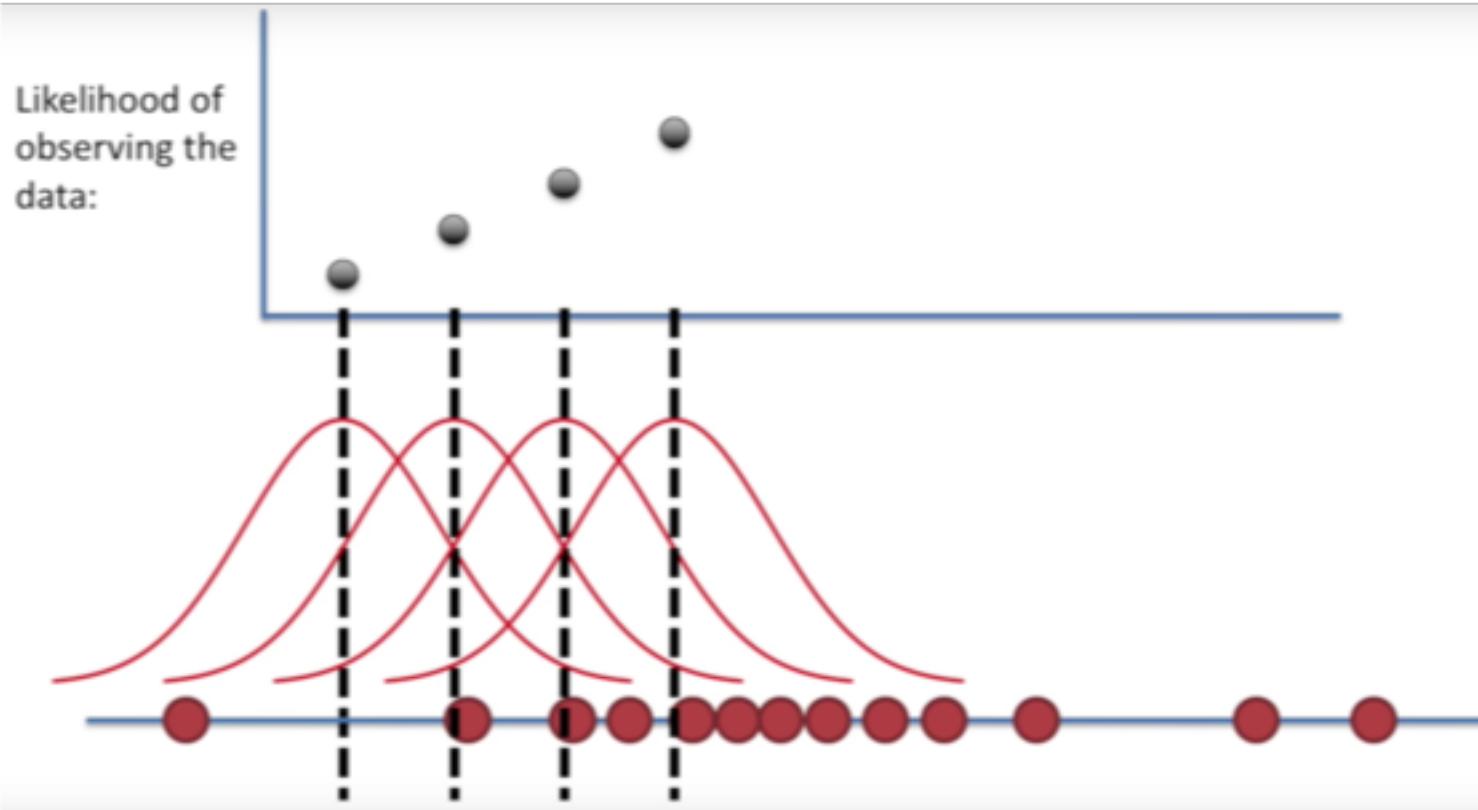
분포의 평균 (mean value of the distribution)



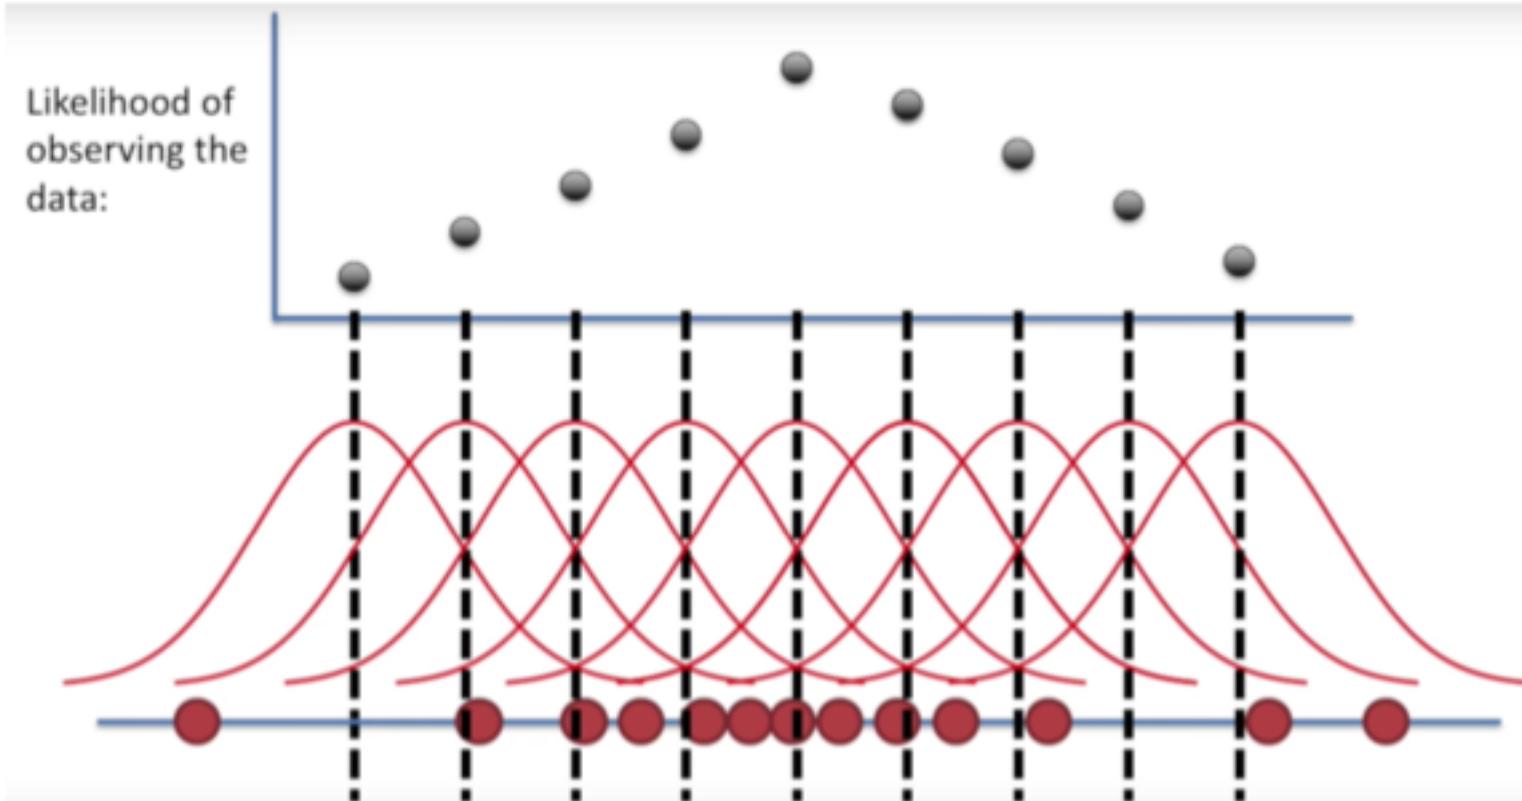
Maximum Likelihood



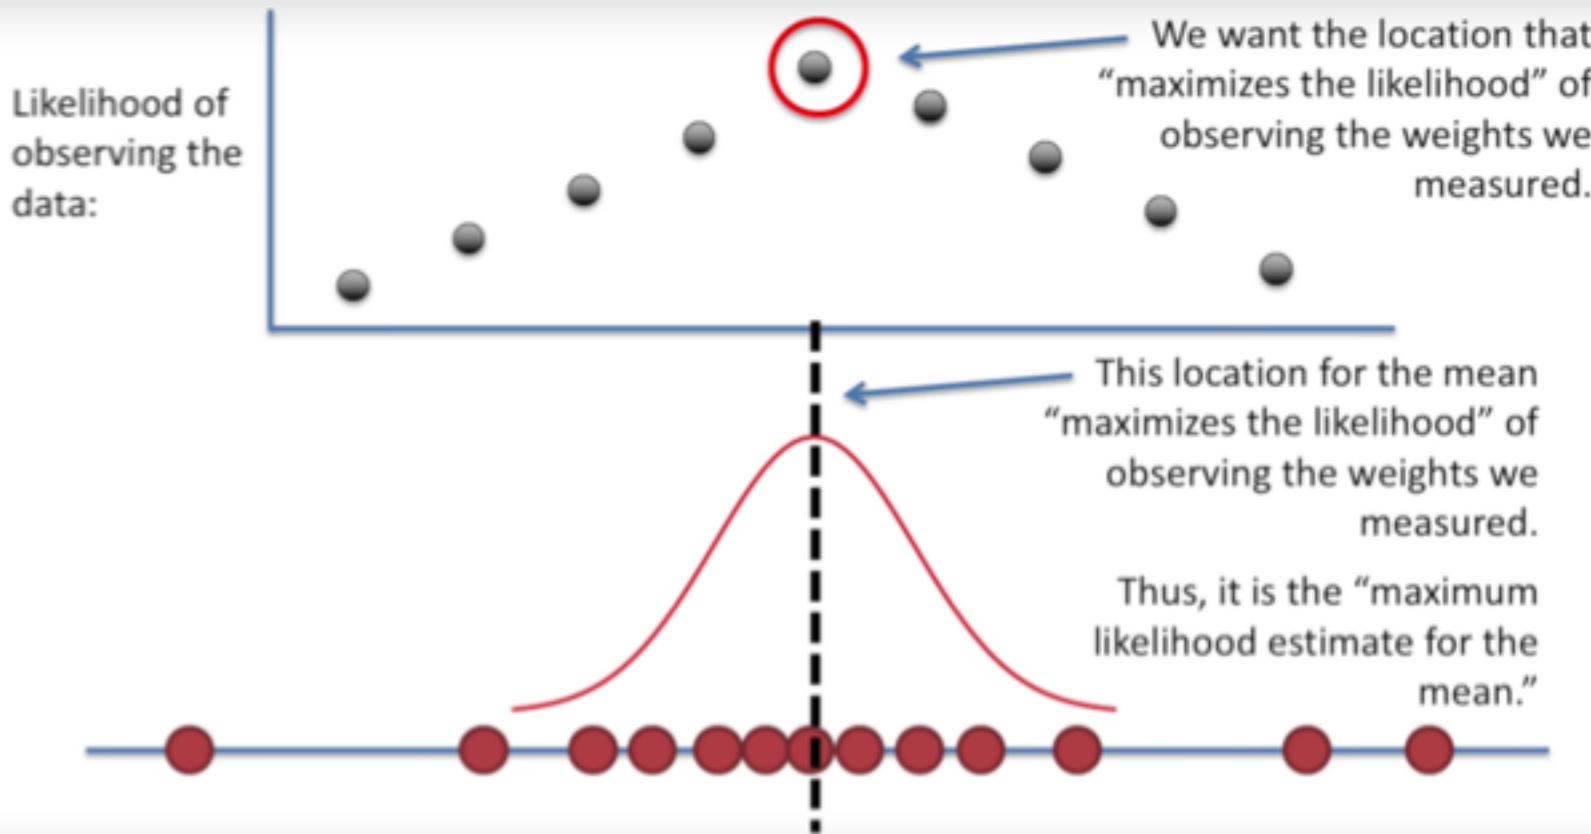
Maximum Likelihood



Maximum Likelihood

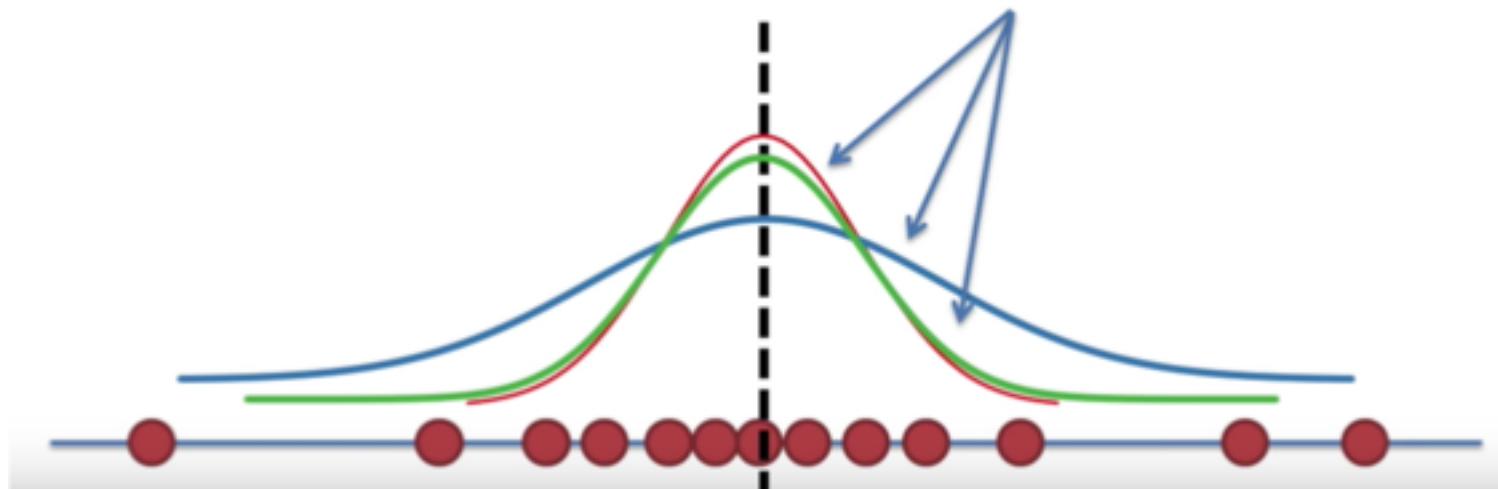


Maximum Likelihood

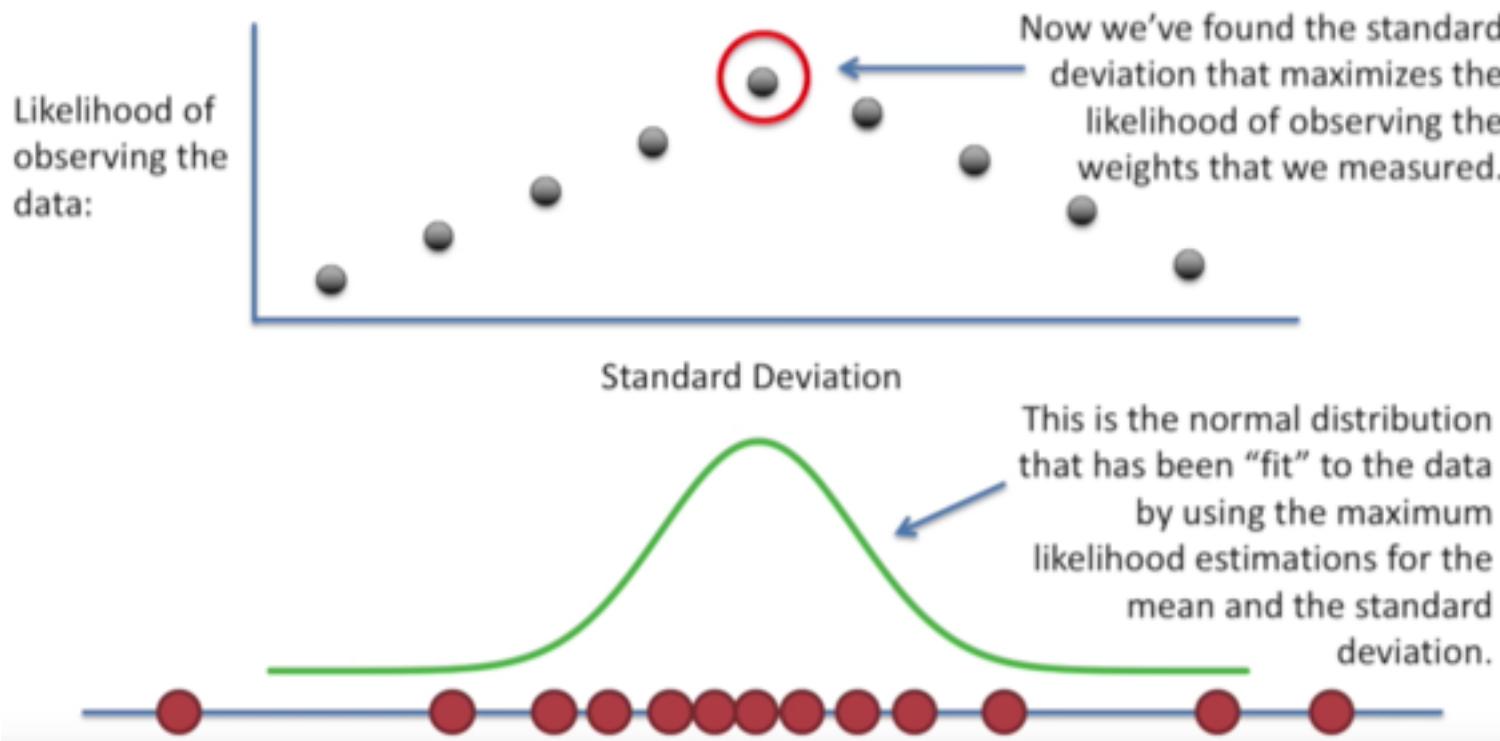


Maximum Likelihood

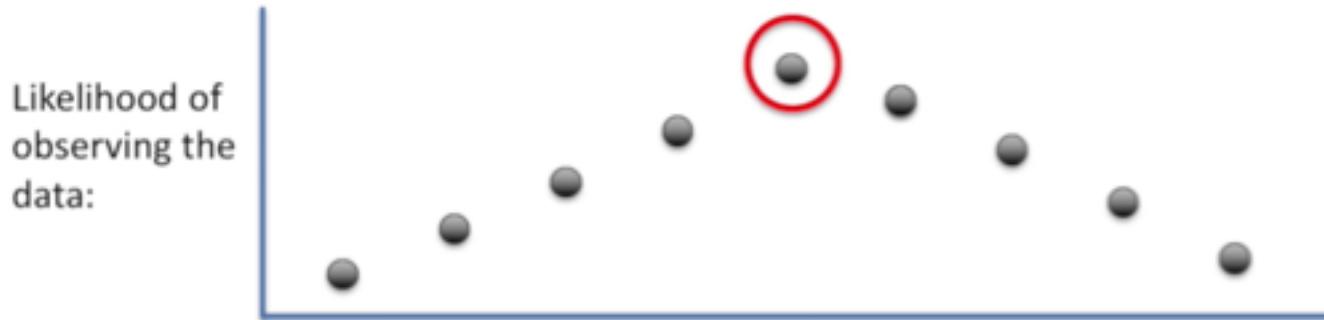
Now we have to figure out the
“maximum likelihood estimate for
the standard deviation....”



Maximum Likelihood

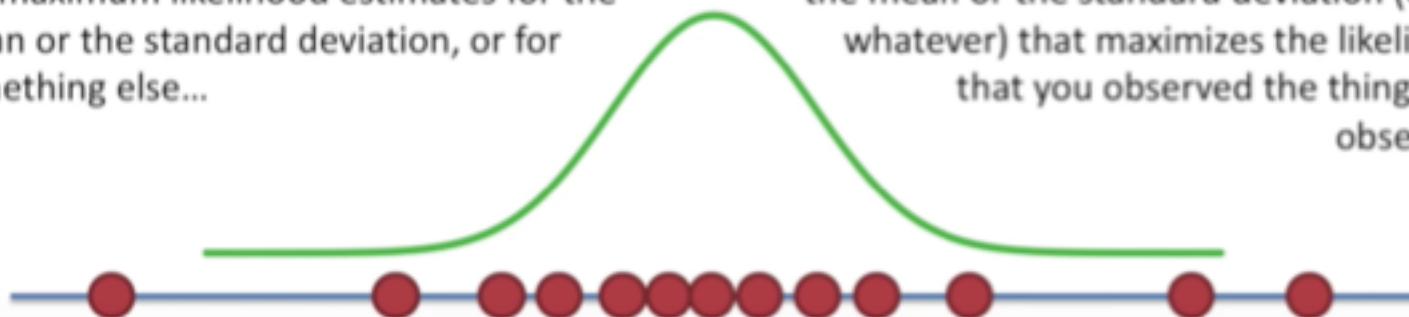


Maximum Likelihood



Now when someone says that they have the maximum likelihood estimates for the mean or the standard deviation, or for something else...

... you know that they found the value for the mean or the standard deviation (or for whatever) that maximizes the likelihood that you observed the things you observed.



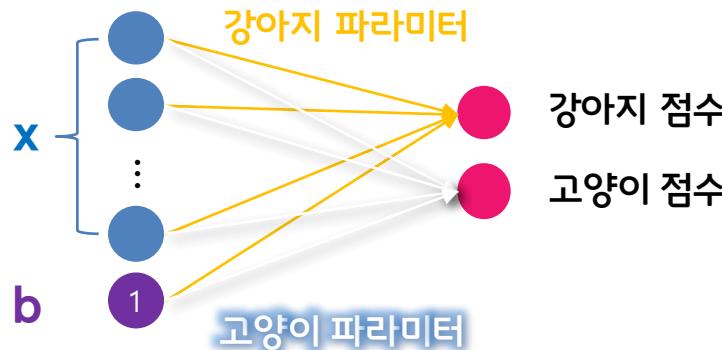
Maximum Likelihood



Terminology alert!

In everyday conversation, “probability” and “likelihood” mean the same thing. However, in Stats-Land, “likelihood” specifically refers to this situation we’ve covered here; where you are trying to find the optimal value for the mean or standard deviation for a distribution given a bunch of observed measurements.

돌아보기 ...



선형분류기

강아지 파라미터
고양이 파라미터

0.2	0.1	...	0.3	0.7
0.7	0.8	...	0.9	0.4

W

32x32x3 영상 =
3,072 픽셀



$$f(x, W) = Wx + b$$

2×1 2×1
 $3,072 \times 1$ $2 \times 3,072$

x

155
200
...
110
78

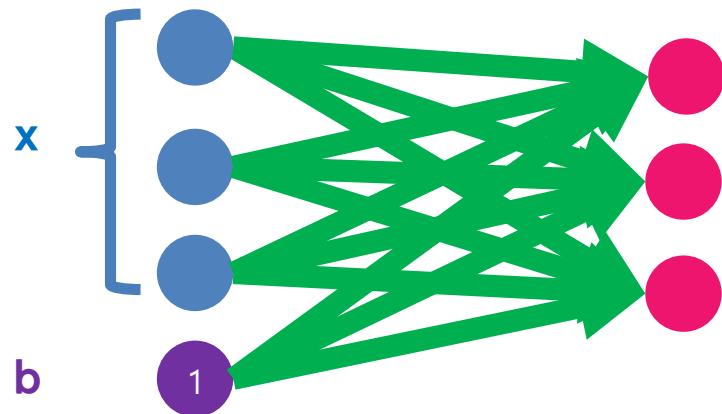
0.1
0.2

0.3
0.7

강아지 점수
고양이 점수

돌아보기 ..

$$(3 \times 4) \quad (4 \times 1)$$
$$W_1 X^T$$

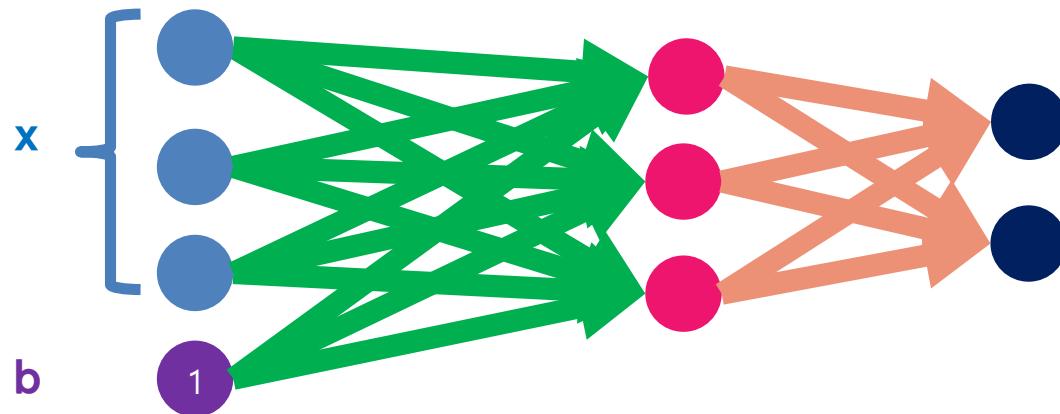


돌아보기 ..

(2x3) (3x4) (4x1)

$W_2 \quad W_1 X^T$

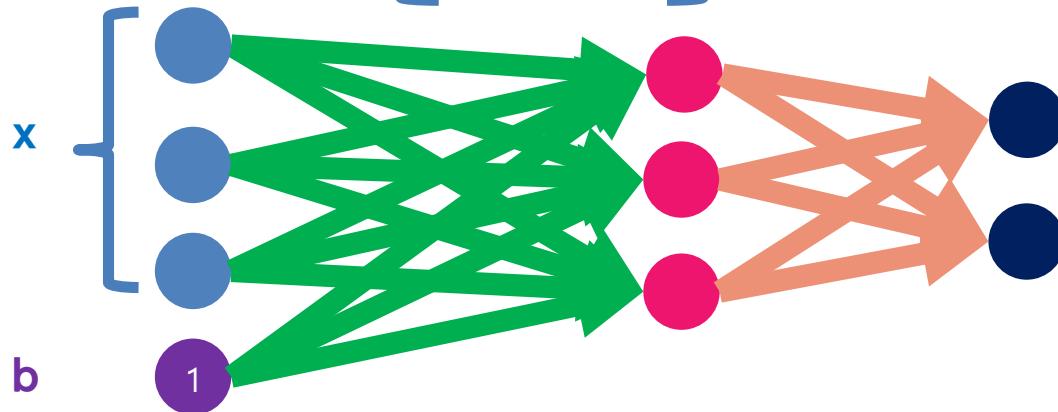
레이어추가



돌아보기 ..

핫 ~

$$\begin{matrix} (2 \times 3) & (3 \times 4) \\ W_2 & W_1 \end{matrix} = \begin{matrix} (2 \times 4) \\ W \end{matrix}$$
$$\left[\begin{matrix} (2 \times 3) & (3 \times 4) \\ W_2 & W_1 \end{matrix} \right]^{(4 \times 1)} X^T$$

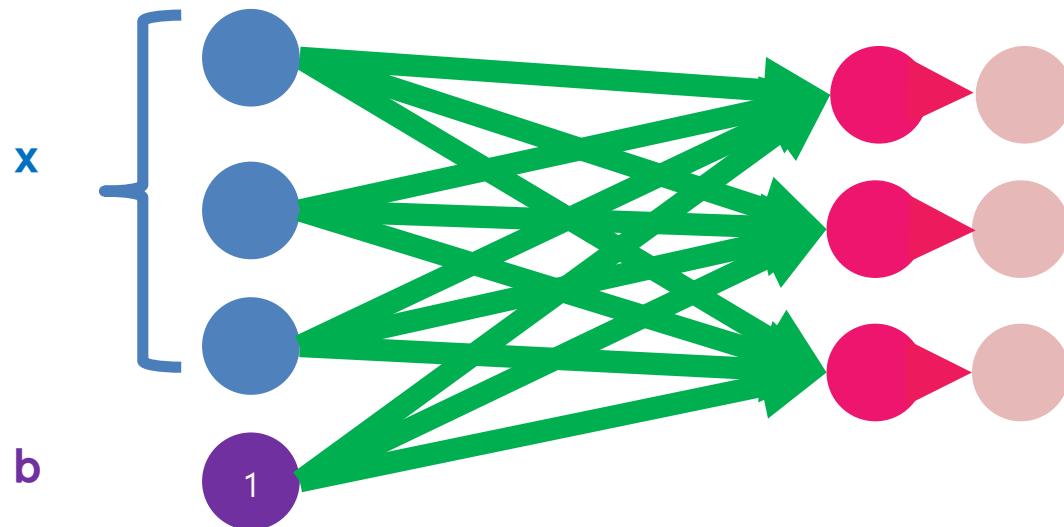


- 두개의 레이어가 하나로 표현
- 레이어를 쌓는 효과가 없어짐

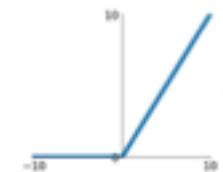
비선형 연산이 필요

돌아보기 ..

$$\max(0, W_1 X^T)$$



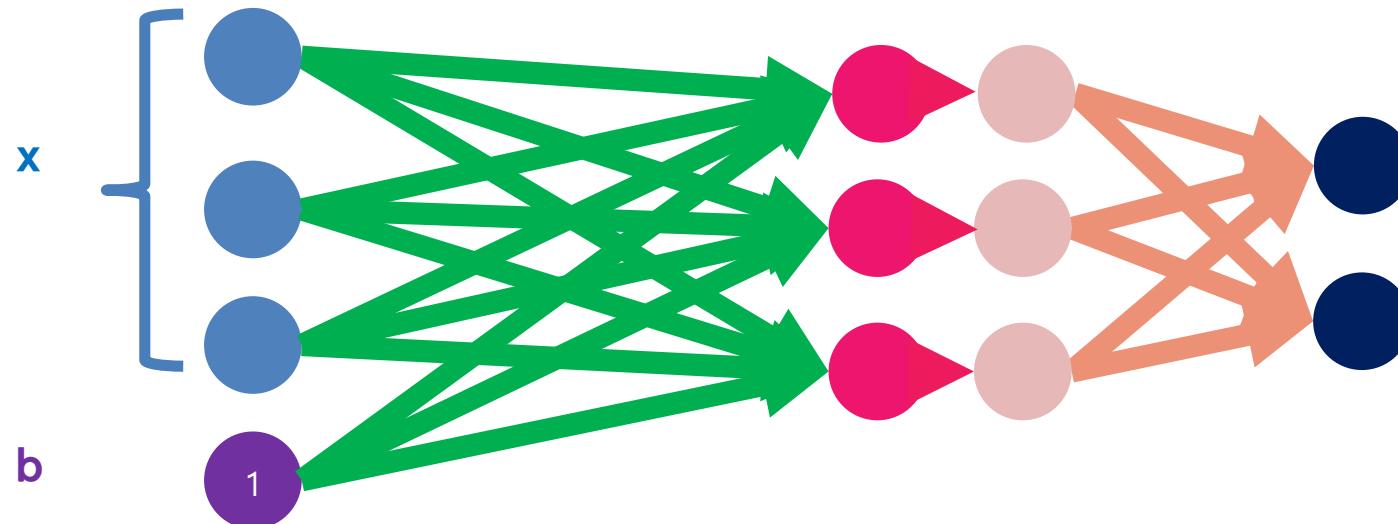
ReLU
 $\max(0, x)$



돌아보기 ..

레이어 추가

$$W_2 \times \max(0, W_1 X^T)$$

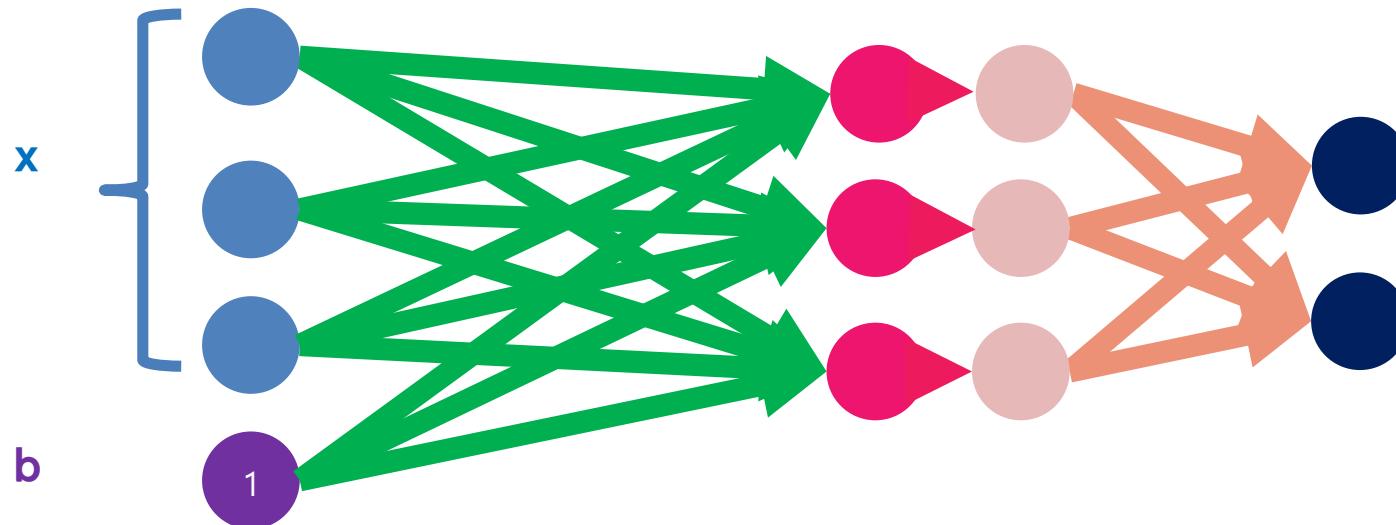


돌아보기 ..

1. Score function

레이어 추가

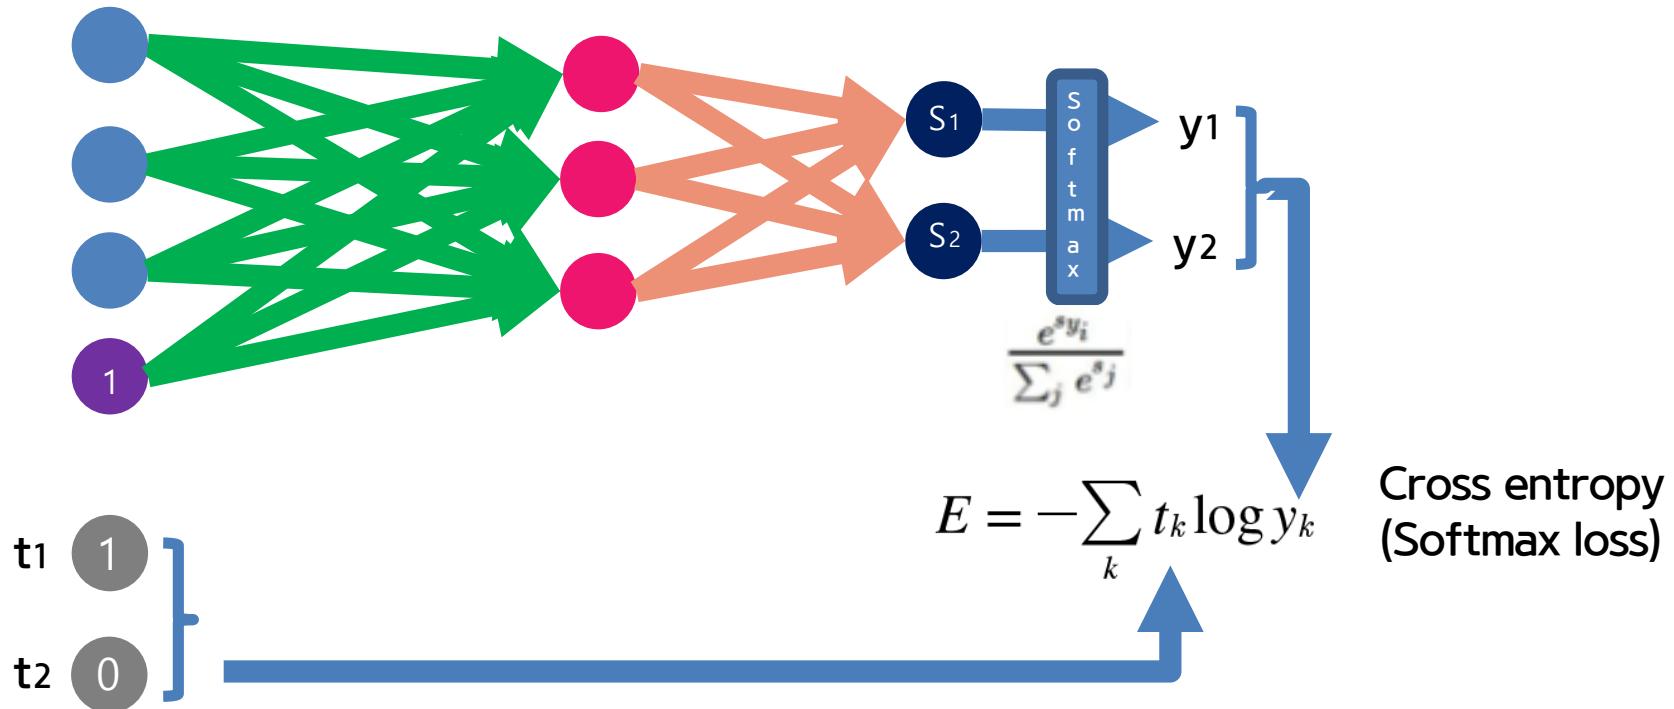
$$W_2 \times \max(0, W_1 X^T)$$



돌아보기 ..

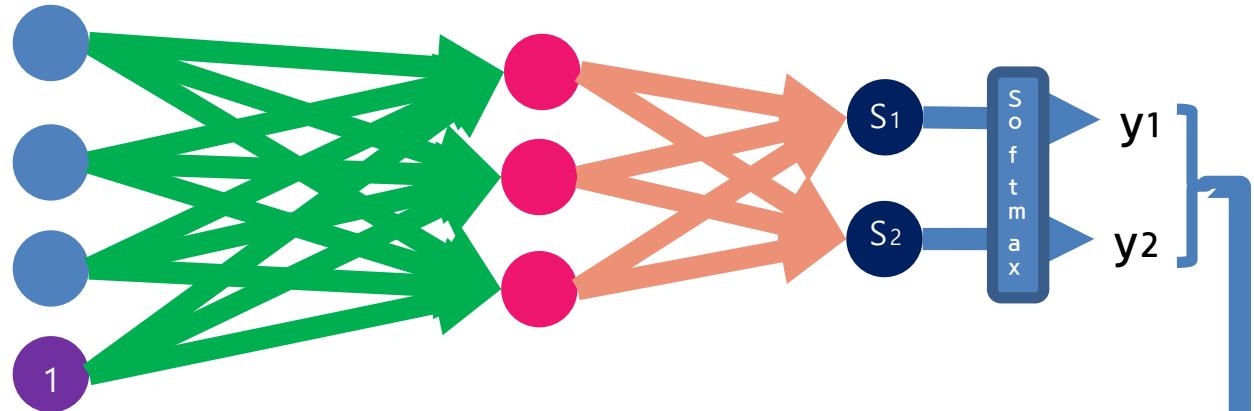
1. Score function
2. Loss function

Loss function



돌아보기 ..

1. Score function
2. Loss function
3. Optimization



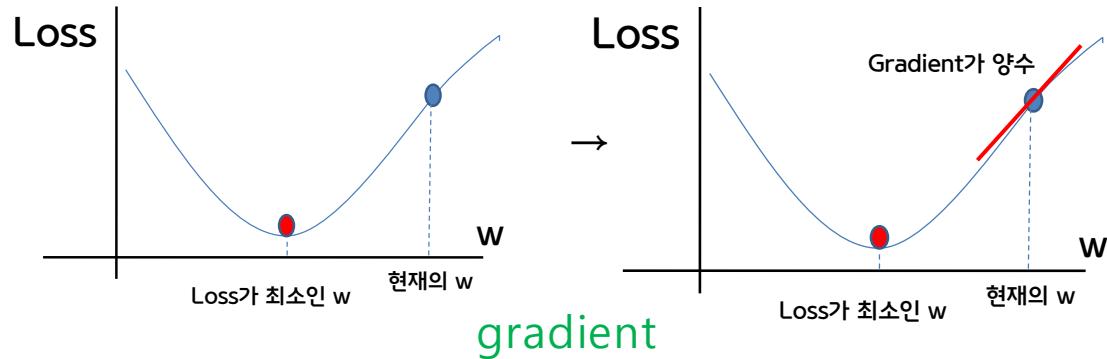
Loss를 줄이는 방법

Gradient Descent

$$Loss \leftarrow E = -\sum_k t_k \log y_k$$

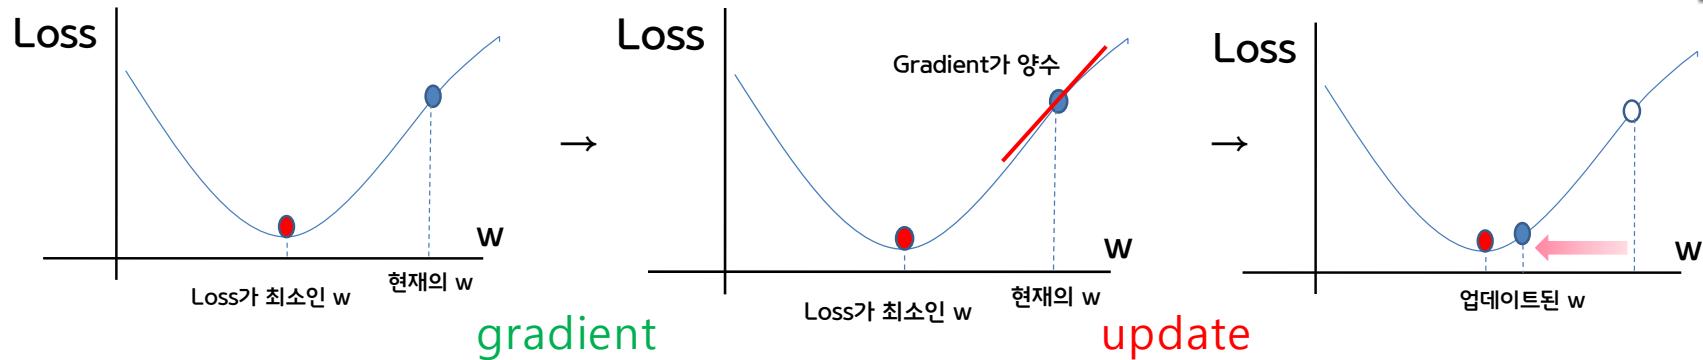
Cross entropy
(Softmax loss)

Gradient Descent



$$w = w - \eta \frac{\partial L}{\partial w}$$

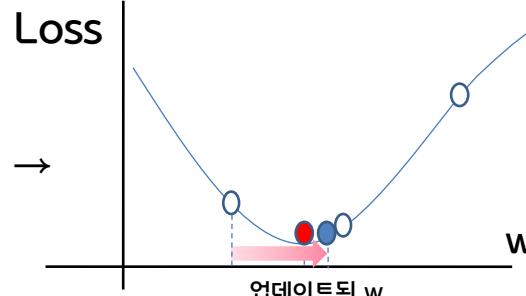
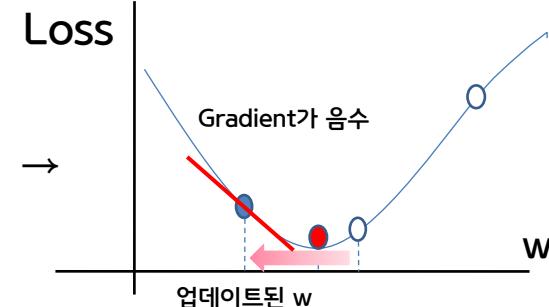
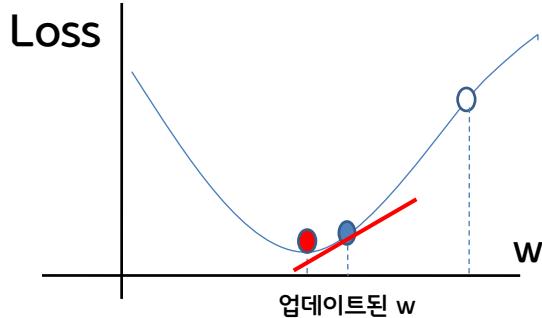
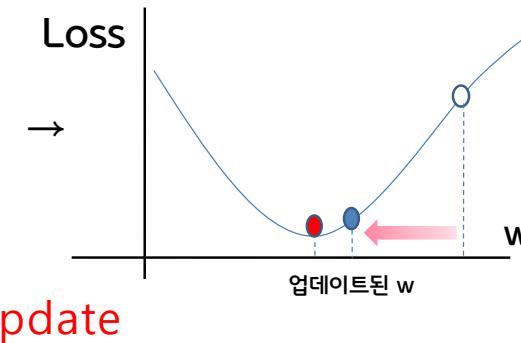
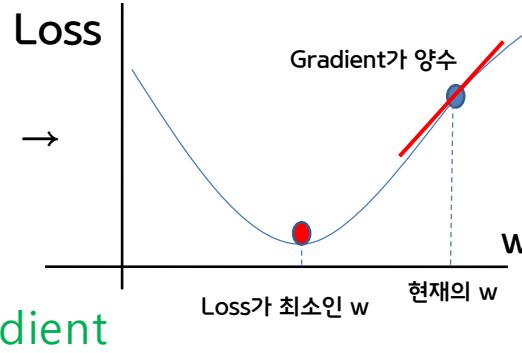
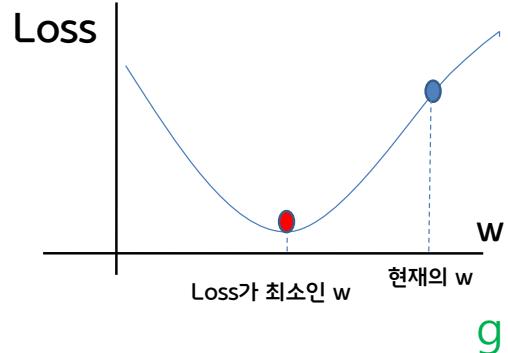
Gradient Descent



$$w = w - \eta \frac{\partial L}{\partial w}$$

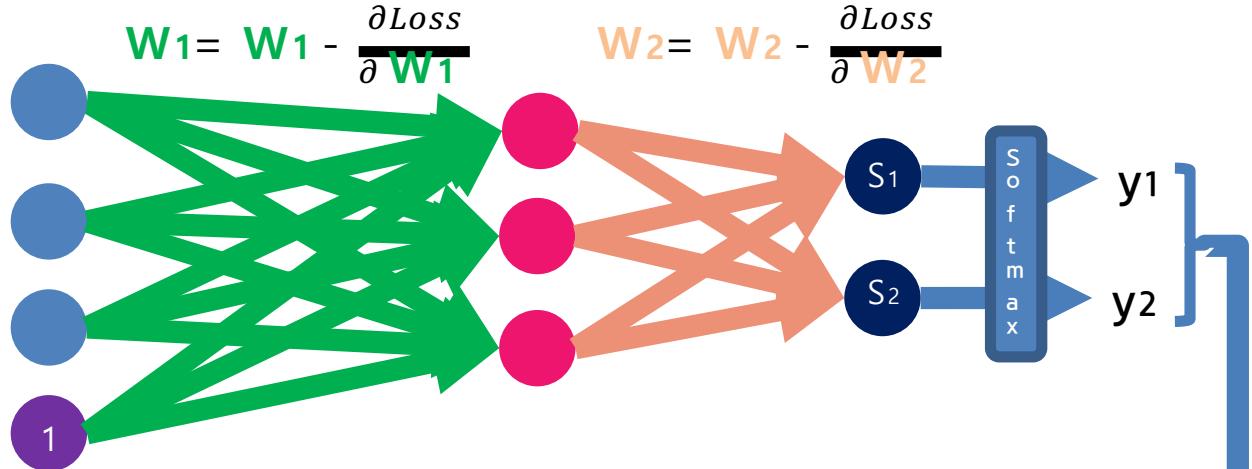
Gradient Descent :

$$w = w - \eta \frac{\partial L}{\partial w}$$



1. Score function
2. Loss function
3. Optimization

돌아보기 ..



Loss를 줄이는 방법

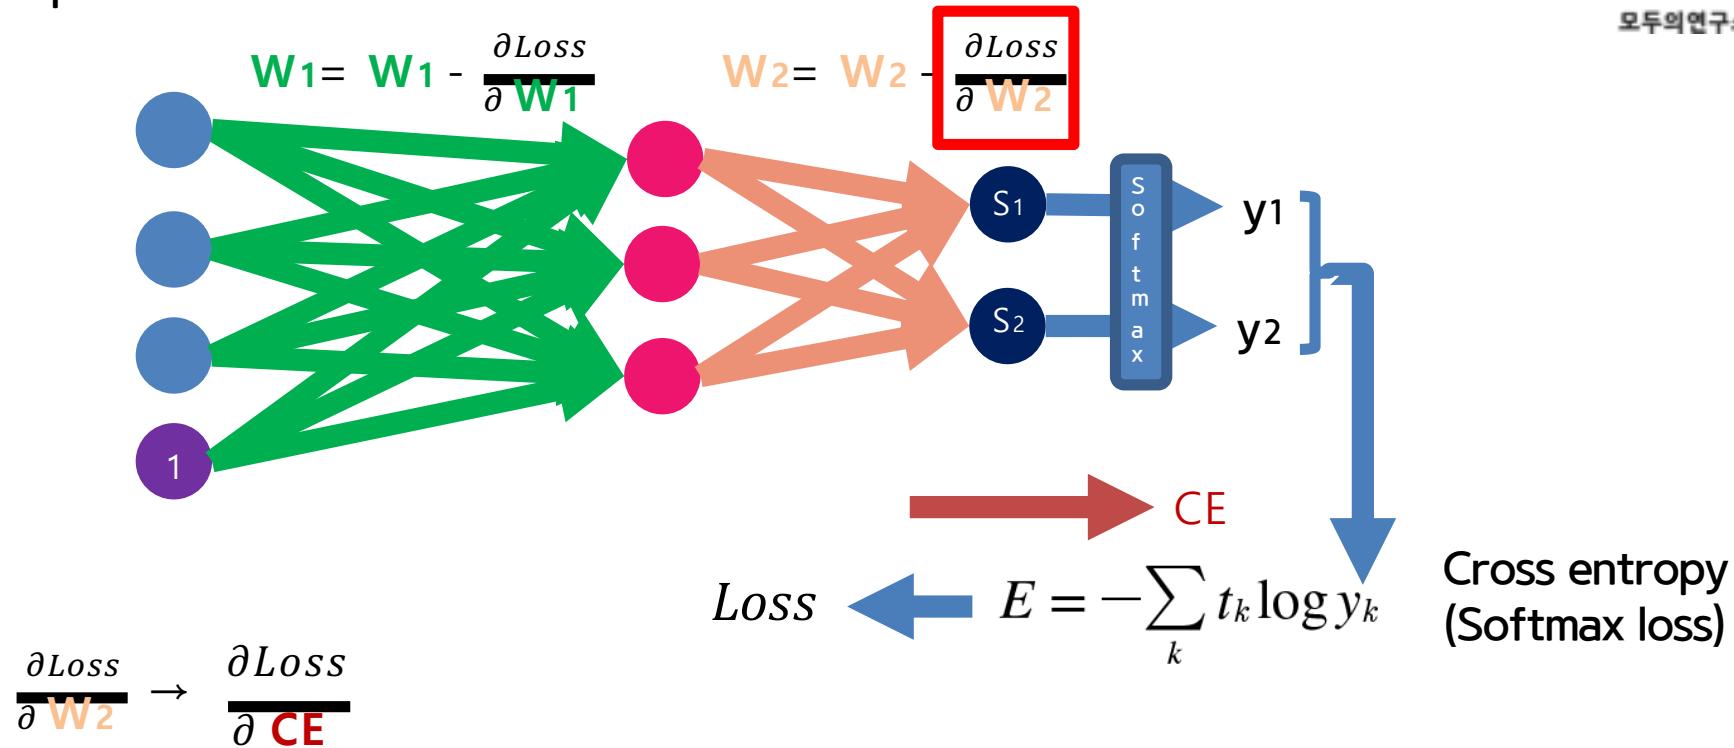
$$Loss \leftarrow E = -\sum_k t_k \log y_k$$

Cross entropy
(Softmax loss)

Gradient 계산방법 : Backpropagation

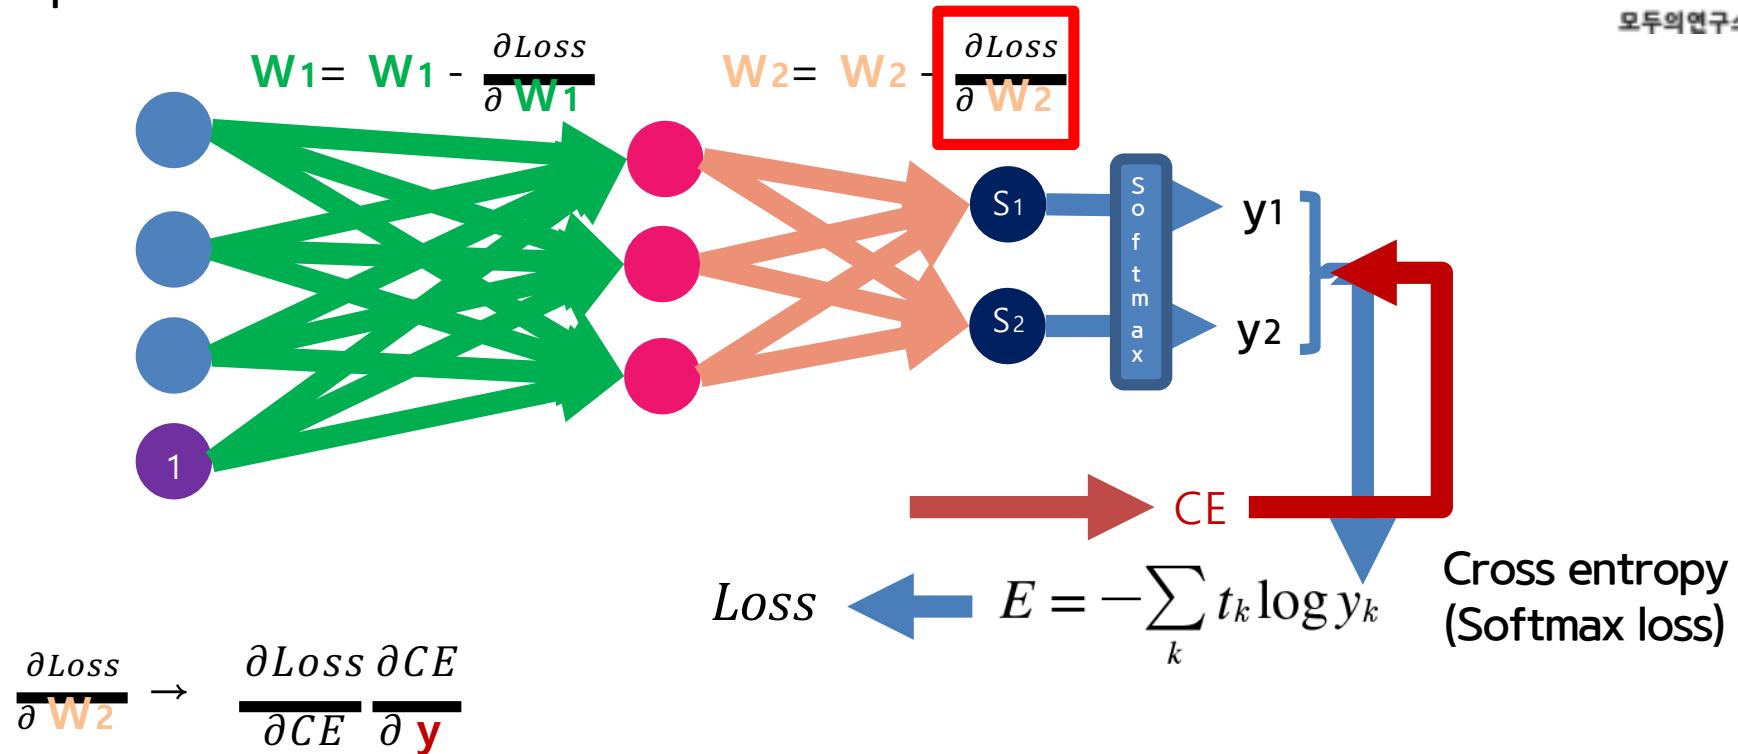
1. Score function
2. Loss function
3. Optimization

돌아보기 ..



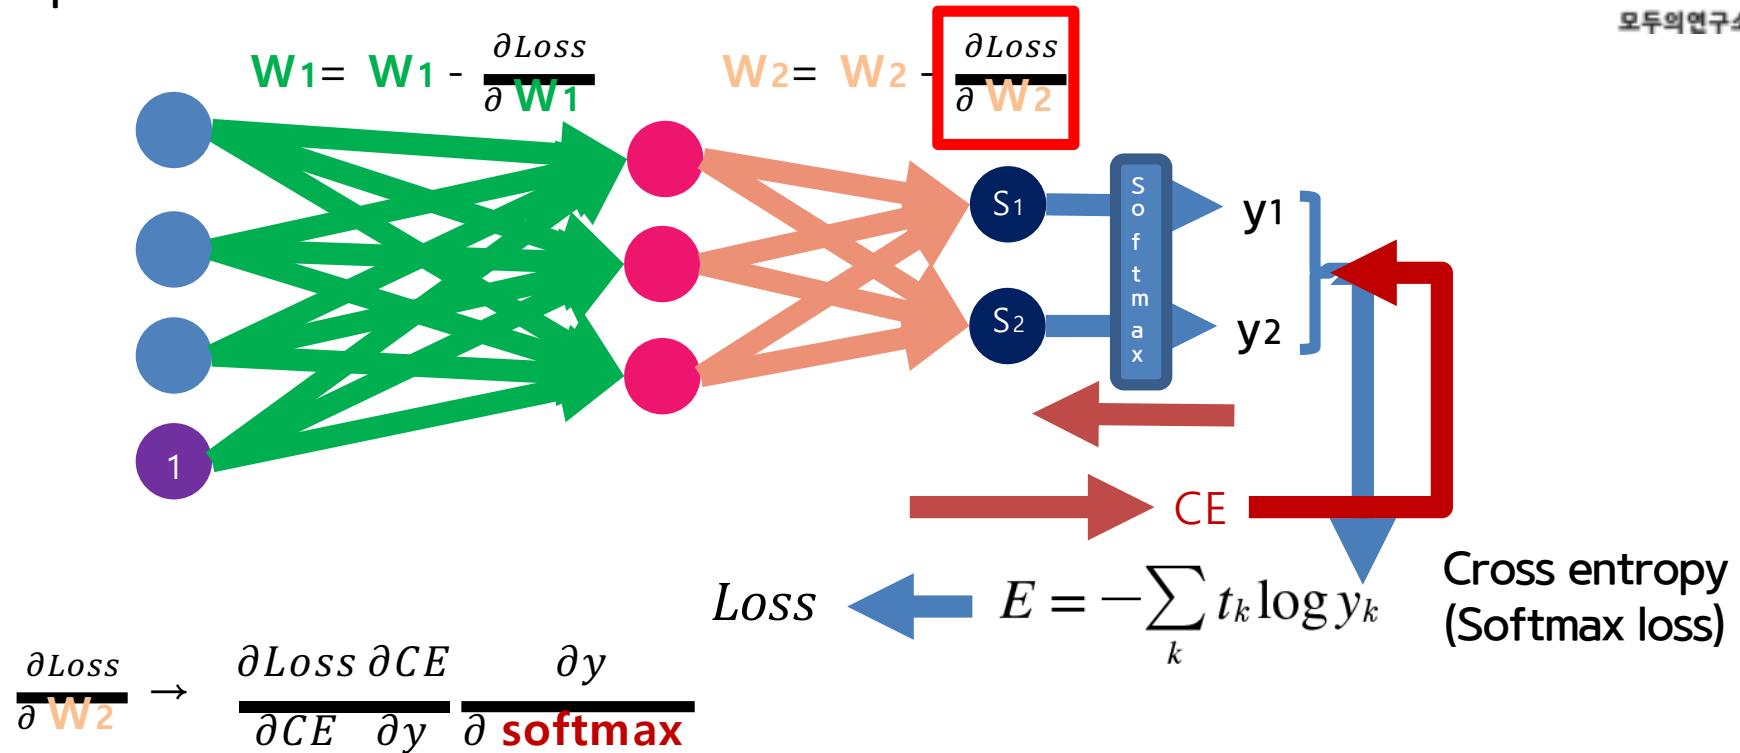
1. Score function
2. Loss function
3. Optimization

돌아보기 ..



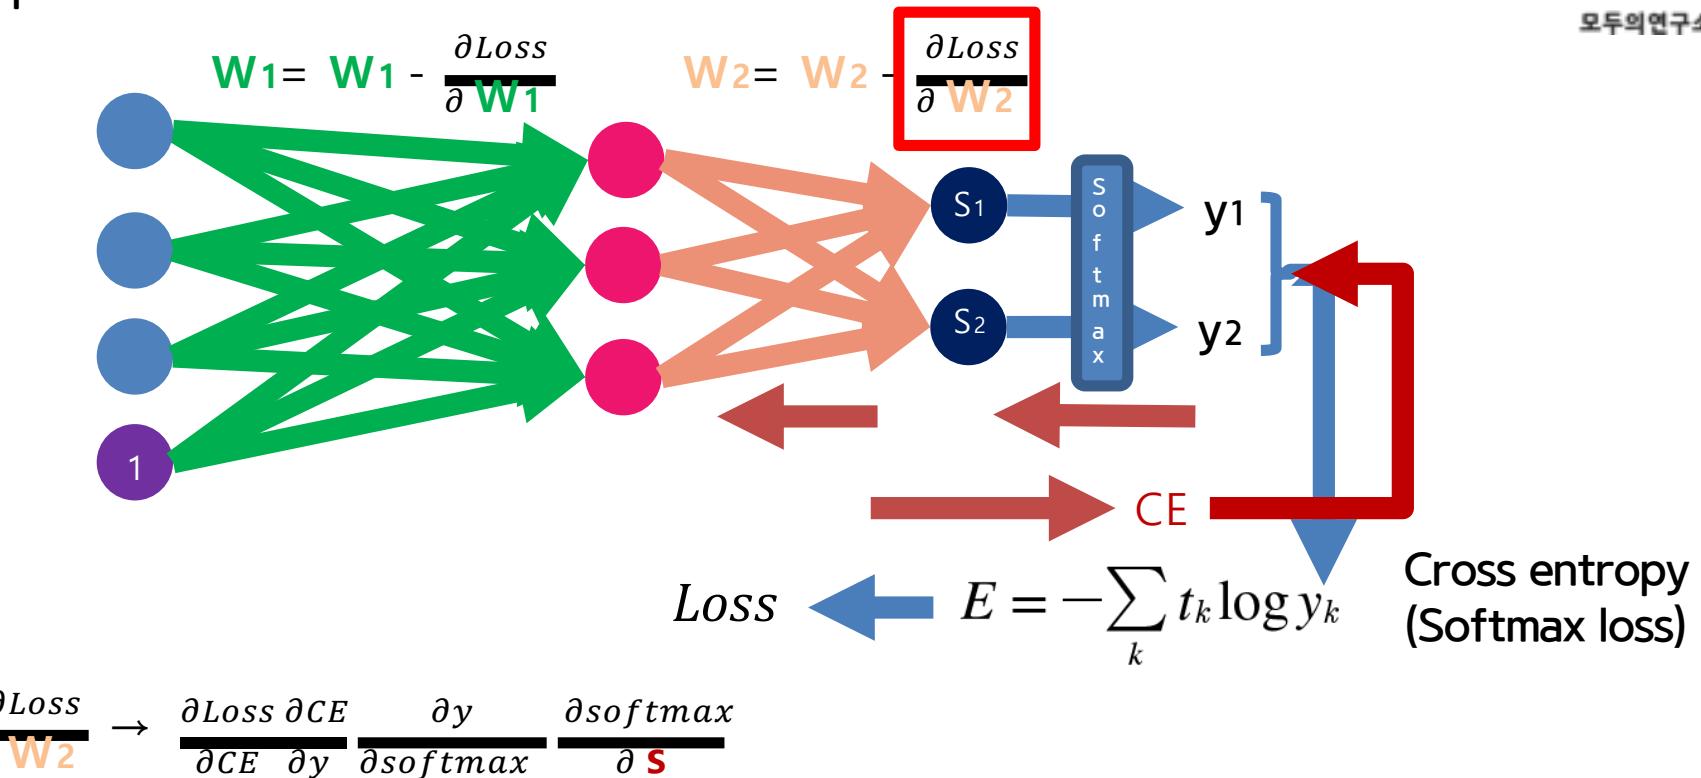
1. Score function
2. Loss function
3. Optimization

돌아보기 ..



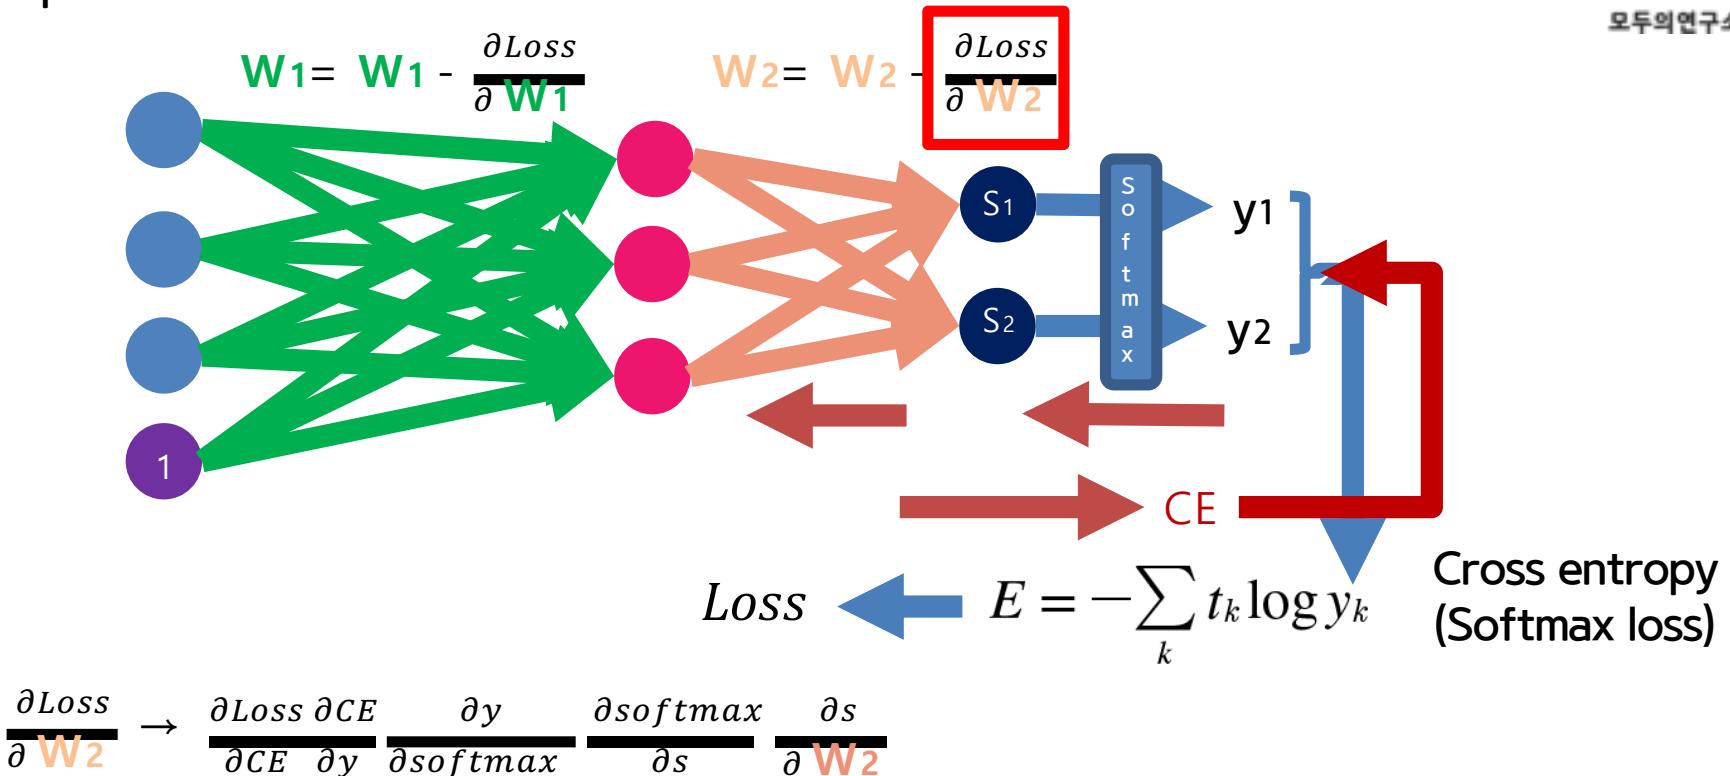
1. Score function
2. Loss function
3. Optimization

돌아보기 ..



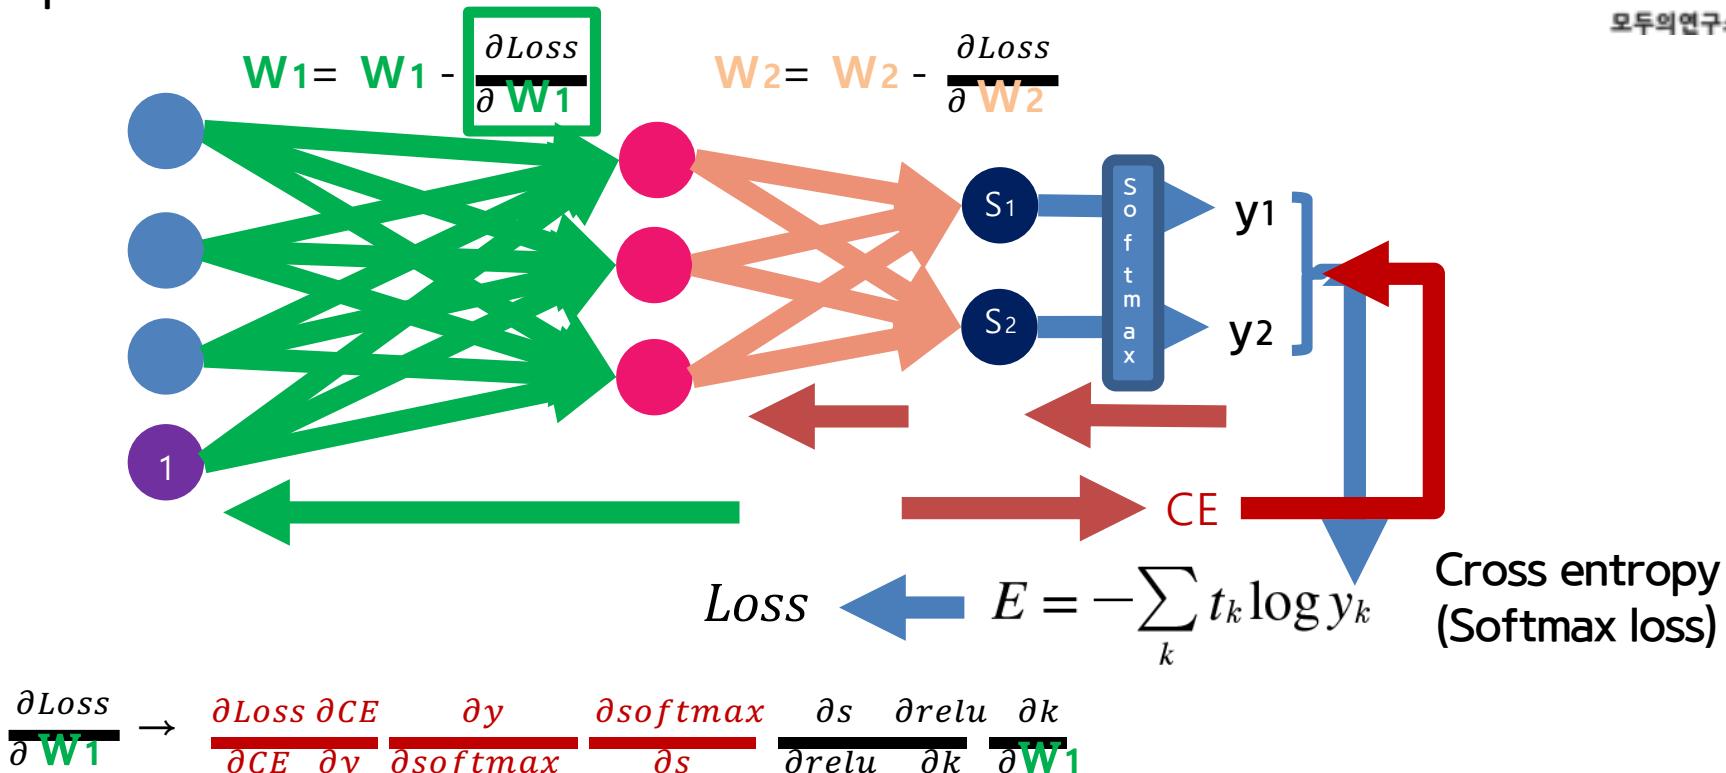
1. Score function
2. Loss function
3. Optimization

돌아보기 ..



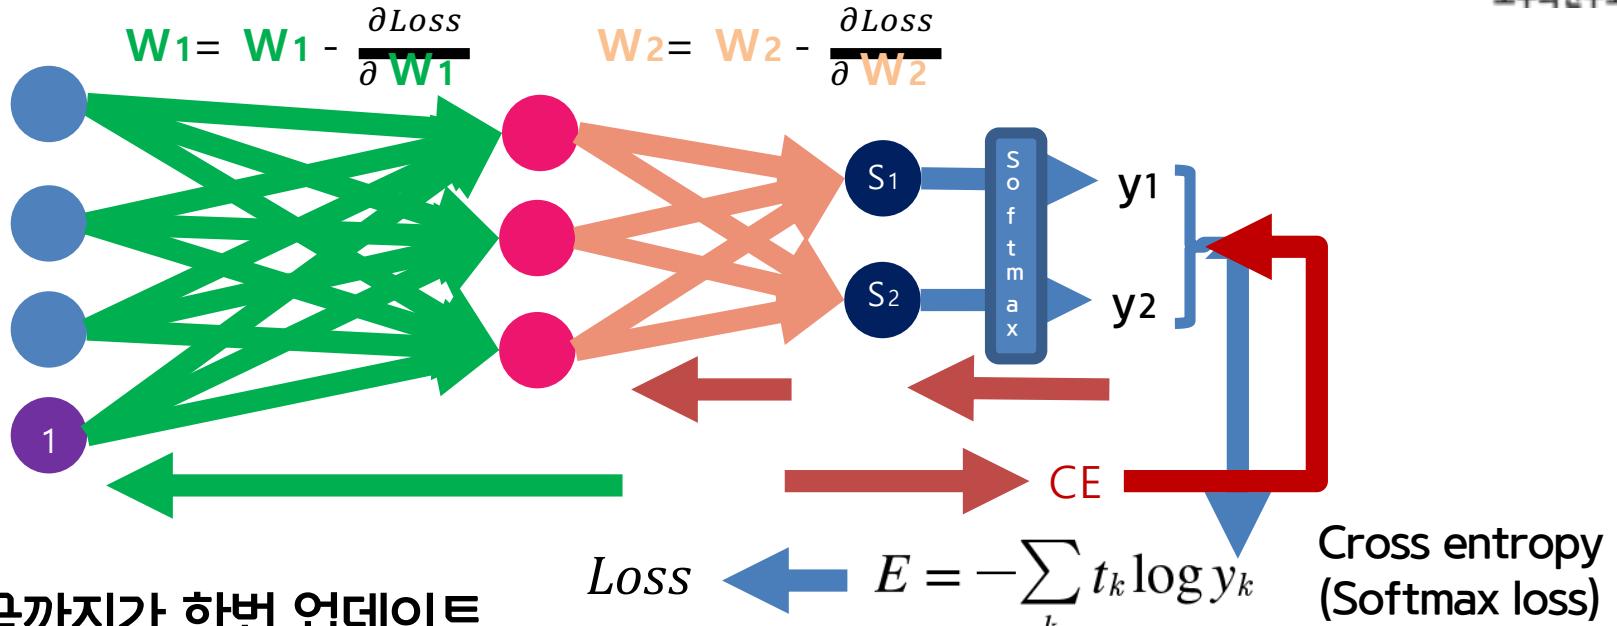
1. Score function
2. Loss function
3. Optimization

돌아보기 ..



1. Score function
2. Loss function
3. Optimization

돌아보기 ..



- 지금까지가 한번 업데이트
- 모든 이미지에 대해서 반복

Multi Layer Perceptron 을 Keras로 구현해 봅시다

? 표 채우기

- 3day/prac3_mlp_mnist.py

- Categorical 변환

- ? 를 채우시오

```
8 from keras.layers import Dense, Dropout
9 from keras.optimizers import RMSprop
10
11 batch_size = 128
12 num_classes = 10
13 epochs = 20
14
15 # the data, shuffled and split between train and test sets
16 (x_train, y_train), (x_test, y_test) = mnist.load_data()
17
18 x_train = x_train.reshape(60000, 784)
19 x_test = x_test.reshape(10000, 784)
20 x_train = x_train.astype('float32')
21 x_test = x_test.astype('float32')
22 x_train /= 255
23 x_test /= 255
24 print(x_train.shape[0], 'train samples')
25 print(x_test.shape[0], 'test samples')
26
27 # convert class vectors to binary class matrices
28 y_train = keras.utils.to_categorical(y_train, num_classes)
29 y_test = keras.utils.to_categorical(y_test, num_classes)
30
31 model = Sequential()
32
33 ##### 주석에 알게 ? 부분을 채우세요.
34 ##### Code #####
35
36 #### 주석에 알게 ? 부분을 채우세요.
37 # 512개의 뉴런, activation : 'relu'의 Dense Layer
38 model.add(Dense(? , activation='?' , input_shape=(784,)))
39 # 512개의 뉴런, activation : 'relu'
40 model.add(Dense(? , activation='?'))
41 # Classify 수 많음의 뉴런, activation : 'softmax'
42 model.add(Dense(? , activation='?'))
43
44 model.summary()
45
46 # loss 는 크로스 엔트로피, Optimizer는 RMSprop(), metric : 평가지 아닌 평가할 때 사용하는 것.
47 model.compile(loss='categorical_crossentropy',
48                 optimizer=RMSprop(),
49                 metrics=['accuracy'])
```

필요한 것 Import 하기

- 모양변환 및 타입 변환
- Scale 0~1

Convolutional Neural Networks

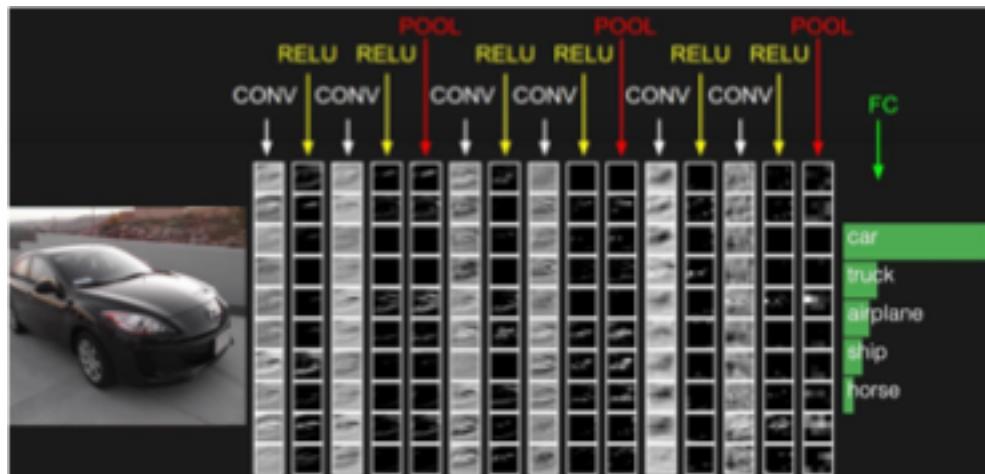
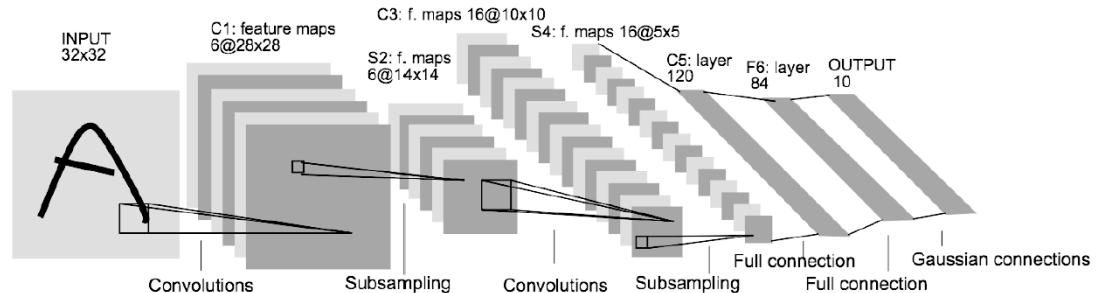
모두의연구소
박은수 책임연구원



큰 그림

큰 그림 : 구조

- LeNet (1998년)



전형적인 구조

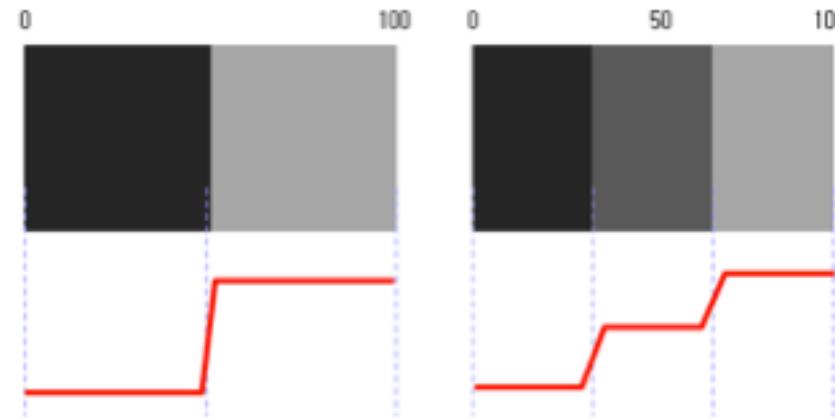
CONV POOL FC

큰 그림 : 컨볼루션은 ...

컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

-1	0	1
-1	0	1
-1	0	1

수직 에지를 찾는
컨볼루션 필터



큰 그림 : 컨볼루션은 ...

컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

수직 에지를 찾는
컨볼루션 필터

0

100

-1	0	1			
-1	0	1			
-1	0	1			



0			

큰 그림 : 컨볼루션은 ...

컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

수직 에지를 찾는
컨볼루션 필터

0

100

	-1	0	1		
	-1	0	1		
	-1	0	1		



0	300		

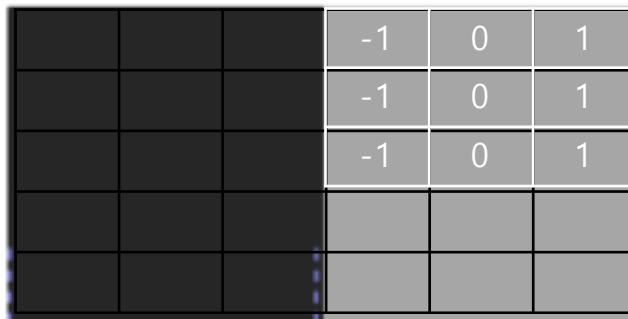
큰 그림 : 컨볼루션은 ...

컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

수직 에지를 찾는
컨볼루션 필터

0

100

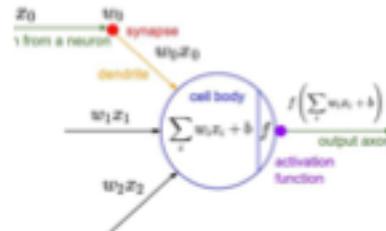


0	300	300	0

큰 그림 : 컨볼루션은 ...

컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

수직 에지를 찾는
컨볼루션 필터



0

100

A 4x6 input image grid. The bottom-left 3x3 subgrid contains values -1, 0, 1 repeated. A 3x3 convolutional filter is applied to this subgrid, resulting in a value of 300 at the bottom-right position of the output feature map. The rest of the output grid is zero.

-1	0	1			
-1	0	1			
-1	0	1			



0	300	300	0
0			

큰 그림 : 컨볼루션은 ...

컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

수직 에지를 찾는
컨볼루션 필터

0

100

	-1	0	1		
	-1	0	1		
	-1	0	1		



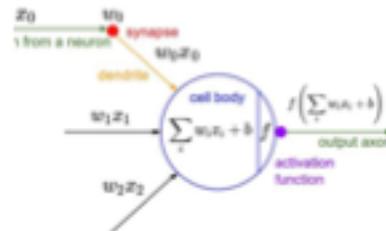
0	300	300	0
0	300		

큰 그림 : 컨볼루션은 ...



컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

수직 에지를 찾는 컨볼루션 필터



0

100

	-1	0	1
	-1	0	1
	-1	0	1

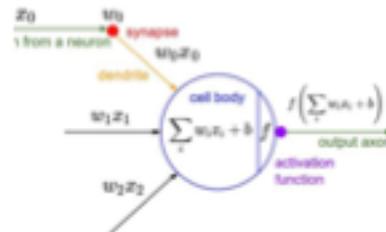


0	300	300	0
0	300	300	

큰 그림 : 컨볼루션은 ...

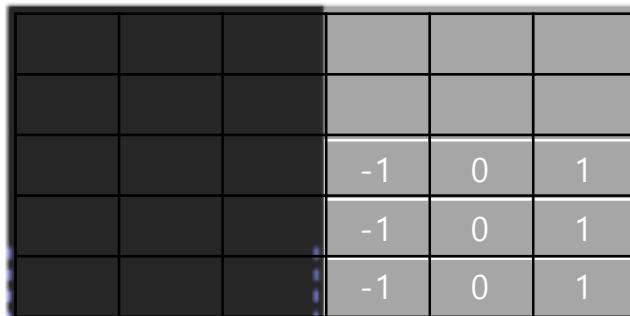
컨볼루션 필터 : 이미지의 특징을 추출할 수 있다

수직 에지를 찾는
컨볼루션 필터



0

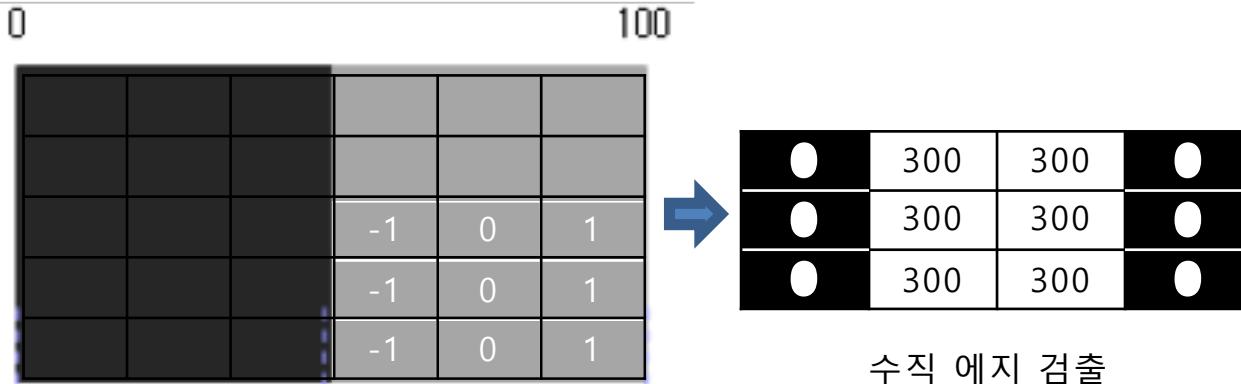
100



0	300	300	0
0	300	300	0
0	300	300	0

큰 그림 : 컨볼루션은 ...

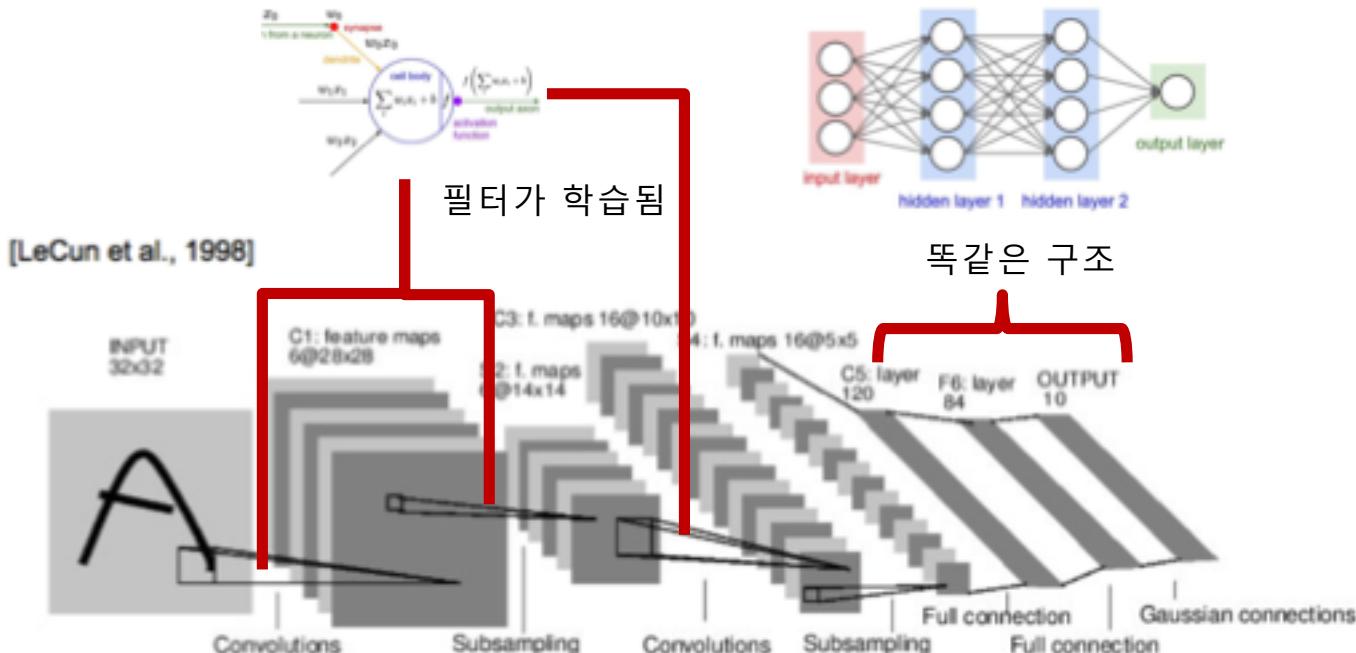
컨볼루션 필터 : 이미지의 특징을 추출할 수 있다



다양한 필터로 다양한
특징을 추출할 수 있다

컨볼루셔널 뉴럴넷 (Convolutional Neural Network)

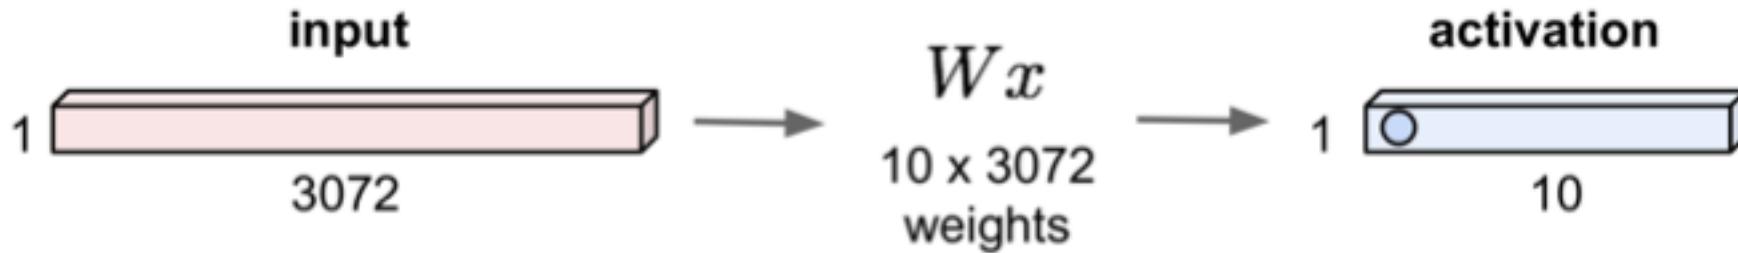
- 필터를 학습하는 구조다



이제부터
자세히 살펴봅시다

Fully Connected Layer

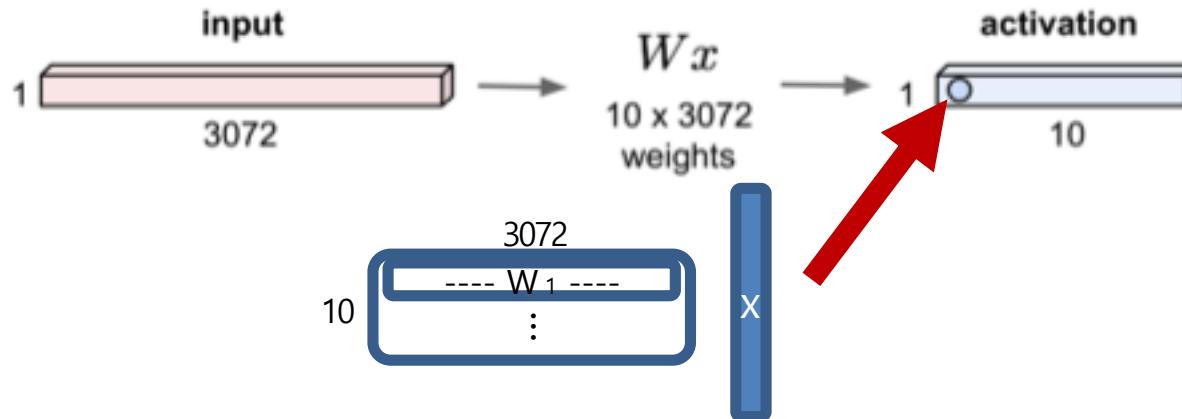
32x32x3 image -> stretch to 3072 x 1



지금까지 살펴봤던 것 입니다

Fully Connected Layer

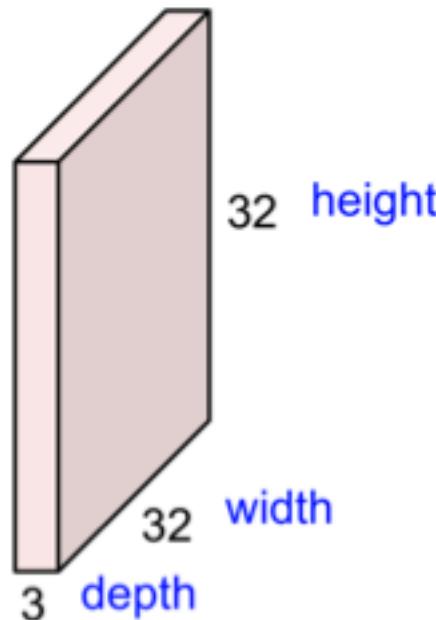
32x32x3 image -> stretch to 3072 x 1



- 1) 전체 영상 각 픽셀에 W_1 의 weighted sum으로 계산한 score
- 2) Activation 적용

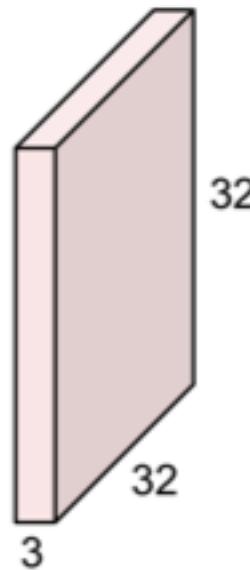
Convolution Layer

32x32x3 image -> preserve spatial structure



Convolution Layer

32x32x3 image



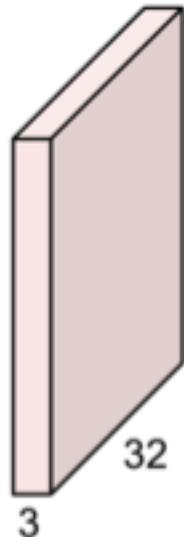
5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer

32x32x3 image



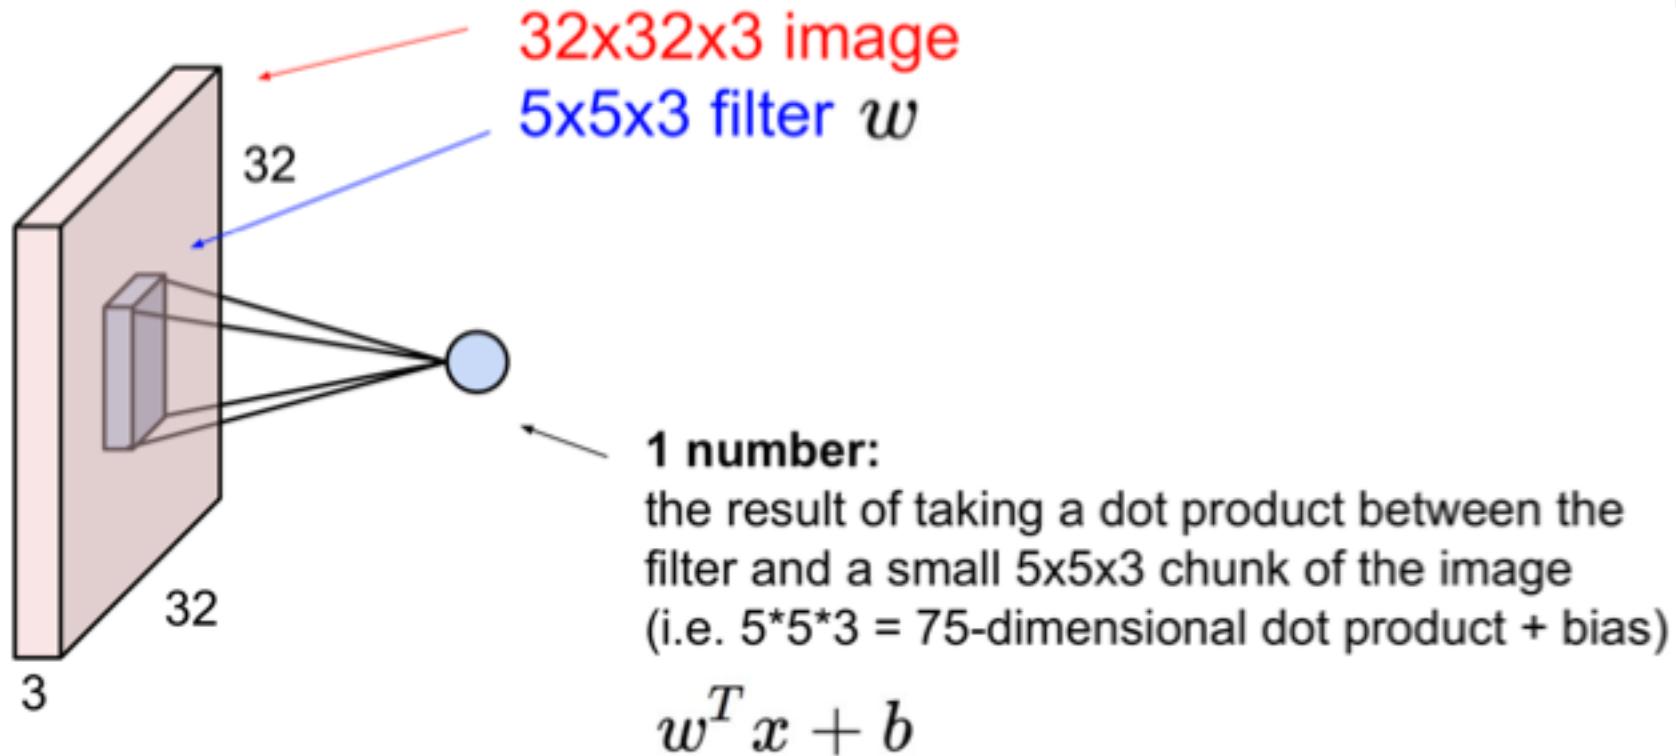
Filters always extend the full depth of the input volume

5x5x3 filter



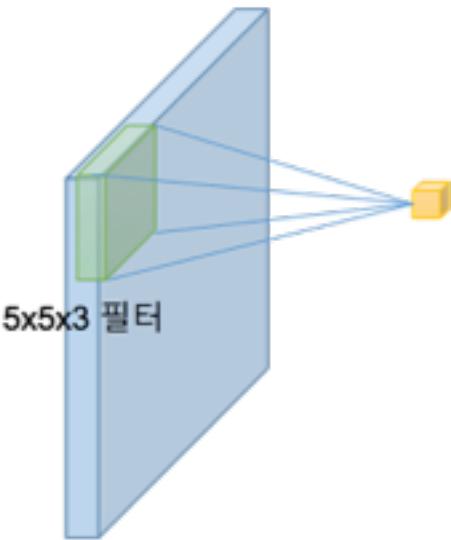
Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer

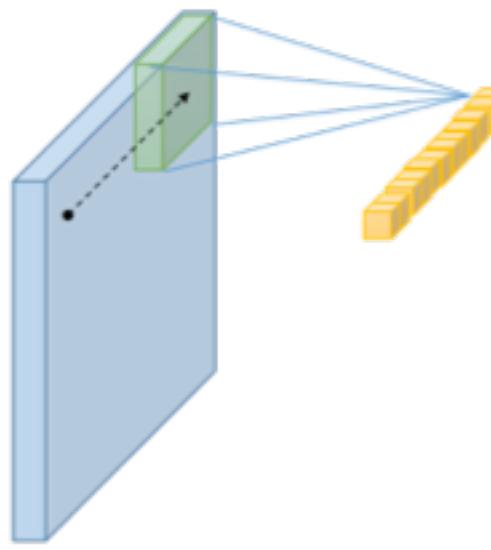


Convolution Layer

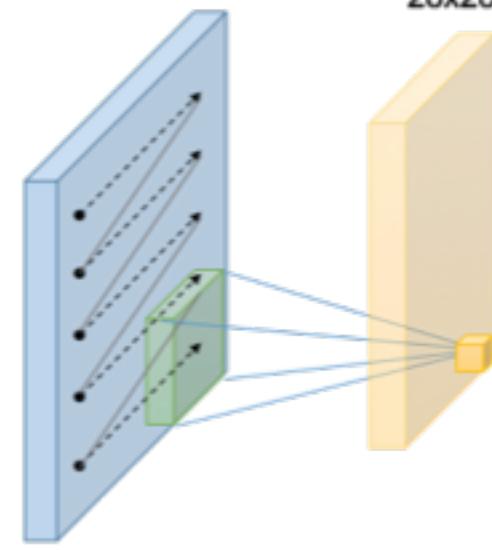
32x32x3 영상



(a)



(b)



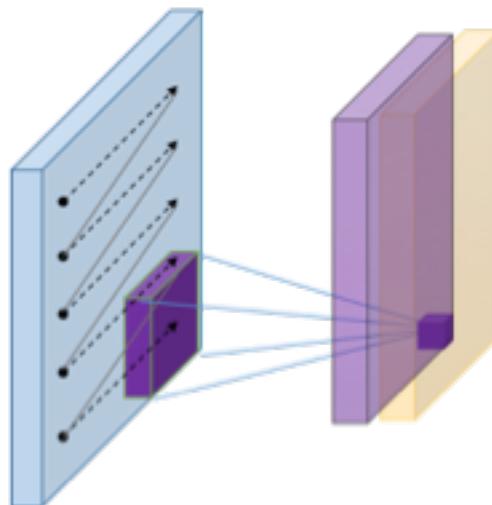
(c)

28x28x1

Convolution Layer

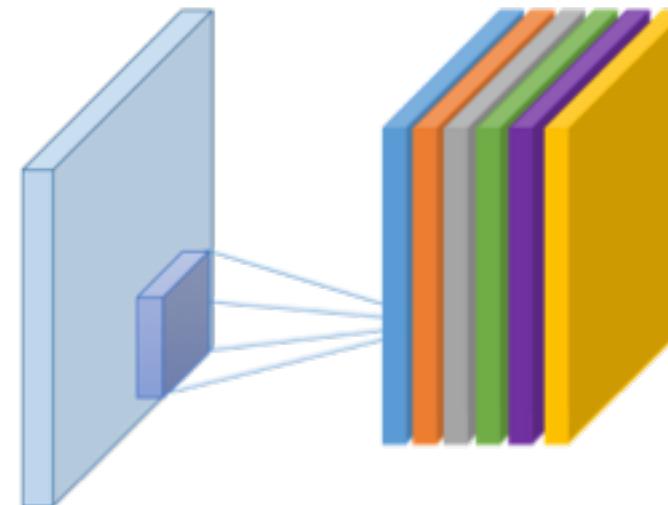
똑같은 크기의 필터 6개를 더 만들어 봅시다

32x32x3 영상



(a)

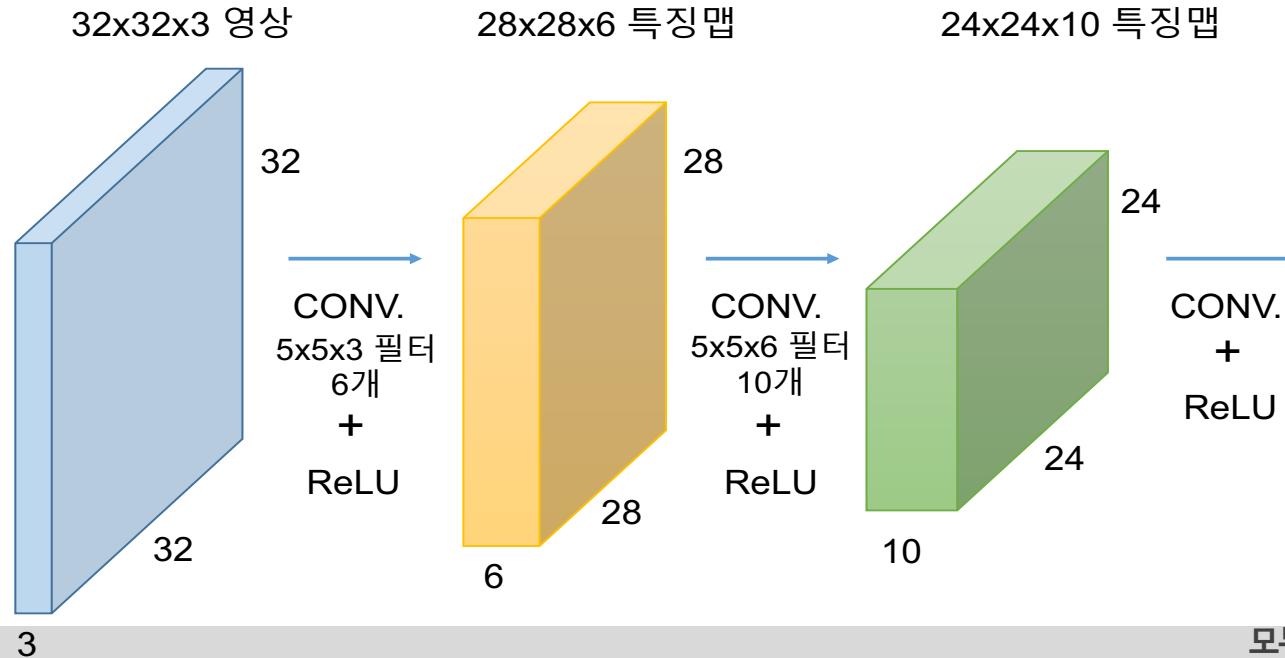
32x32x3 영상



(b)

Convolution Layer

컨볼루션 네트워크는 활성화 함수를 포함한 컨볼루션 레이어의 연결입니다

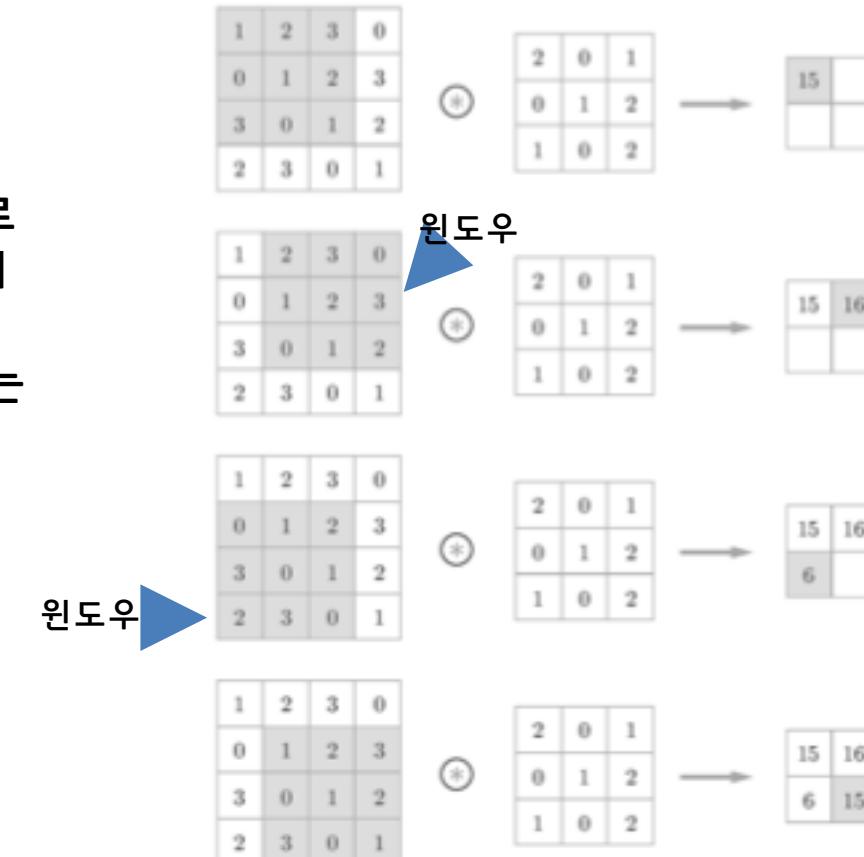


Convolution Layer

- 컨볼루션 연산

- 1) 윈도우를 일정 간격으로 이동해가며 입력 데이터에 적용
- 2) 입력과 필터에 대응하는 원소끼리 곱한 후 그 총합을 함 (단일 곱셈-누산 (fused multiply-add, FMA))

윈도우

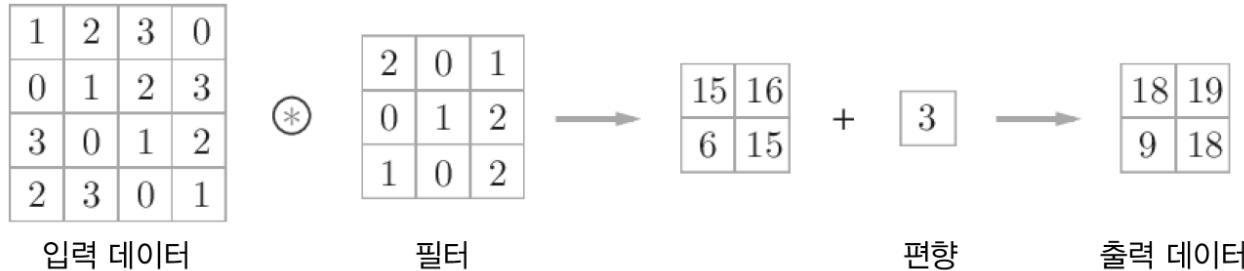


- Bias 파라미터 존재함

Convolution Layer

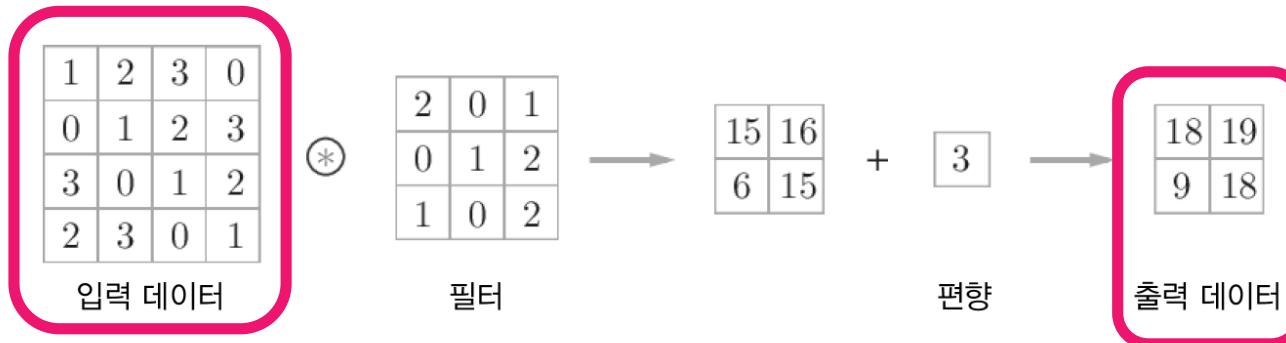


- ## • 컨볼루션 연산



Convolution Layer

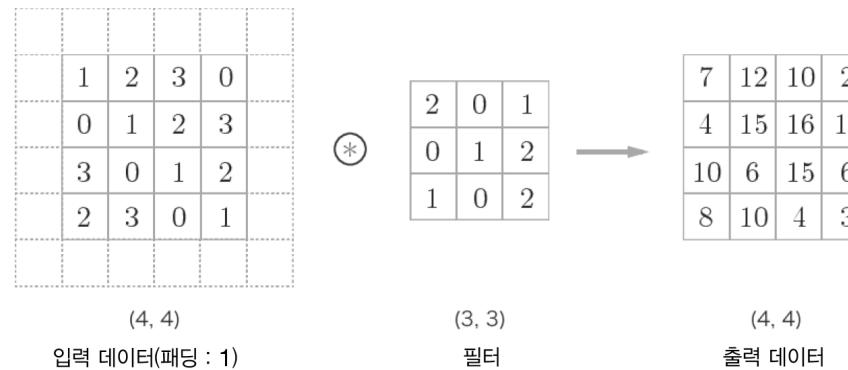
- 컨볼루션 연산



크기가 줄어드는 군요

Convolution Layer

- 패딩 (padding)
 - 컨볼루션 연산의 패딩 처리 : 입력 데이터 주위에 0을 채운다 (패딩은 점선으로 표시했으며 그 안의 값 '0'은 생략함)

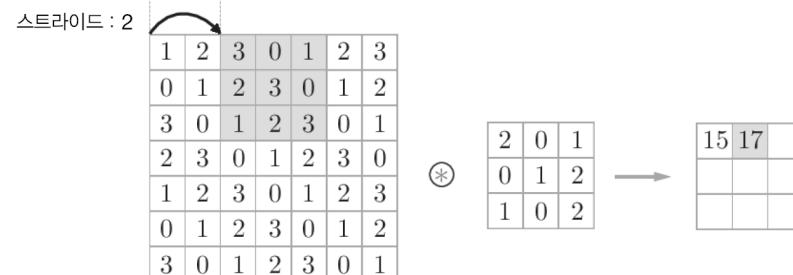
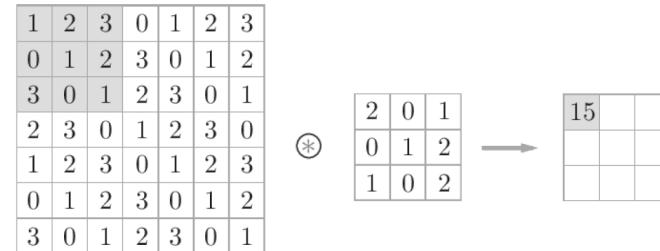


- 패딩은 출력(특징맵(feature map))의 크기를 유지 시키고자 할 때 주로 사용함

Convolution Layer

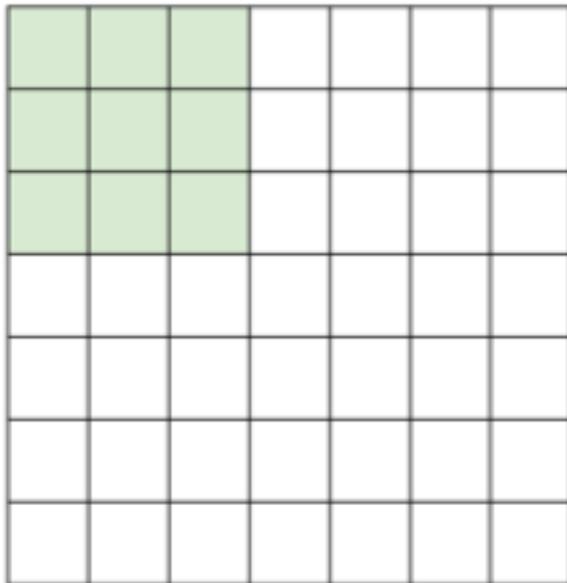
- 스트라이드 (stride)
 - 필터 적용하는 위치의 간격
 - 스트라이드를 2로 하면 필터를 적용하는 윈도우가 두 칸씩 이동함

- 스트라이드를 2로 하니 출력이 3×3 이 됨



특징맵의 크기변화

7

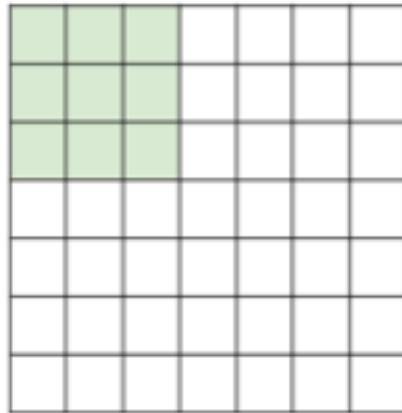


7x7 input (spatially)
assume 3x3 filter

7

특징맵의 크기변화

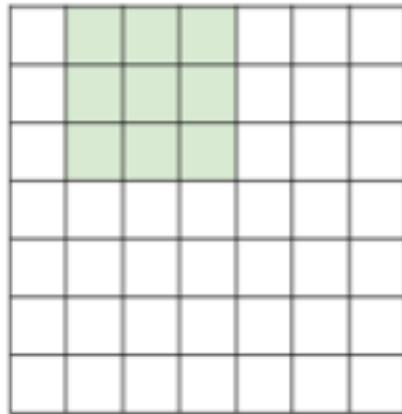
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1

특징맵의 크기변화

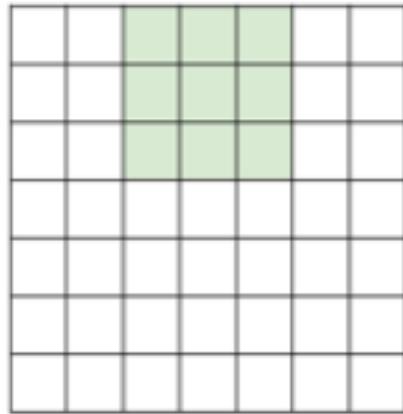
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1

특징맵의 크기변화

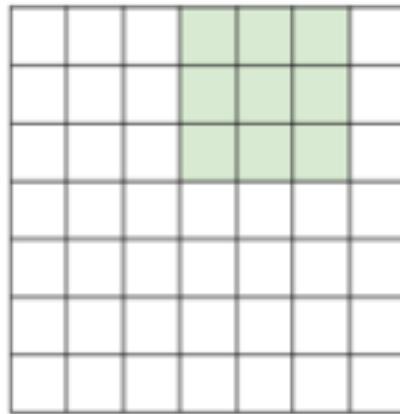
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1

특징맵의 크기변화

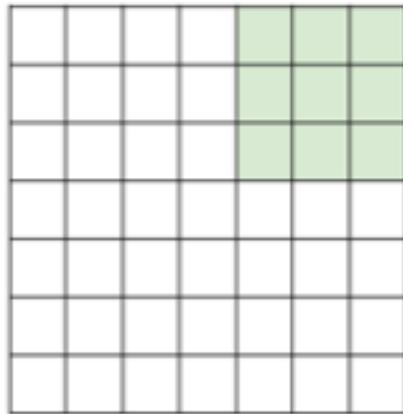
Replicate this column of hidden neurons across space, with some **stride**.



7x7 input
assume 3x3 connectivity, stride 1

특징맵의 크기변화

Replicate this column of hidden neurons across space, with some **stride**.



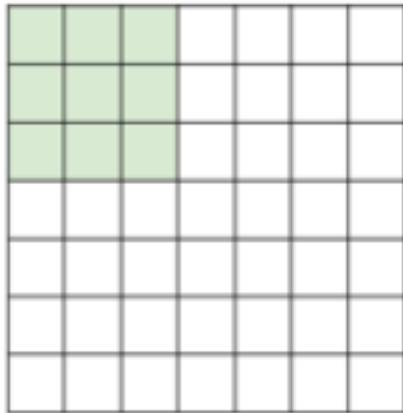
7x7 input

assume 3x3 connectivity, stride 1

=> **5x5 output**

특징맵의 크기변화

Replicate this column of hidden neurons across space, with some **stride**.

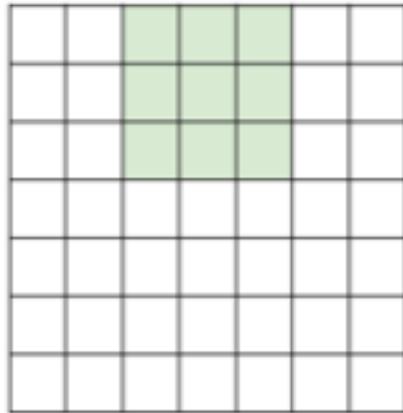


7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?

특징맵의 크기변화

Replicate this column of hidden neurons across space, with some **stride**.

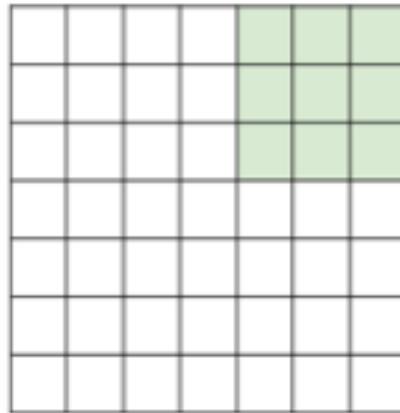


7x7 input
assume 3x3 connectivity, stride 1
=> **5x5 output**

what about stride 2?

특징맵의 크기변화

Replicate this column of hidden neurons across space, with some **stride**.



7x7 input

assume 3x3 connectivity, stride 1

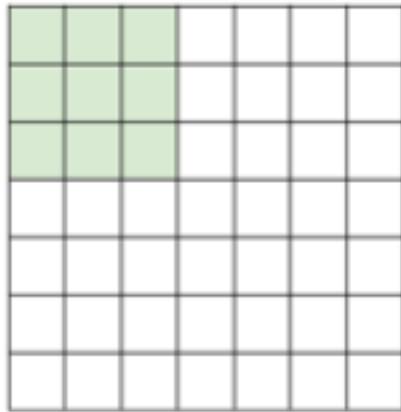
=> **5x5 output**

what about stride 2?

=> **3x3 output**

특징맵의 크기변화

Replicate this column of hidden neurons across space, with some **stride**.



7x7 input

assume 3x3 connectivity, stride 1

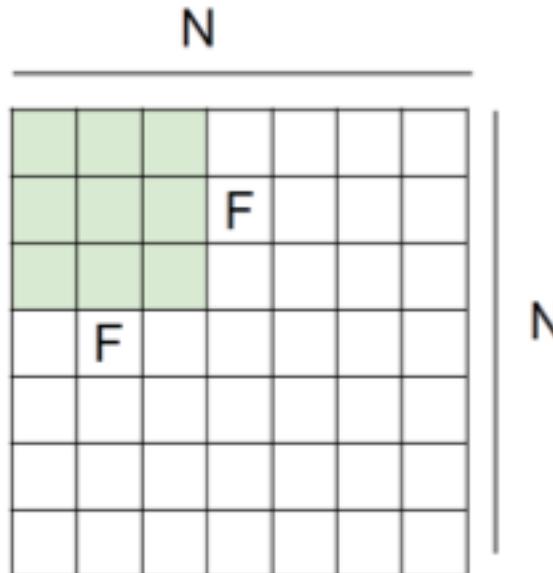
=> **5x5 output**

what about stride 2?

=> **3x3 output**

what about stride 3? **Cannot.**

특징맵의 크기변화



Output size:
 $(N - F) / \text{stride} + 1$

e.g. $N = 7$, $F = 3$:
stride 1 => $(7 - 3)/1 + 1 = 5$
stride 2 => $(7 - 3)/2 + 1 = 3$
stride 3 => $(7 - 3)/3 + 1 = \dots$: \

특징맵의 크기변화

In practice: Common to zero pad the border

0	0	0	0	0	0		
0							
0							
0							
0							

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

(recall:)
$$(N - F) / \text{stride} + 1$$

특징맵의 크기변화



In practice: Common to zero pad the border

e.g. input 7x7

3x3 filter, applied with stride 1

pad with 1 pixel border => what is the output?

7x7 output!

in general, common to see CONV layers with stride 1, filters of size $F \times F$, and zero-padding with $(F-1)/2$. (will preserve size spatially)

e.g. $F = 3 \Rightarrow$ zero pad with 1

$F = 5 \Rightarrow$ zero pad with 2

$F = 7 \Rightarrow$ zero pad with 3

특징맵의 크기변화

- 출력크기 계산해 보기

- 입력크기 : (H, W)
- 필터크기 : (FH, FW)
- 출력크기 : (OH, OW)
- 패딩 : P
- 스트라이드 : S

$$OH = \frac{H + 2P - FH}{S} + 1$$

$$OW = \frac{W + 2P - FW}{S} + 1$$

특징맵의 크기변화

- 출력크기 계산해 보기

입력 : (4,4), 패딩 : 1 , 스트라이드 : 1, 필터: (3,3)

$$OH = \frac{4 + 2 \cdot 1 - 3}{1} + 1 = 4$$

1	2	3	0	
0	1	2	3	
3	0	1	2	
2	3	0	1	

(4, 4)
입력 데이터(패딩 : 1)

$$OH = \frac{H + 2P - FH}{S} + 1$$

$$OW = \frac{W + 2P - FW}{S} + 1$$



2	0	1
0	1	2
1	0	2



7	12	10	2
4	15	16	10
10	6	15	6
8	10	4	3

(3, 3)
필터
(4, 4)
출력 데이터

특징맵의 크기변화

- 출력크기 계산해 보기

입력 : (7,7), 패딩 : 0 , 스트라이드 : 2, 필터: (3,3)

$$OH = \frac{7 + 2 \cdot 0 - 3}{2} + 1 = 3$$

$$OH = \frac{H + 2P - FH}{S} + 1$$

$$OW = \frac{W + 2P - FW}{S} + 1$$

스트라이드 : 2



A 7x7 input feature map with values ranging from 0 to 3. A curved arrow above the first column highlights the stride of 2, indicating that only the first element of each row is considered for the output calculation.

1	2	3	0	1	2	3
0	1	2	3	0	1	2
3	0	1	2	3	0	1
2	3	0	1	2	3	0
1	2	3	0	1	2	3
0	1	2	3	0	1	2
3	0	1	2	3	0	1

⊗

2	0	1
0	1	2
1	0	2



15	17	

특징맵의 크기변화

- 출력크기 계산해 보기

입력 : (28, 31), 패딩 : 2 , 스트라이드 : 3, 필터: (5, 5)

$$OH = ?$$

$$OH = \frac{H + 2P - FH}{S} + 1$$

$$OW = ?$$

$$OW = \frac{W + 2P - FW}{S} + 1$$

특징맵의 크기변화

- 출력크기 계산해 보기

입력 : (28, 31), 패딩 : 2 , 스트라이드 : 3, 필터: (5, 5)

$$OH = \frac{28 + 2 \cdot 2 - 5}{3} + 1 = 10$$

$$OH = \frac{H + 2P - FH}{S} + 1$$

$$OW = \frac{W + 2P - FW}{S} + 1$$

$$OW = \frac{31 + 2 \cdot 2 - 5}{3} + 1 = 11$$

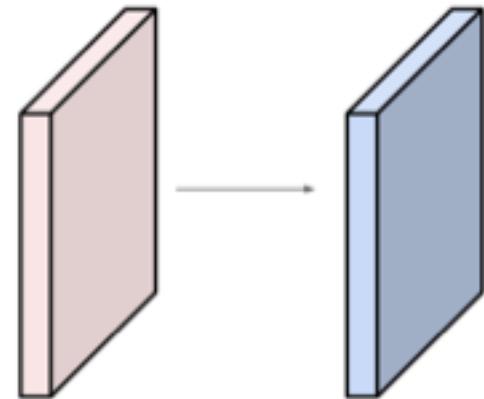
- OH , OW 가 정수가 아니면?
 - 출력이 안나오는 것임. 오류를 내는 등의 대응 필요
 - 딥러닝 프레임워크는 가까운 정수로 내림 하는 등, 특별히 에러를 내지 않고 진행되도록 구현되는 경우가 많음

특징맵의 크기변화

Examples time:

Input volume: **32x32x3**

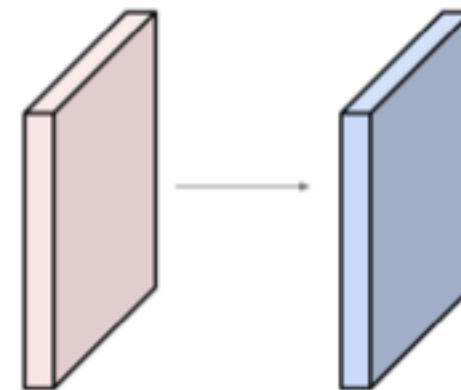
10 5x5 filters with stride 1, pad 2



Number of parameters in this layer?

특징맵의 크기변화

Examples time:



Input volume: **32x32x3**

10 5x5 filters with stride 1, pad 2

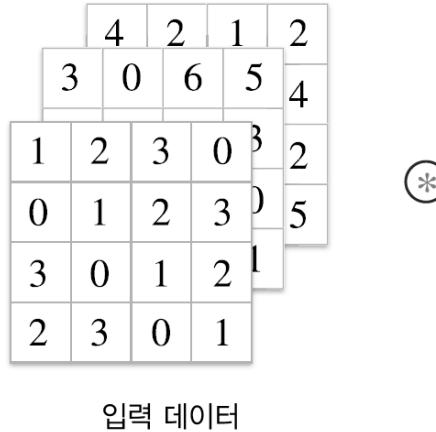
Number of parameters in this layer?

each filter has $5*5*3 + 1 = 76$ params (+1 for bias)

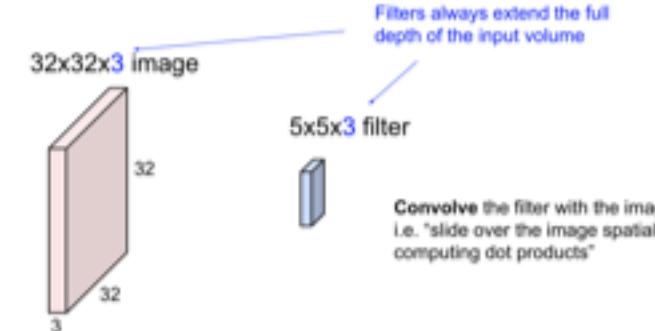
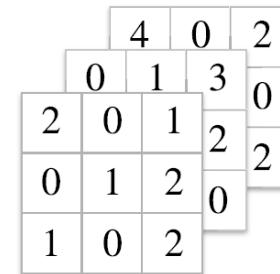
$$\Rightarrow 76 * 10 = 760$$

Convolution Layer

- 필터의 채널 수는 입력되는 특징맵의 채널 수와 일치해야 함을 잊지 마세요

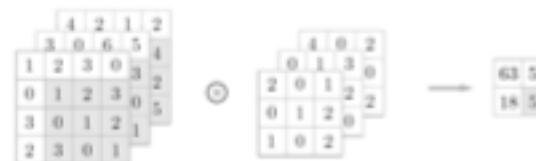
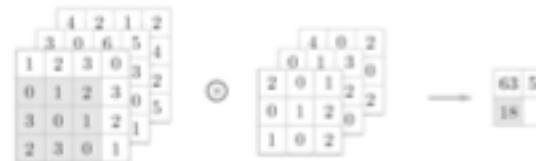
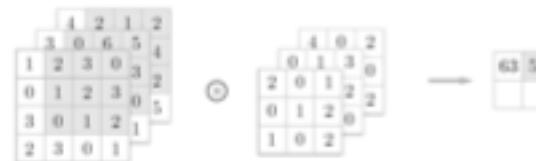
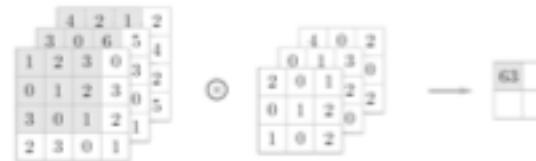


⊗



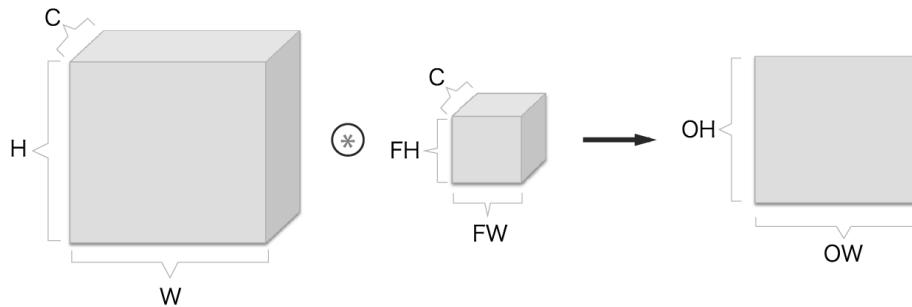
Convolution Layer

- 3차원 데이터의 컨볼루션
 - 주의할 점은 입력데이터의 채널 수와 필터의 채널 수가 같아야 한다는 점
 - 각 필터의 채널크기는 같아야 함



Convolution Layer

- 블록으로 생각하기



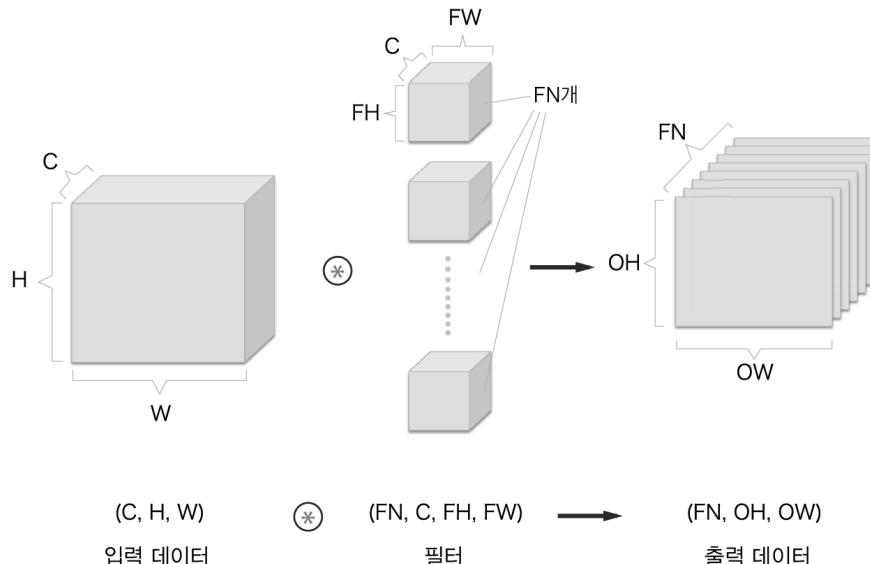
$$\begin{array}{ccc} (C, H, W) & \circledast & (1, OH, OW) \\ \text{입력 데이터} & & \text{출력 데이터} \\ & \text{필터} & \end{array}$$

- 출력 데이터는 한장의 특징맵
- 한장이 아니라 여러장, 즉 다수의 채널을 내보내려면?

Convolution Layer



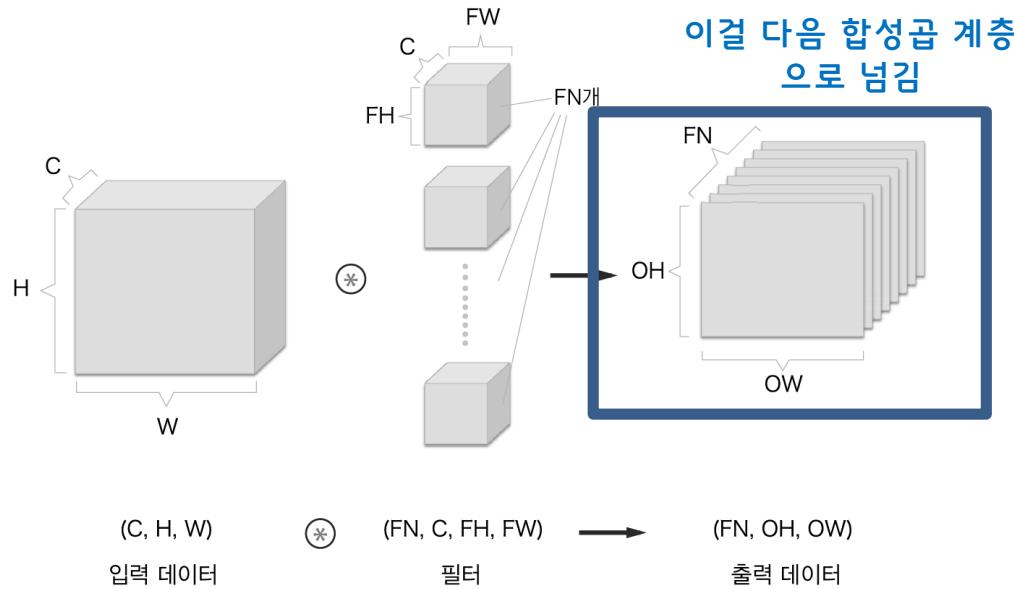
- 블록으로 생각하기
 - 한장이 아니라 여러장, 즉 다수의 채널을 내보내려면?
 - 필터를 여러개 사용하면 됨



Convolution Layer



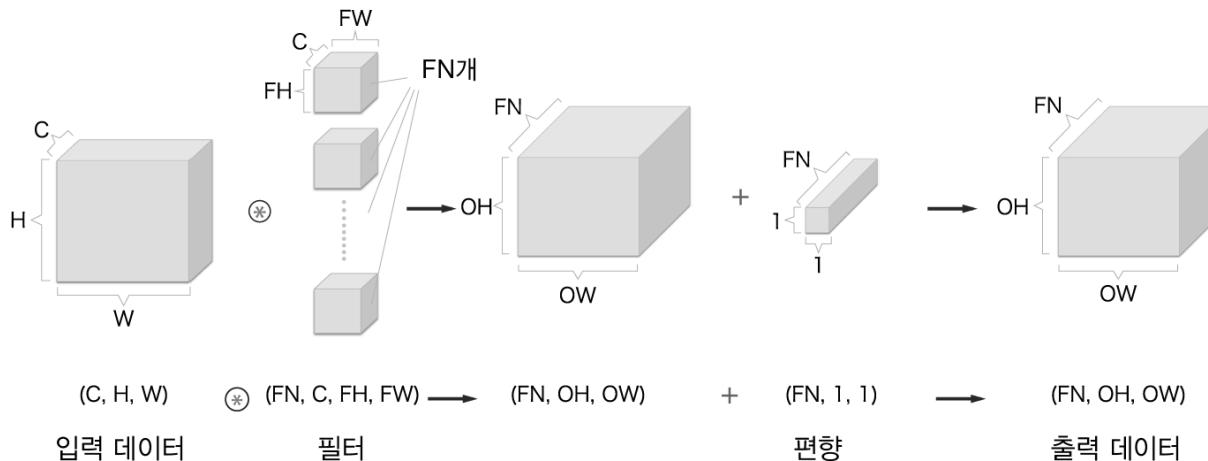
- **블록으로 생각하기**
 - 한장이 아니라 여러장, 즉 다수의 채널을 내보내려면?
 - **필터를 여러개 사용하면 됨**



Convolution Layer

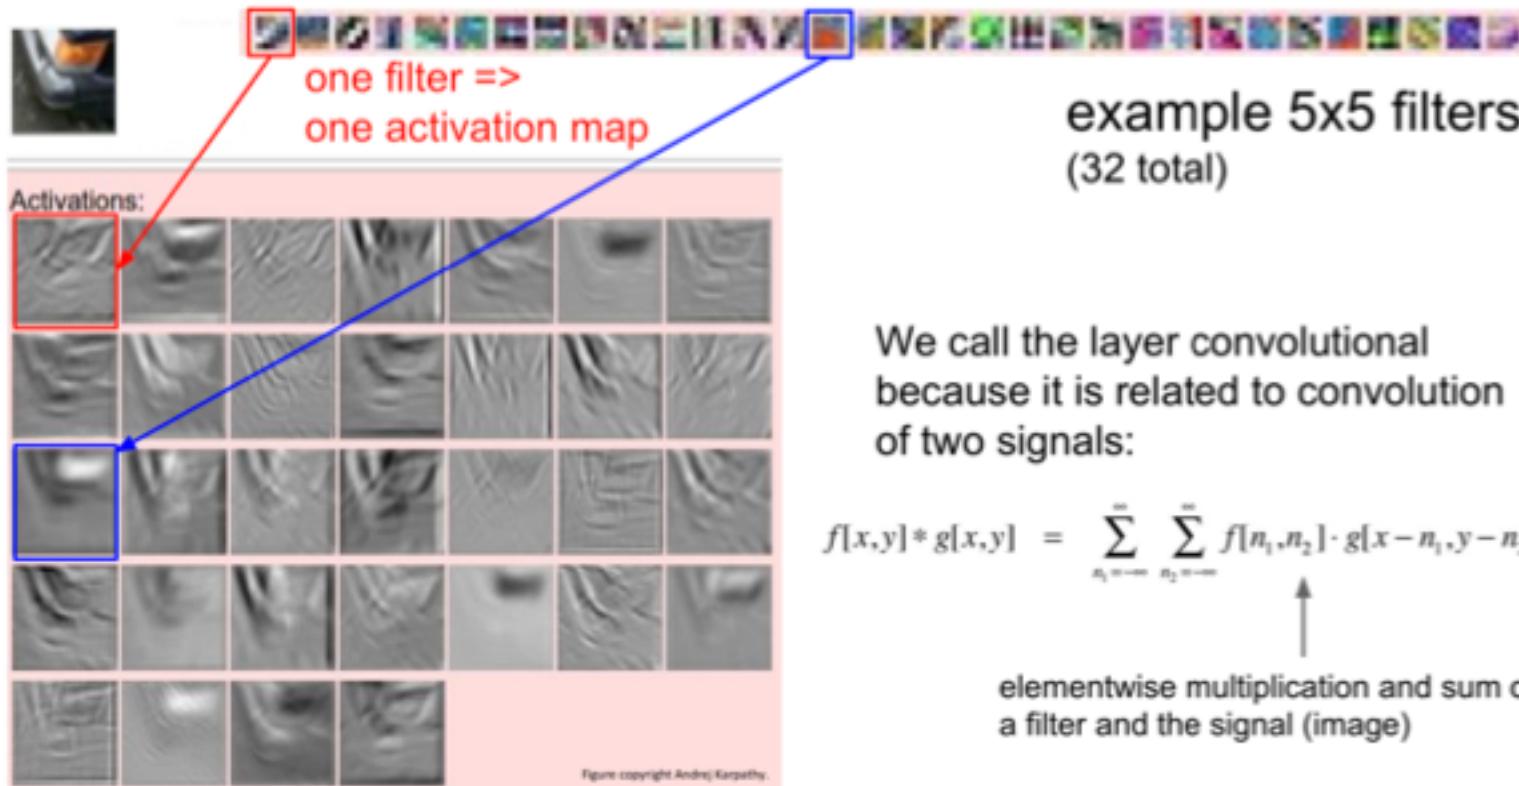


- 편향 (bias)



- 필터마다 편향이 하나씩 존재

Convolution Layer



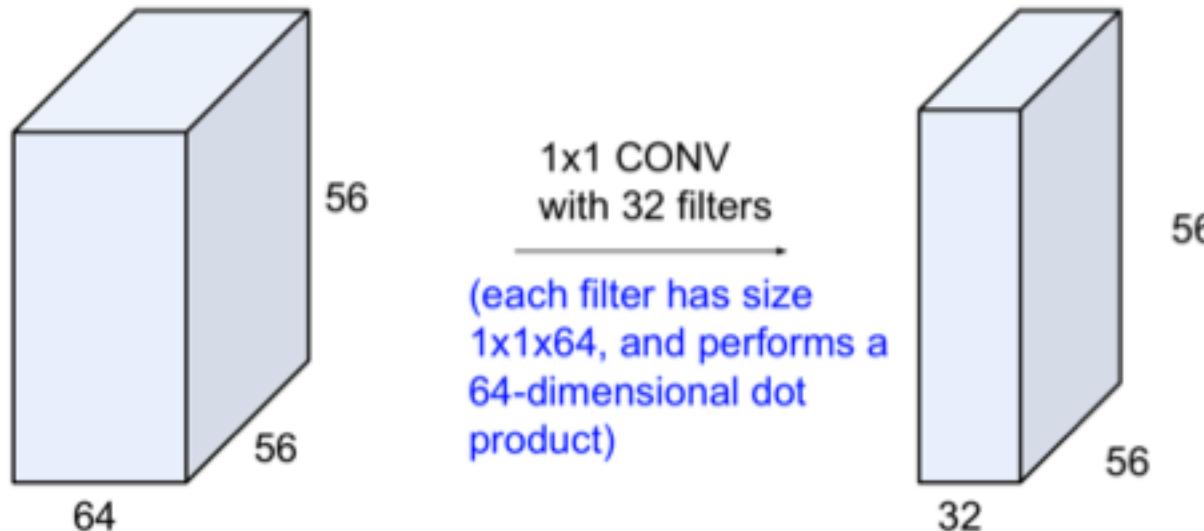
We call the layer convolutional because it is related to convolution of two signals:

$$f[x, y] * g[x, y] = \sum_{n_1=-\infty}^{\infty} \sum_{n_2=-\infty}^{\infty} f[n_1, n_2] \cdot g[x - n_1, y - n_2]$$

↑
elementwise multiplication and sum of a filter and the signal (image)

Convolution Layer

(btw, 1x1 convolution layers make perfect sense)



Convolution Layer

Example: CONV layer in Torch

SpatialConvolution

```
module = nn.SpatialConvolution(nInputPlane, nOutputPlane, kW, kH, [dW], [dH], [padW], [padH])
```

Applies a 2D convolution over an input image composed of several input planes. The `input` tensor in `forward(input)` is expected to be a 3D tensor (`nInputPlane x height x width`).

The parameters are the following:

- `nInputPlane` : The number of expected input planes in the image given into `forward()`.
- `nOutputPlane` : The number of output planes the convolution layer will produce.
- `kW` : The kernel width of the convolution
- `kH` : The kernel height of the convolution
- `dW` : The step of the convolution in the width dimension. Default is `1`.
- `dH` : The step of the convolution in the height dimension. Default is `1`.
- `padW` : The additional zeros added per width to the input planes. Default is `0`, a good number is `(kW-1)/2`.
- `padH` : The additional zeros added per height to the input planes. Default is `padW`, a good number is `(kH-1)/2`.

Summary To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .

Note that depending of the size of your kernel, several (of the last) columns or rows of the input image might be lost. It is up to the user to add proper padding in images.

If the input image is a 3D tensor `nInputPlane x height x width`, the output image size will be `nOutputPlane x oheight x owidth` where

```
owidth = Floor((width + 2*padW - kW) / dW + 1)
oheight = Floor((height + 2*padH - kH) / dH + 1)
```

Torch is licensed under [BSD 3-clause](#).

Convolution Layer

Example: CONV
layer in Caffe

```

layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  # learning rate and decay multipliers for the filters
  param { lr_mult: 1 decay_mult: 1 }
  # learning rate and decay multipliers for the biases
  param { lr_mult: 2 decay_mult: 0 }
  convolution_param {
    num_output: 96      # learn 96 filters
    kernel_size: 11     # each filter is 11x11
    stride: 4           # step 4 pixels between each filter application
    weight_filler {
      type: "gaussian" # initialize the filters from a Gaussian
      std: 0.01         # distribution with stdev 0.01 (default mean: 0)
    }
    bias_filler {
      type: "constant" # initialize the biases to zero (0)
      value: 0
    }
  }
}

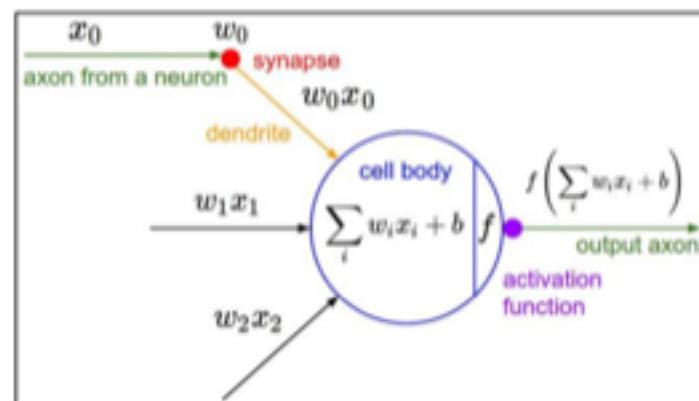
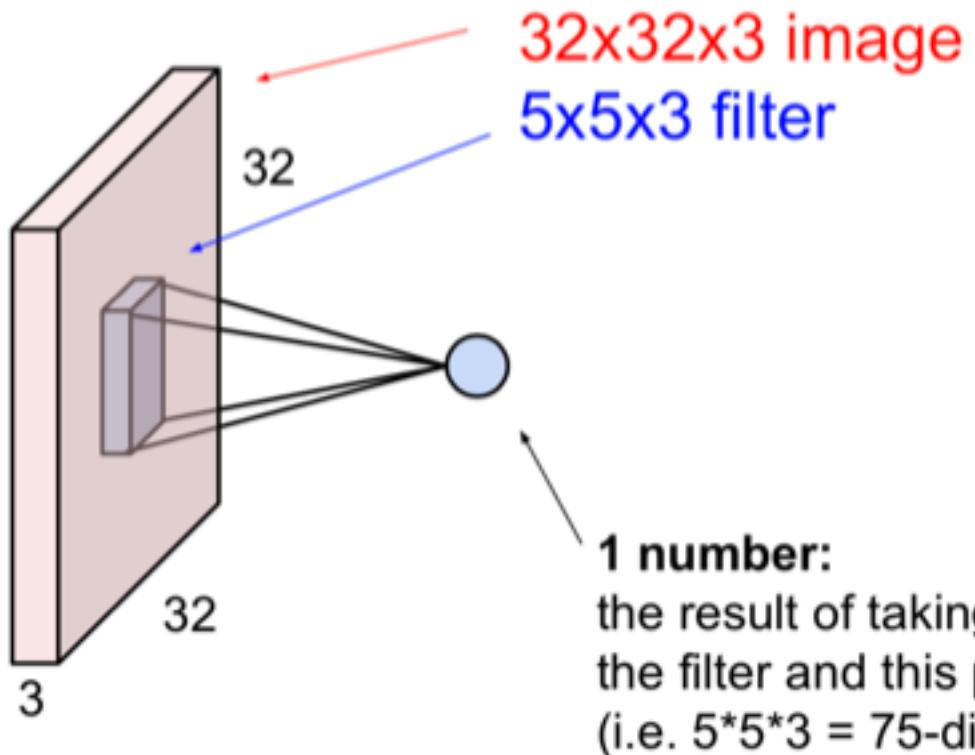
```

Summary. To summarize, the Conv Layer:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires four hyperparameters:
 - Number of filters K ,
 - their spatial extent F ,
 - the stride S ,
 - the amount of zero padding P .

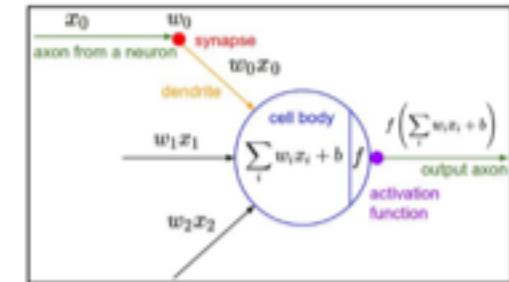
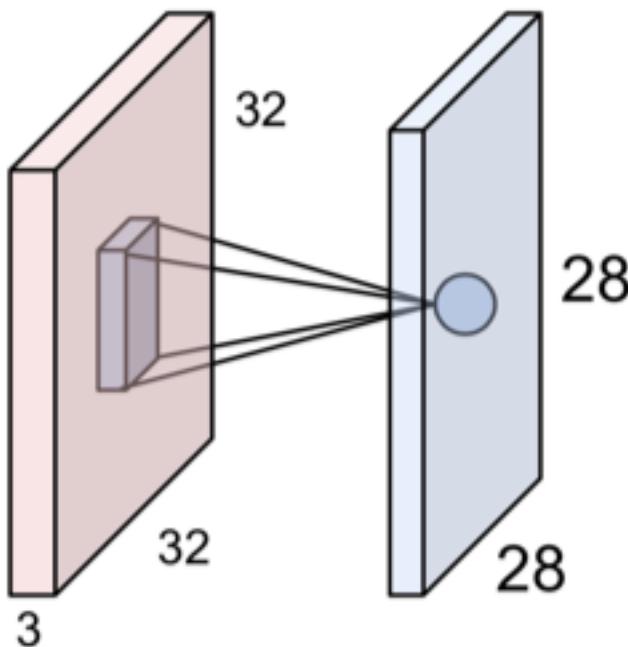
Caffe is licensed under BSD 2-Clause.

The brain/neuron view of CONV Layer



It's just a neuron with local connectivity...

The brain/neuron view of CONV Layer

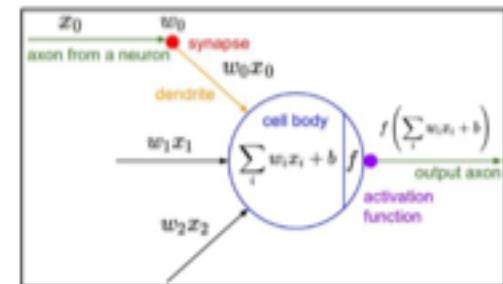
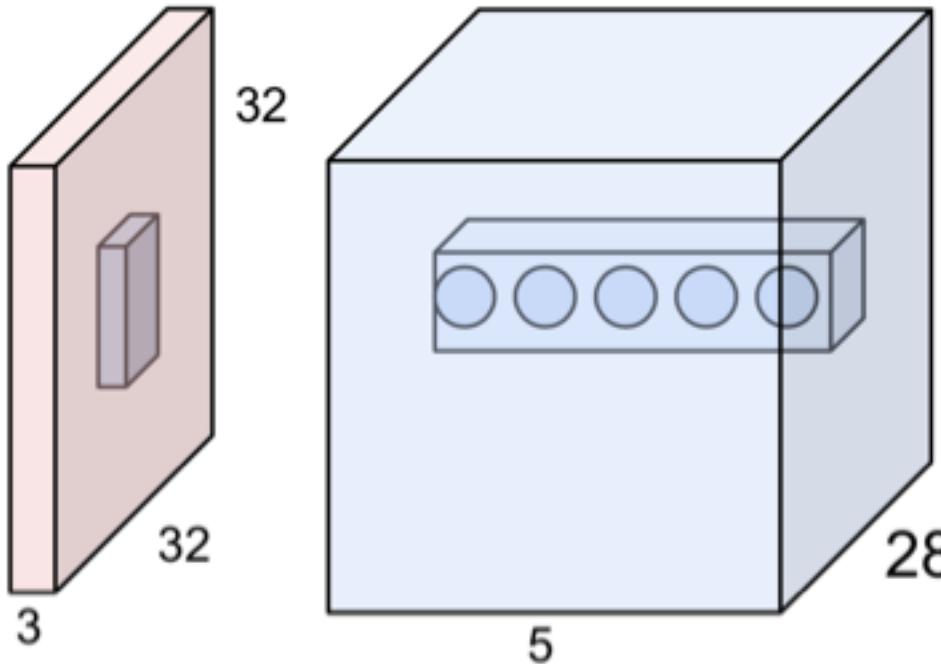


An activation map is a 28x28 sheet of neuron outputs:

1. Each is connected to a small region in the input
2. All of them share parameters

“5x5 filter” -> “5x5 receptive field for each neuron”

The brain/neuron view of CONV Layer



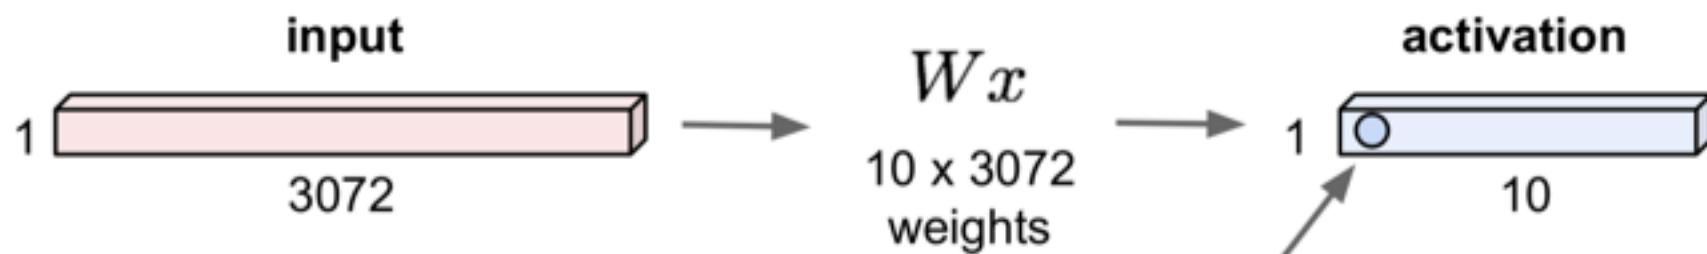
E.g. with 5 filters,
CONV layer consists of
neurons arranged in a 3D grid
(28x28x5)

There will be 5 different
neurons all looking at the same
region in the input volume

Reminder: Fully Connected Layer

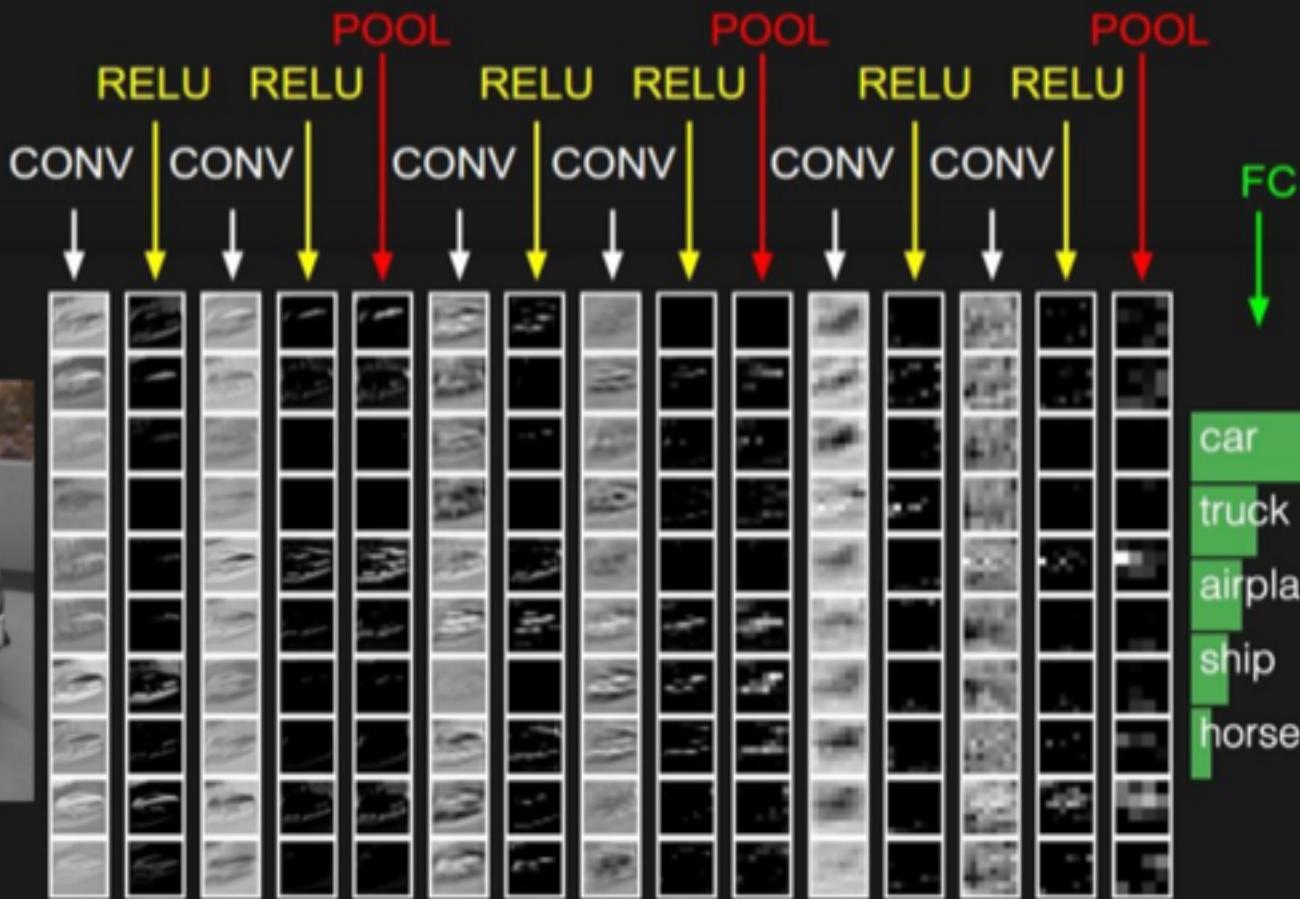
32x32x3 image -> stretch to 3072 x 1

Each neuron
looks at the full
input volume



1 number:
the result of taking a dot product
between a row of W and the input
(a 3072-dimensional dot product)

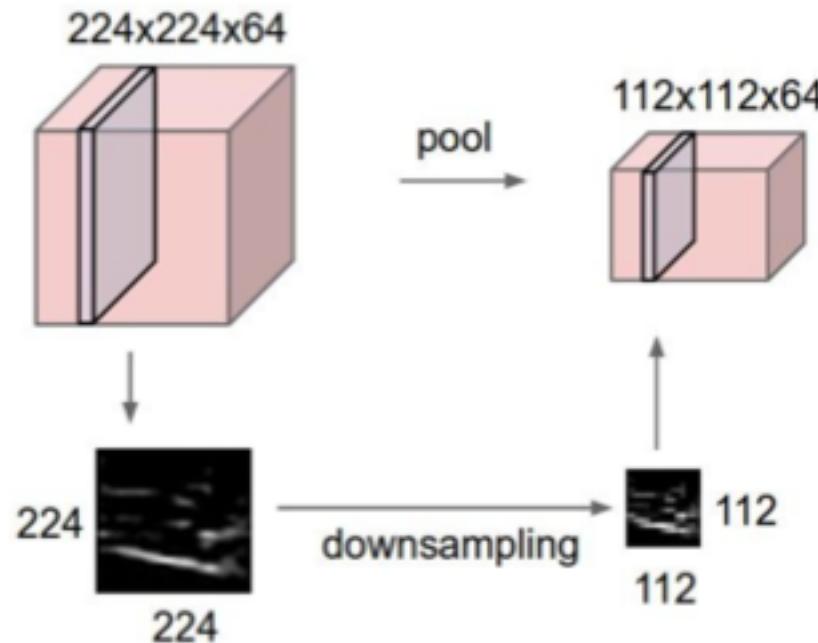
two more layers to go: POOL/FC



car
truck
airplane
ship
horse

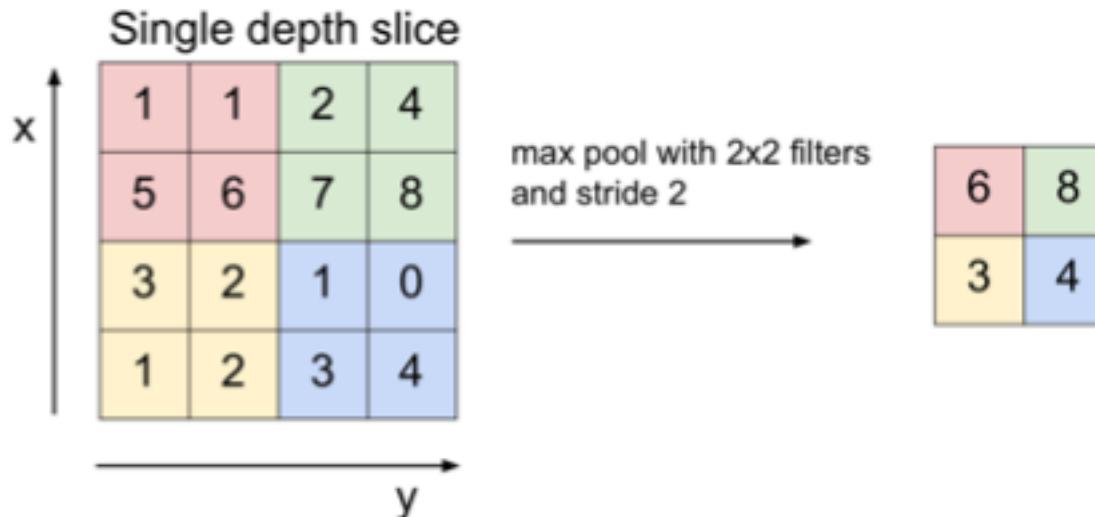
Pooling Layer

- makes the representations smaller and more manageable
- operates over each activation map independently:



Pooling Layer

- 세로·가로 방향의 공간을 줄이는 연산
 - 2x2 최대 풀링(max pooling)을 스트라이드 2로

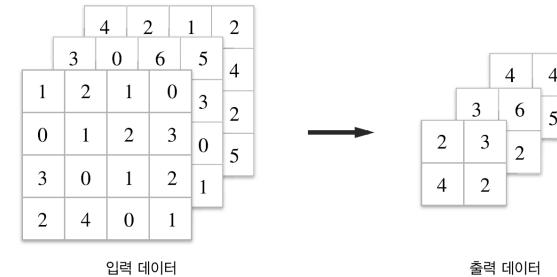


- 평균 풀링(Average pooling)도 있습니다

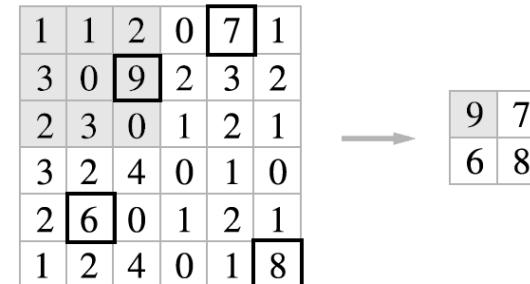
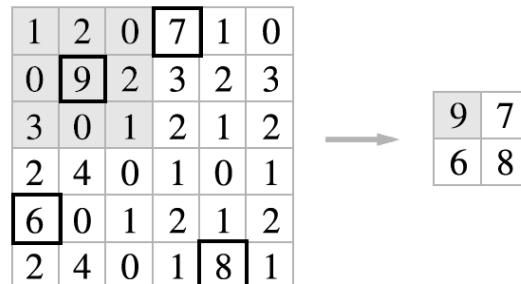
Pooling Layer

- 풀링 계층의 특징

- 학습해야 할 매개변수가 없음
- 채널 수가 변하지 않음
- 입력의 변화에 영향을 적게 받음 (강건하다)



데이터가 오른쪽으로 1칸씩 이동한 경우



Pooling 후 특징맵 크기 변화

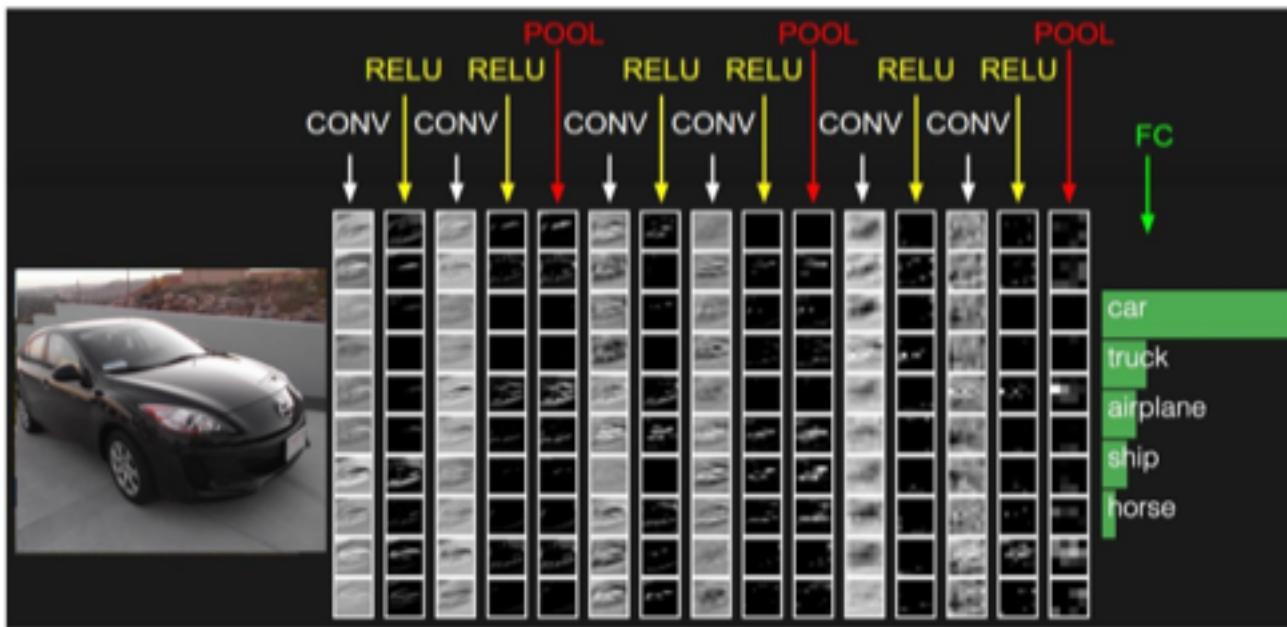


Common settings:

- Accepts a volume of size $W_1 \times H_1 \times D_1$
- Requires three hyperparameters:
 - their spatial extent F ,
 - the stride S ,
- Produces a volume of size $W_2 \times H_2 \times D_2$ where:
 - $W_2 = (W_1 - F)/S + 1$
 - $H_2 = (H_1 - F)/S + 1$
 - $D_2 = D_1$
- Introduces zero parameters since it computes a fixed function of the input
- Note that it is not common to use zero-padding for Pooling layers

Fully Connected Layer (FC layer)

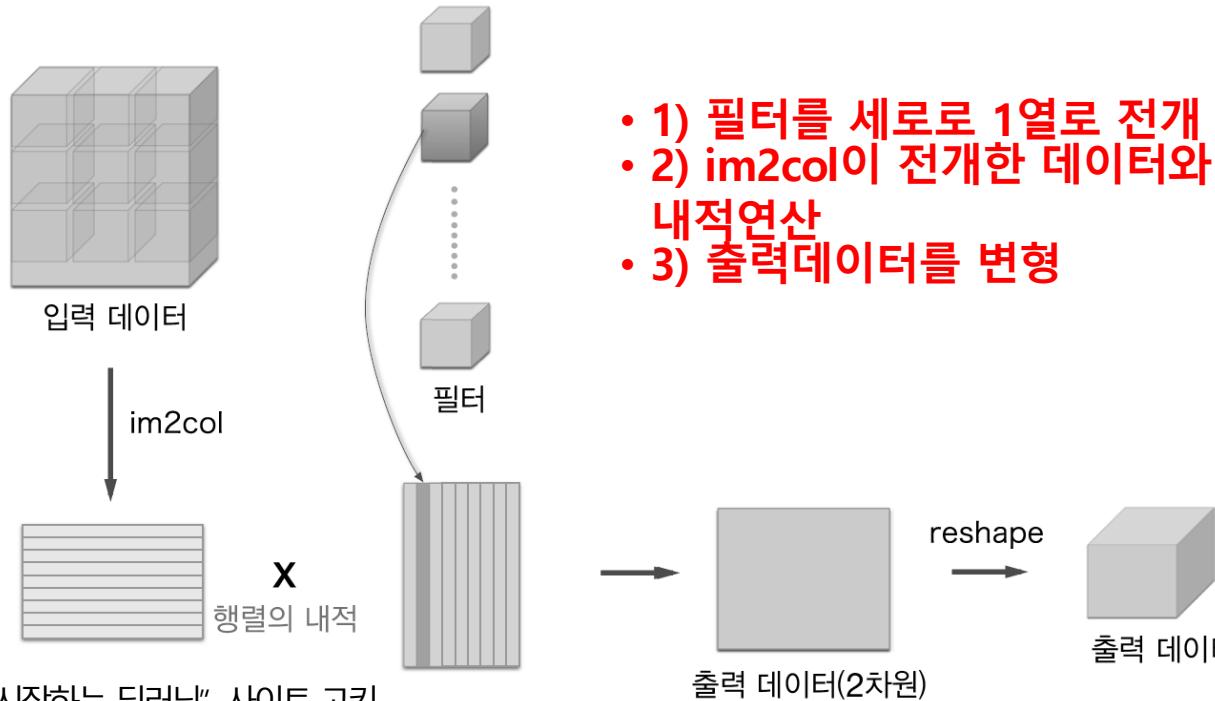
- Contains neurons that connect to the entire input volume, as in ordinary Neural Networks



Convolutional Networks 구현

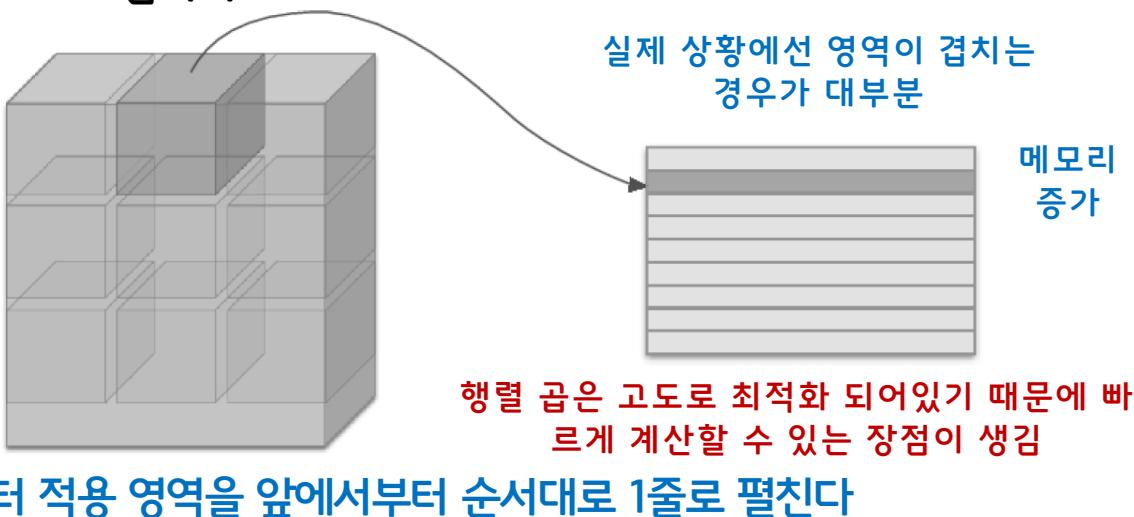


- im2col로 데이터 전개하기
 - 합성곱을 행렬 곱 연산으로 (Affine 계층연산으로)



Convolutional Networks 구현

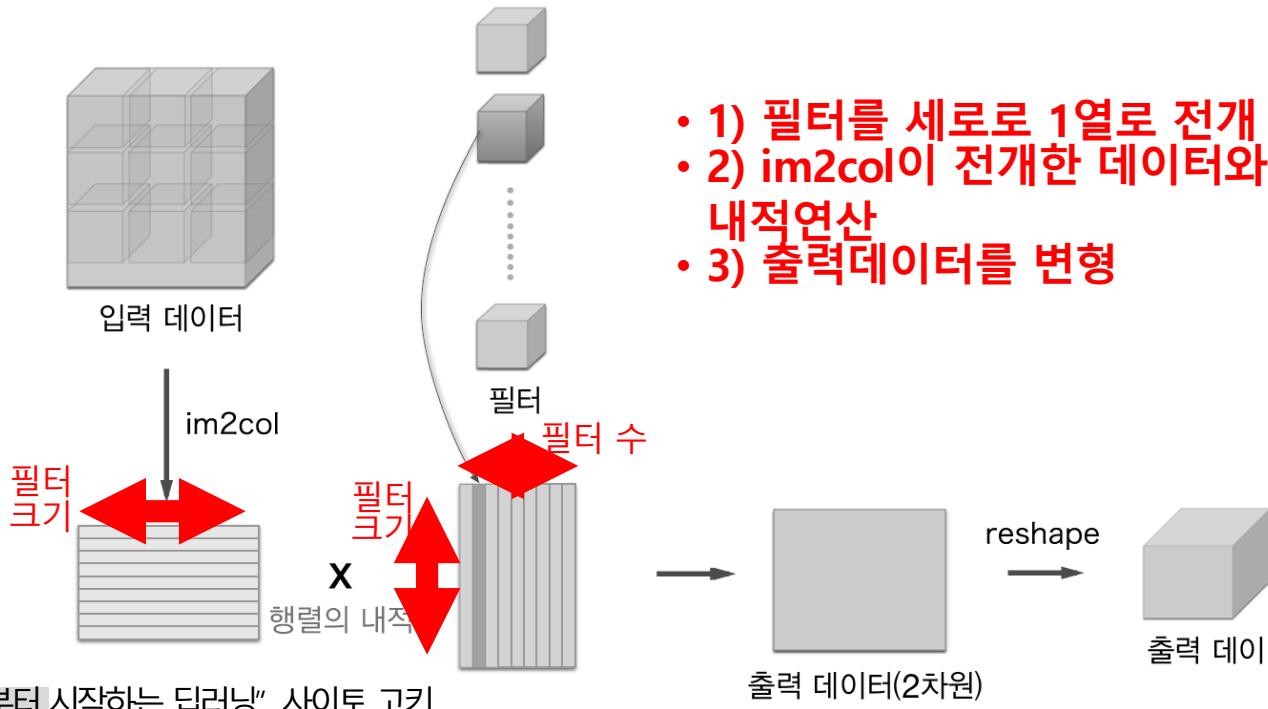
- im2col로 데이터 전개하기
 - 입력데이터에서 필터를 적용하는 영역(3차원 블록)을 한줄로 늘어 놓습니다
 - 이 전개를 필터를 적용하는 모든 영역에서 수행하는게 im2col입니다



Convolutional Networks 구현



- im2col로 데이터 전개하기
 - 합성곱을 행렬 곱 연산으로 (Affine 계층연산으로)



Convolution 레이어 구현하기

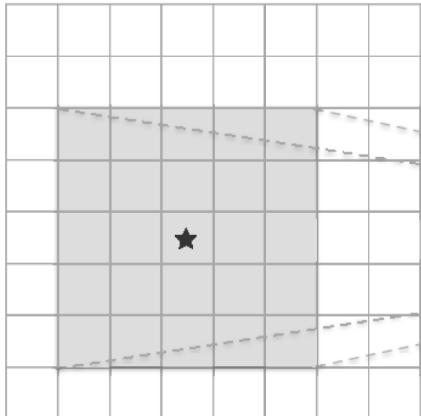
- im2col로 데이터 전개하기
 - 합성곱을 행렬 곱 연산으로 (Affine 계층연산으로)



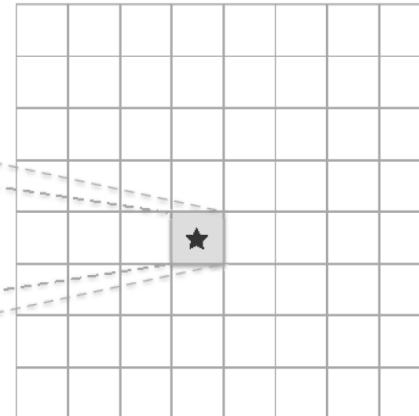
더 작은 필터

- 작은 필터로 더 깊게
 - 5x5 한번과 3x3 두번의 비교

입력 데이터

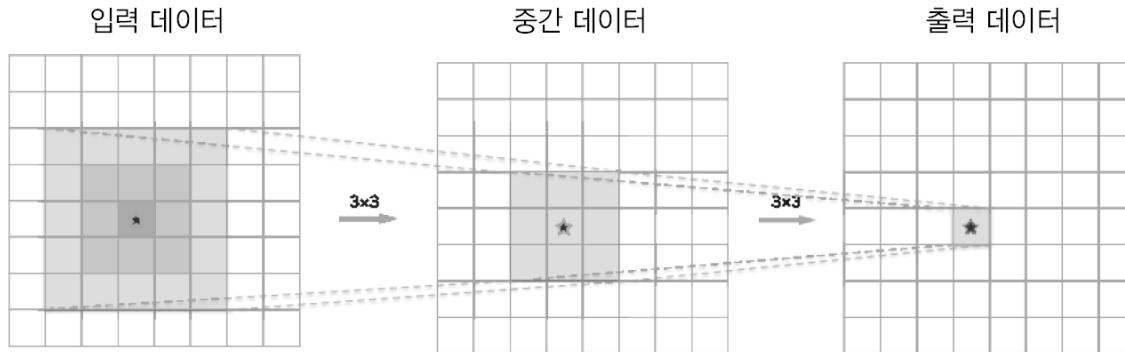


출력 데이터



더 작은 필터

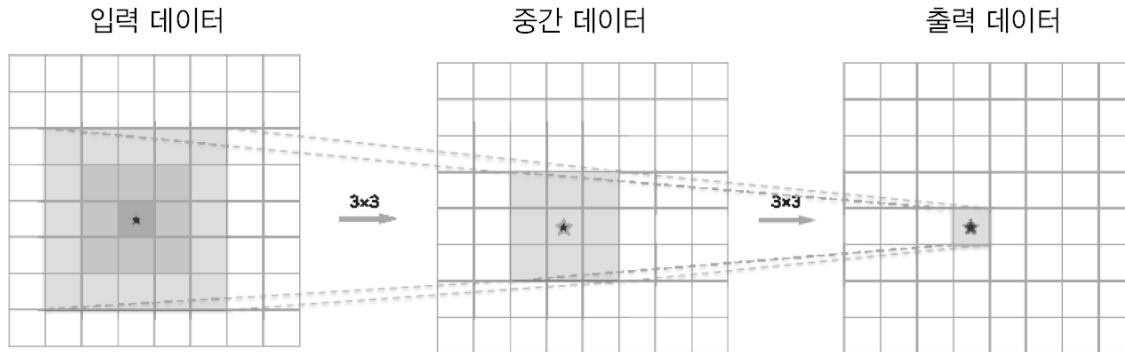
- 작은 필터로 더 깊게
 - 5x5 한번과 3x3 두번의 비교



- 5x5와 같은 크기의 영역을 처리 (receptive field)
- 층이 깊어지기에 ReLU와 같은 비선형성 추가로 표현력이 개선. 비선형 함수가 겹쳐지면 더 복잡한것도 표현 가능

더 작은 필터

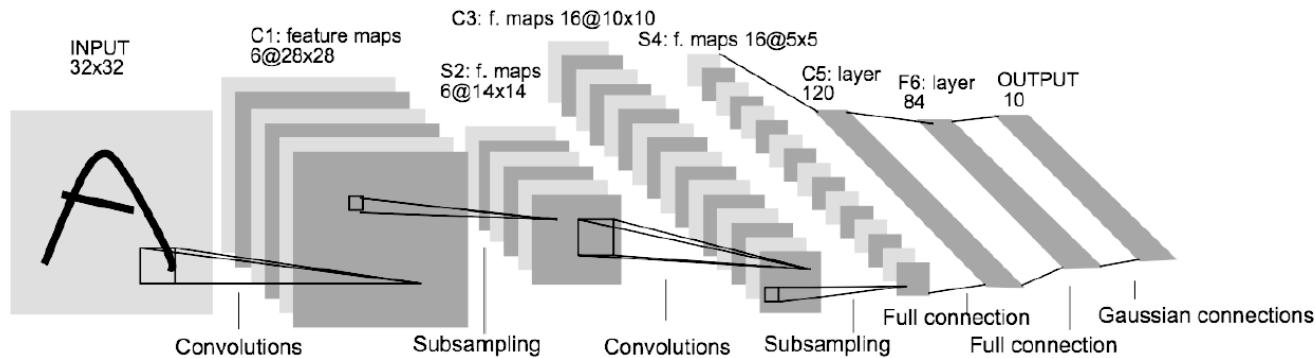
- 작은 필터로 더 깊게
 - 5x5 한번과 3x3 두번의 비교



- 매개변수 수 비교
 - 5x5 필터 1개 : 25개
 - 3x3 필터 2개 : $(3 \times 3) \times 2 = 18$ 개
 - 7x7 필터 1개 : 49개
 - 3x3 필터 3개 : $(3 \times 3) \times 3 = 27$ 개

LeNet

- LeNet (1998년)



- 현재의 CNN과의 가장 큰 차이 : 시그모이드
 - 현재는 주로 ReLU

요약

- ConvNet은 CONV, POOL, FC 레이어로 구성
- 작은 필터로 더 깊은 네트워크를 구성하는게 트렌드
- FC레이어를 제거하는게 트렌드
 - 추후 영상처리 부분에서 다루게 됨
- 전형적 구조
 - $[(\text{CONV-RELU})^*N - \text{POOL}]^*M - (\text{FC-RELU})^*K$, SOFTMAX
 - N : 보통 ~5, M : 크게, $0 \leq K \leq 2$
- 최근엔 전형적인 구조 보다 다양한 CNN 모델 변형들이 많아짐



박 은 수 Research Director

E-mail : es.park@modulabs.co.kr